

Unit 1: Simple Neural Networks

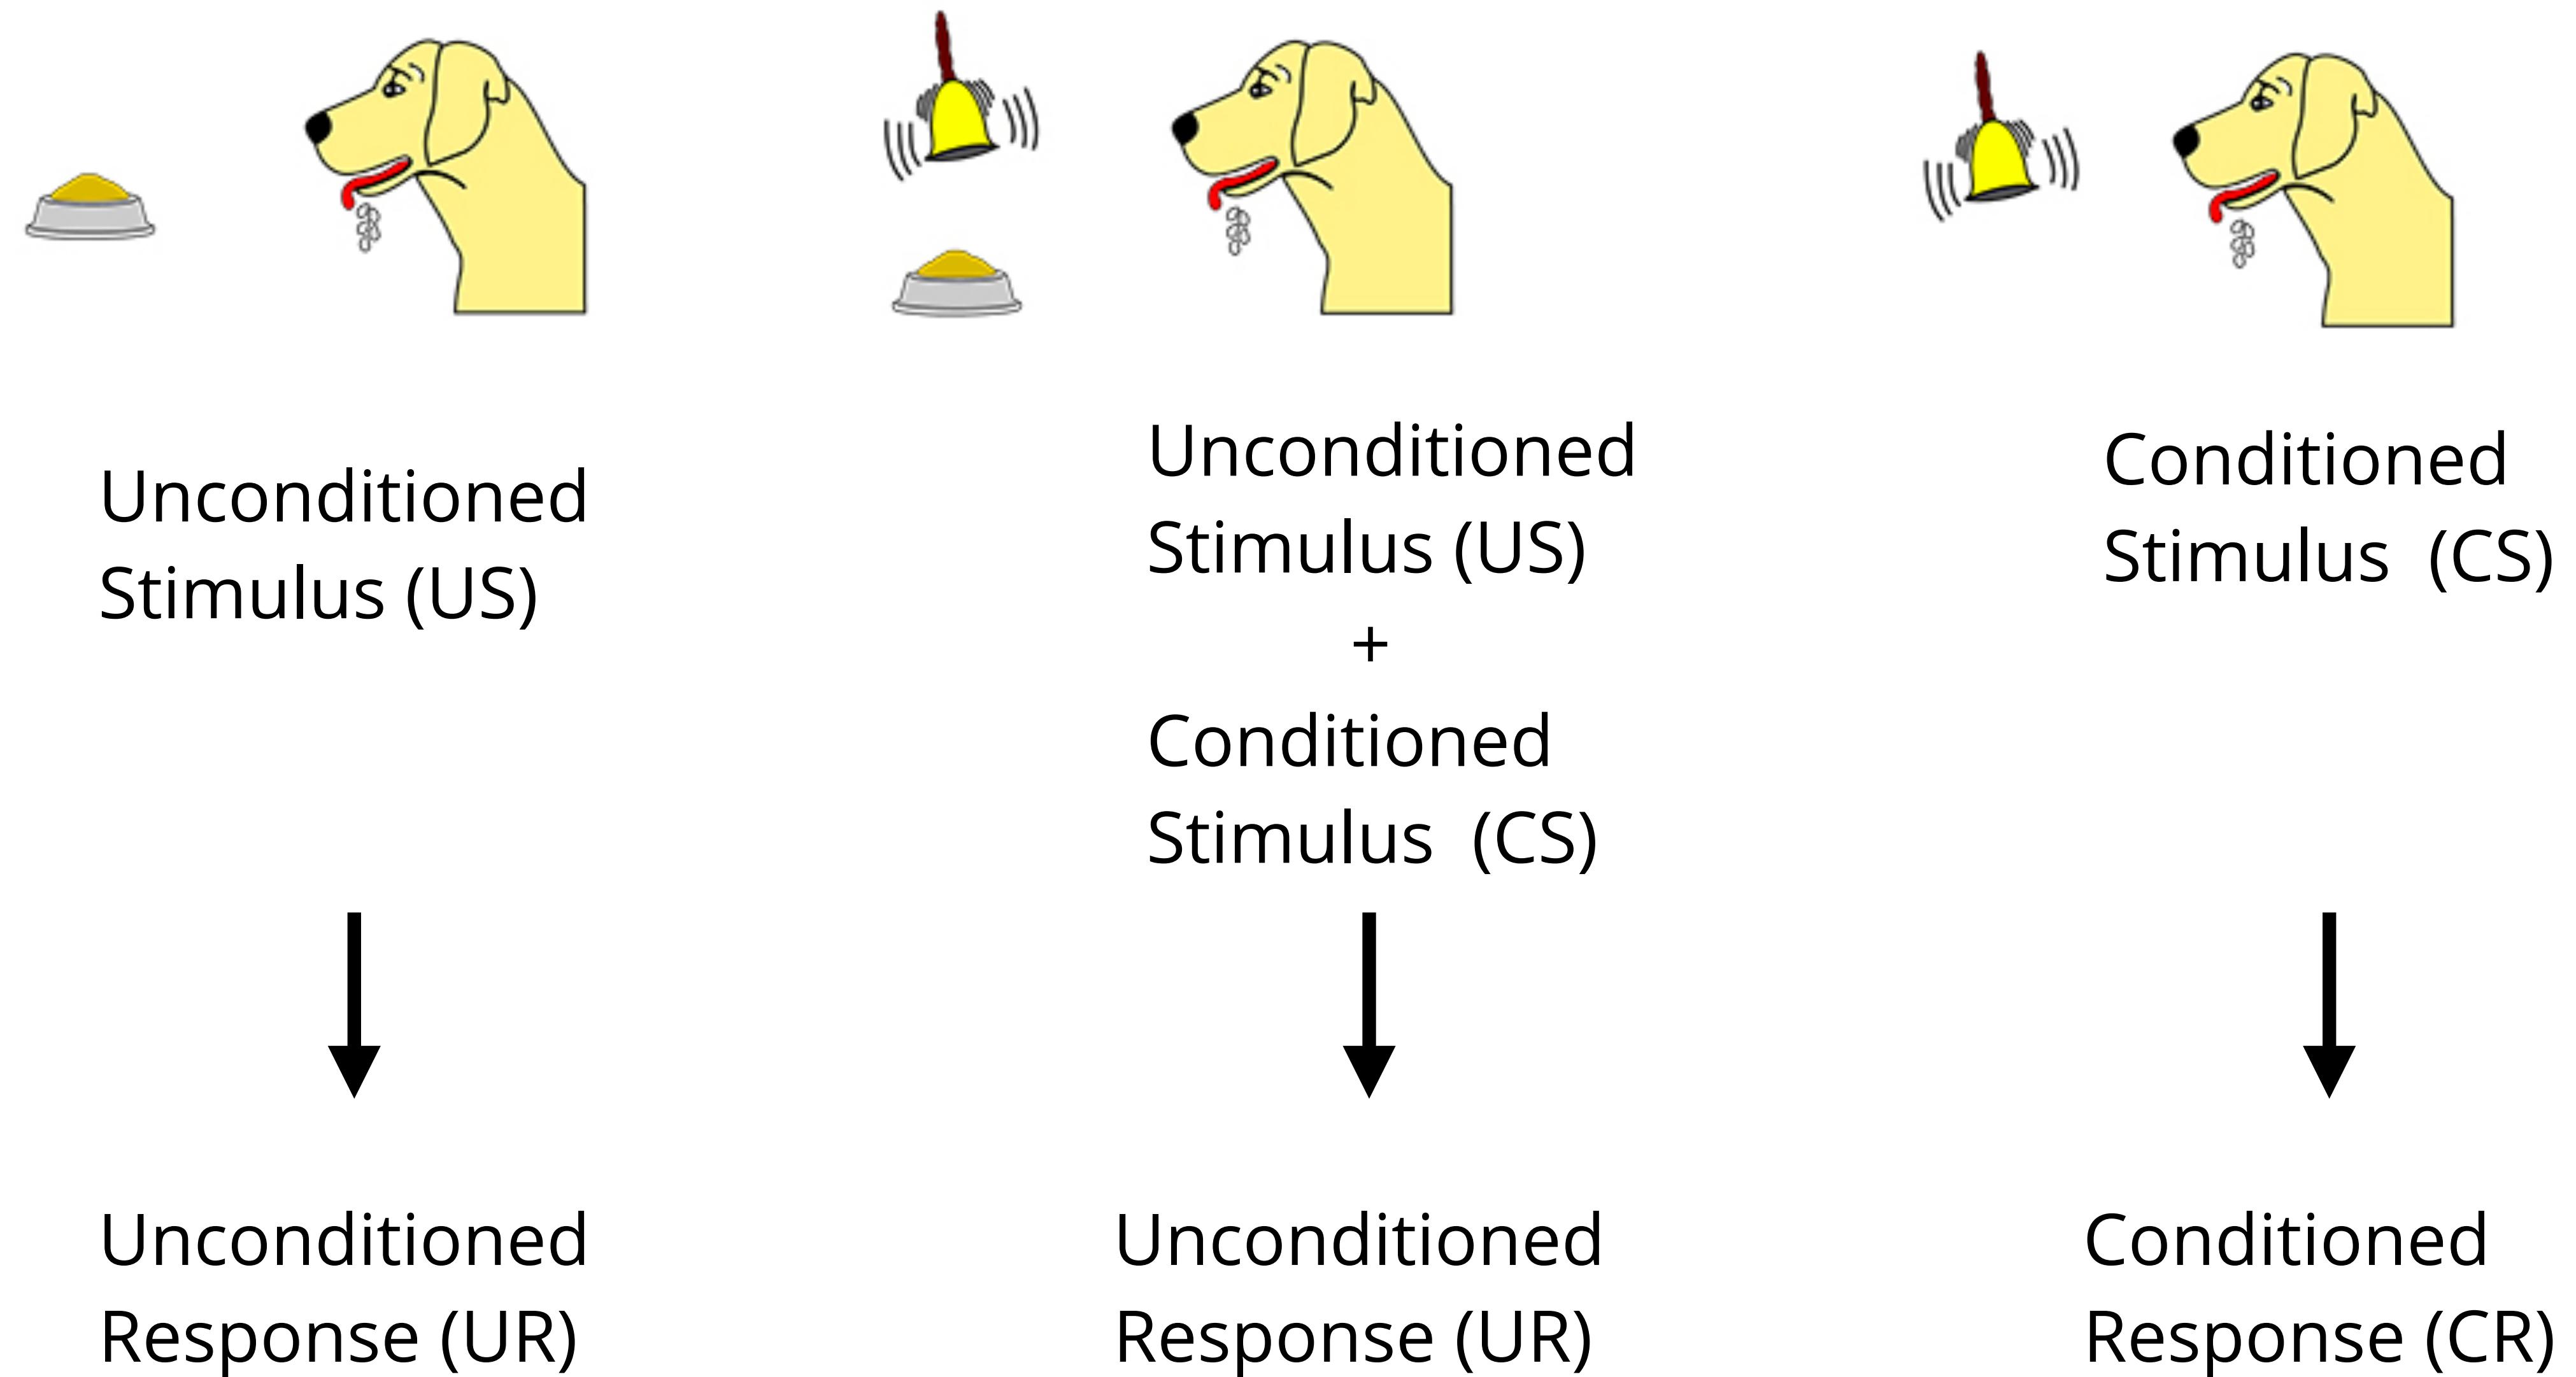
3. Perceptrons

9/15/2020

Perceptrons

- 1. Simple neural networks generalize the Rescorla-Wagner model of associative learning**
- 2. Perceptrons are general-purpose linear classifiers.
They can solve lots of problems**
- 3. But they can't solve all problems...**

Classical conditioning



The Rescorla-Wagner model: Learning is prediction error

$$P(\text{cheese}) = V_{total}$$

$$\Delta V = \alpha \cdot (\lambda - V_{total})$$

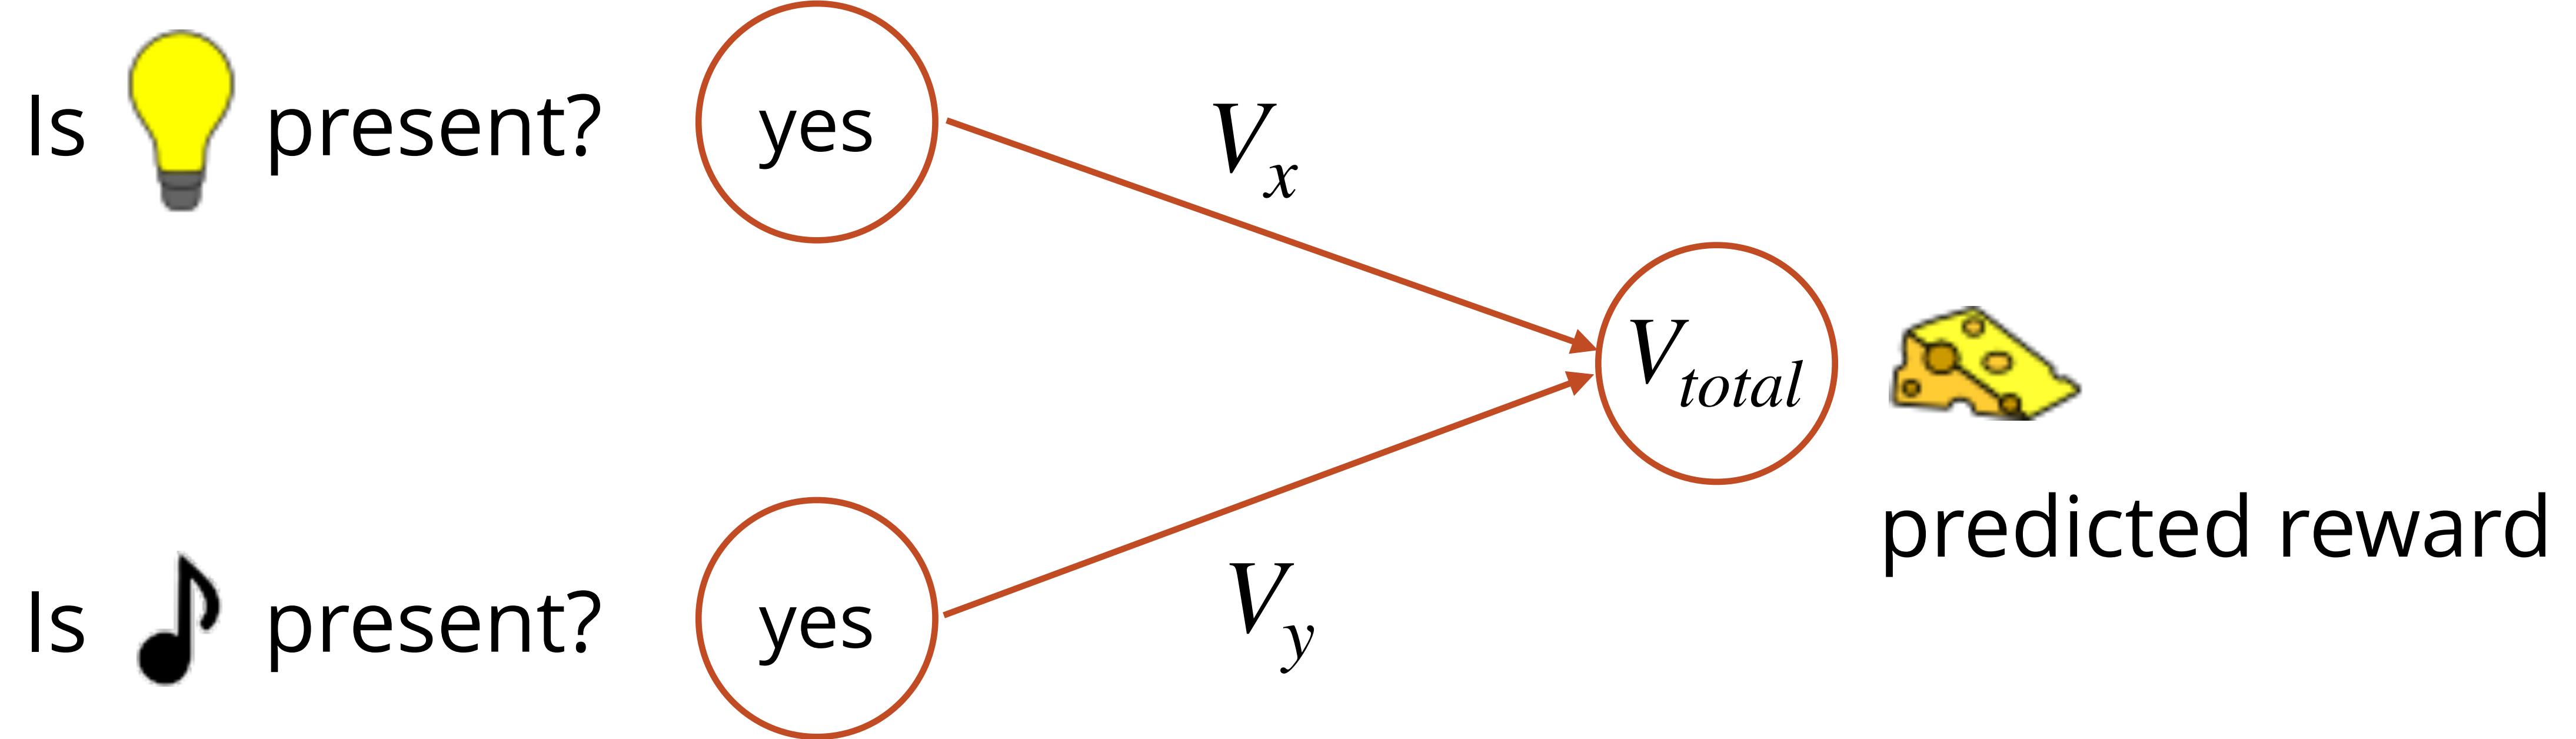
On each trial, the rat predicts whether or not it will get cheese

This prediction comes from the combination of all cues

After each trial, update predictions for each cue

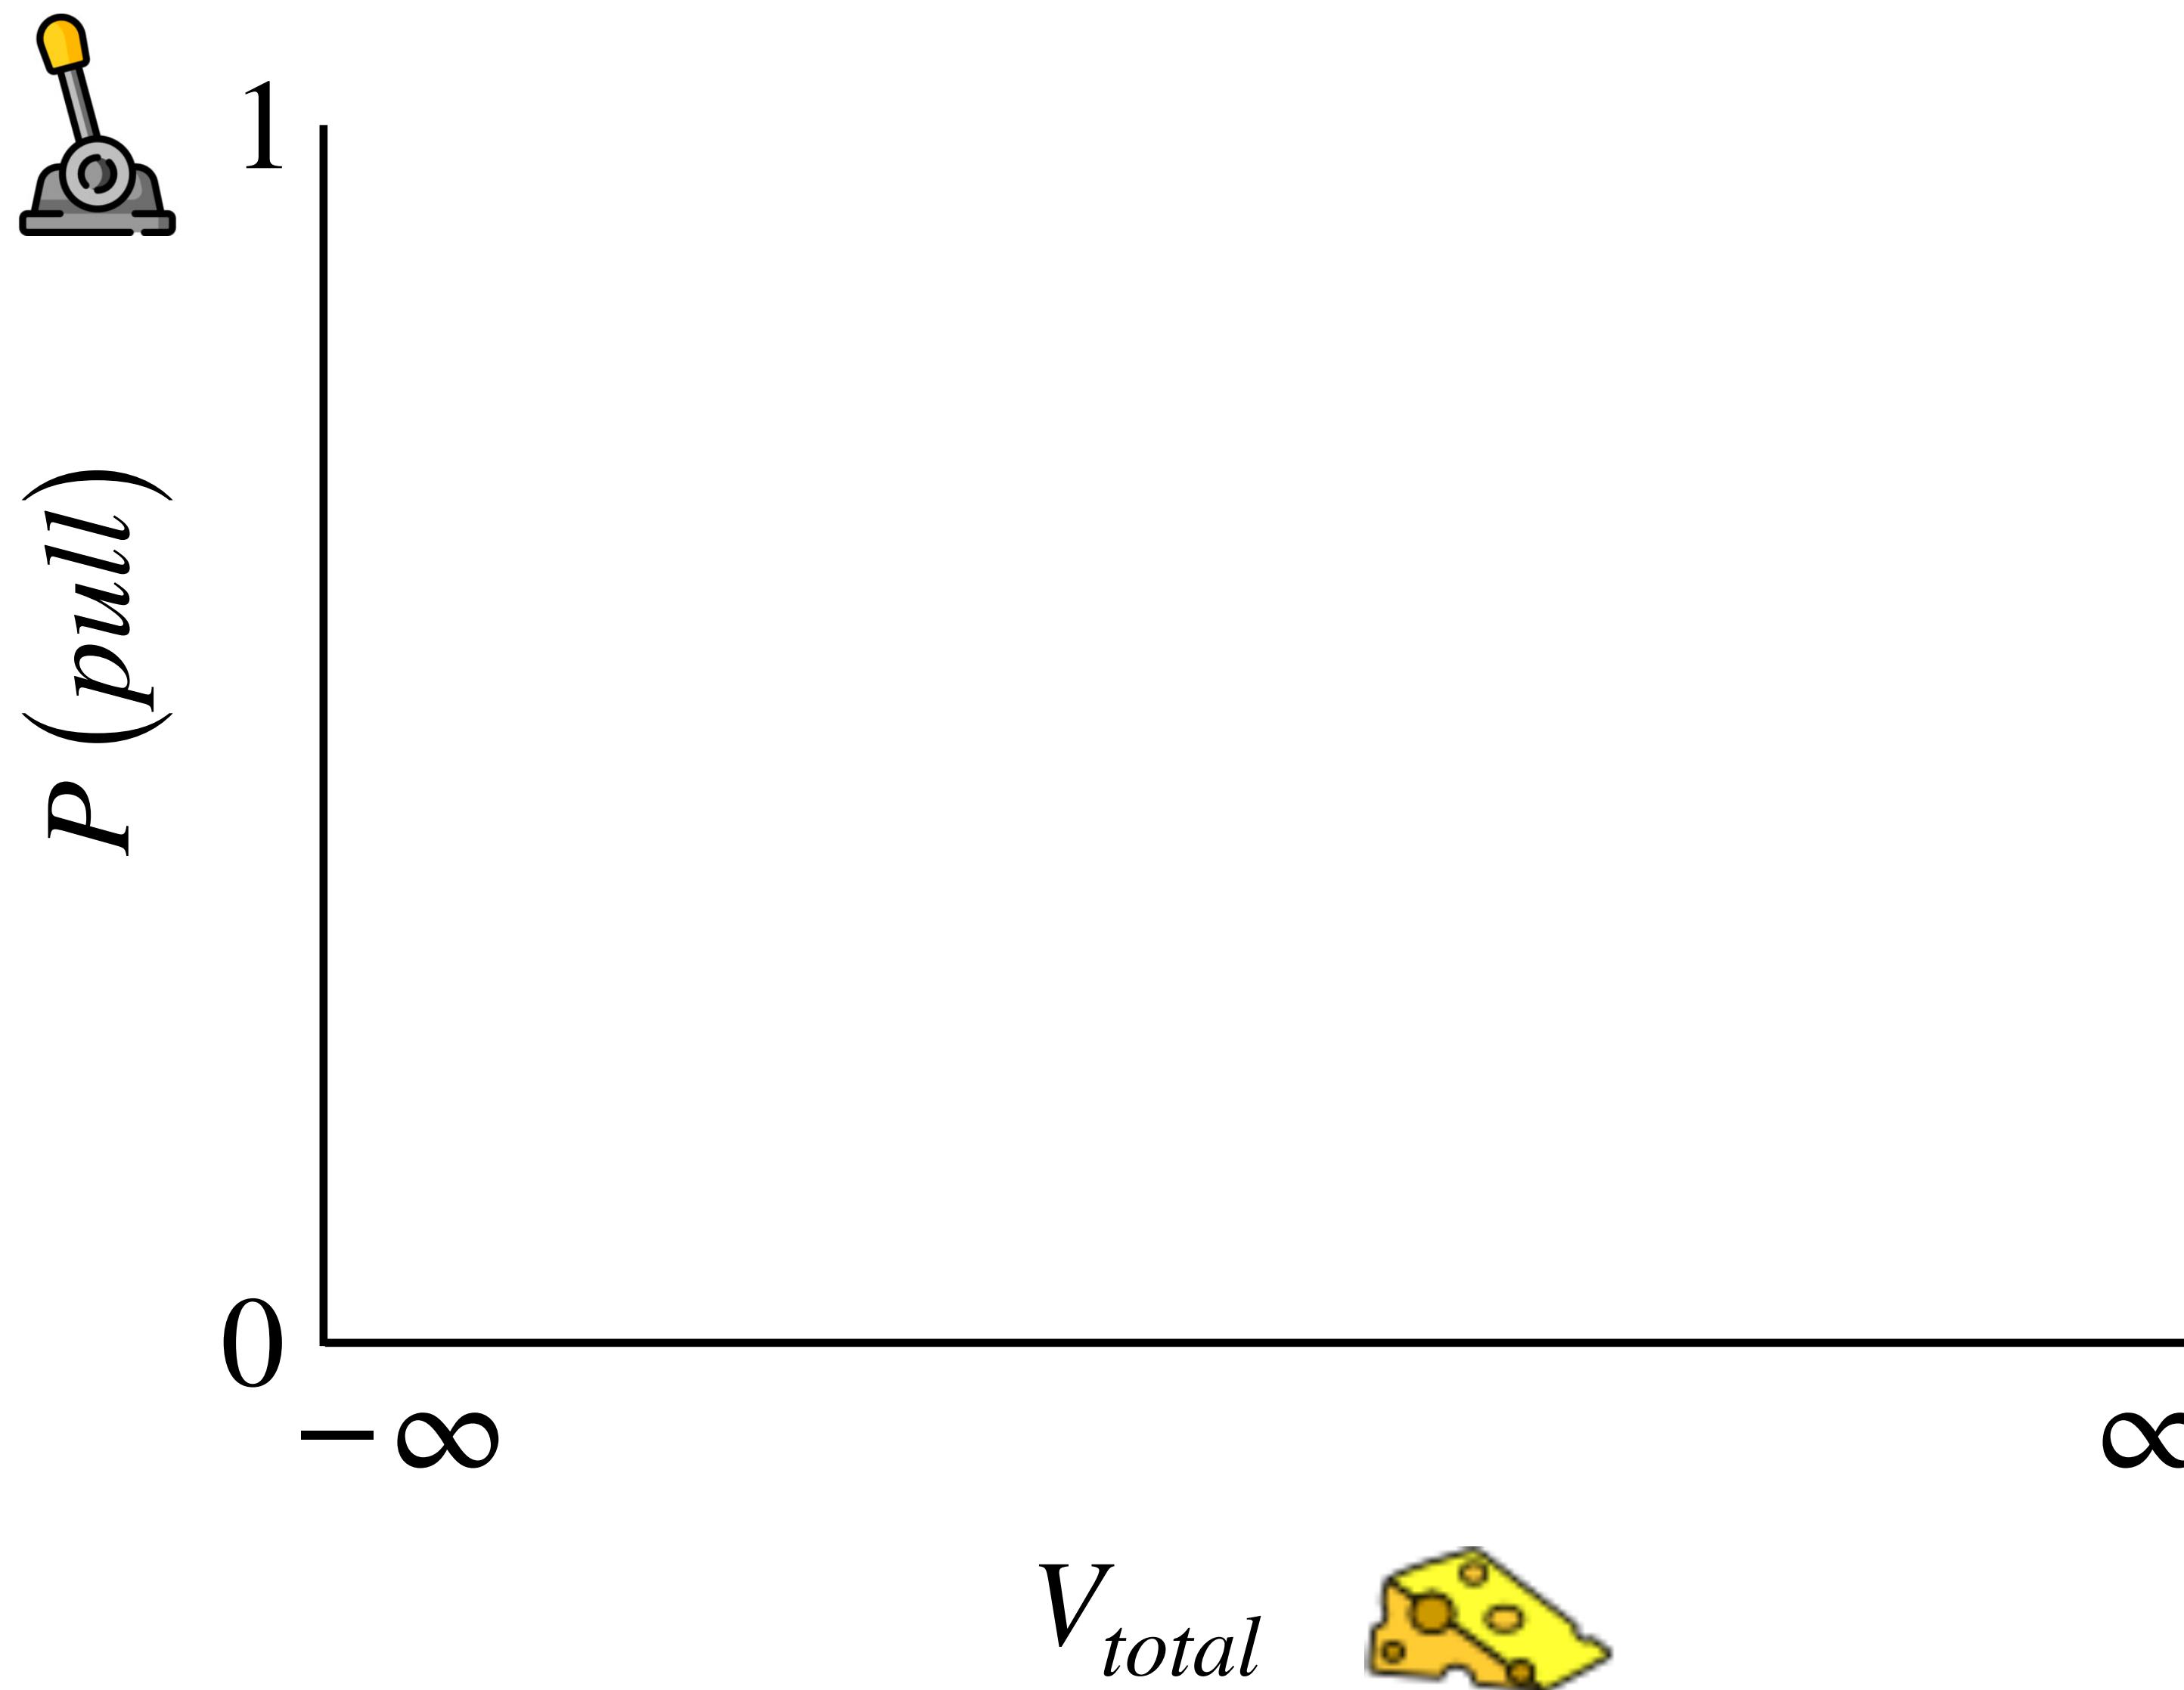
- If the rat gets cheese—but *didn't expect* cheese—
increase prediction for each cue
- If the rat doesn't get cheese—but *expected* cheese—
decrease prediction for each cue
- Otherwise, don't change anything

A network representation of Rescorla-Wagner

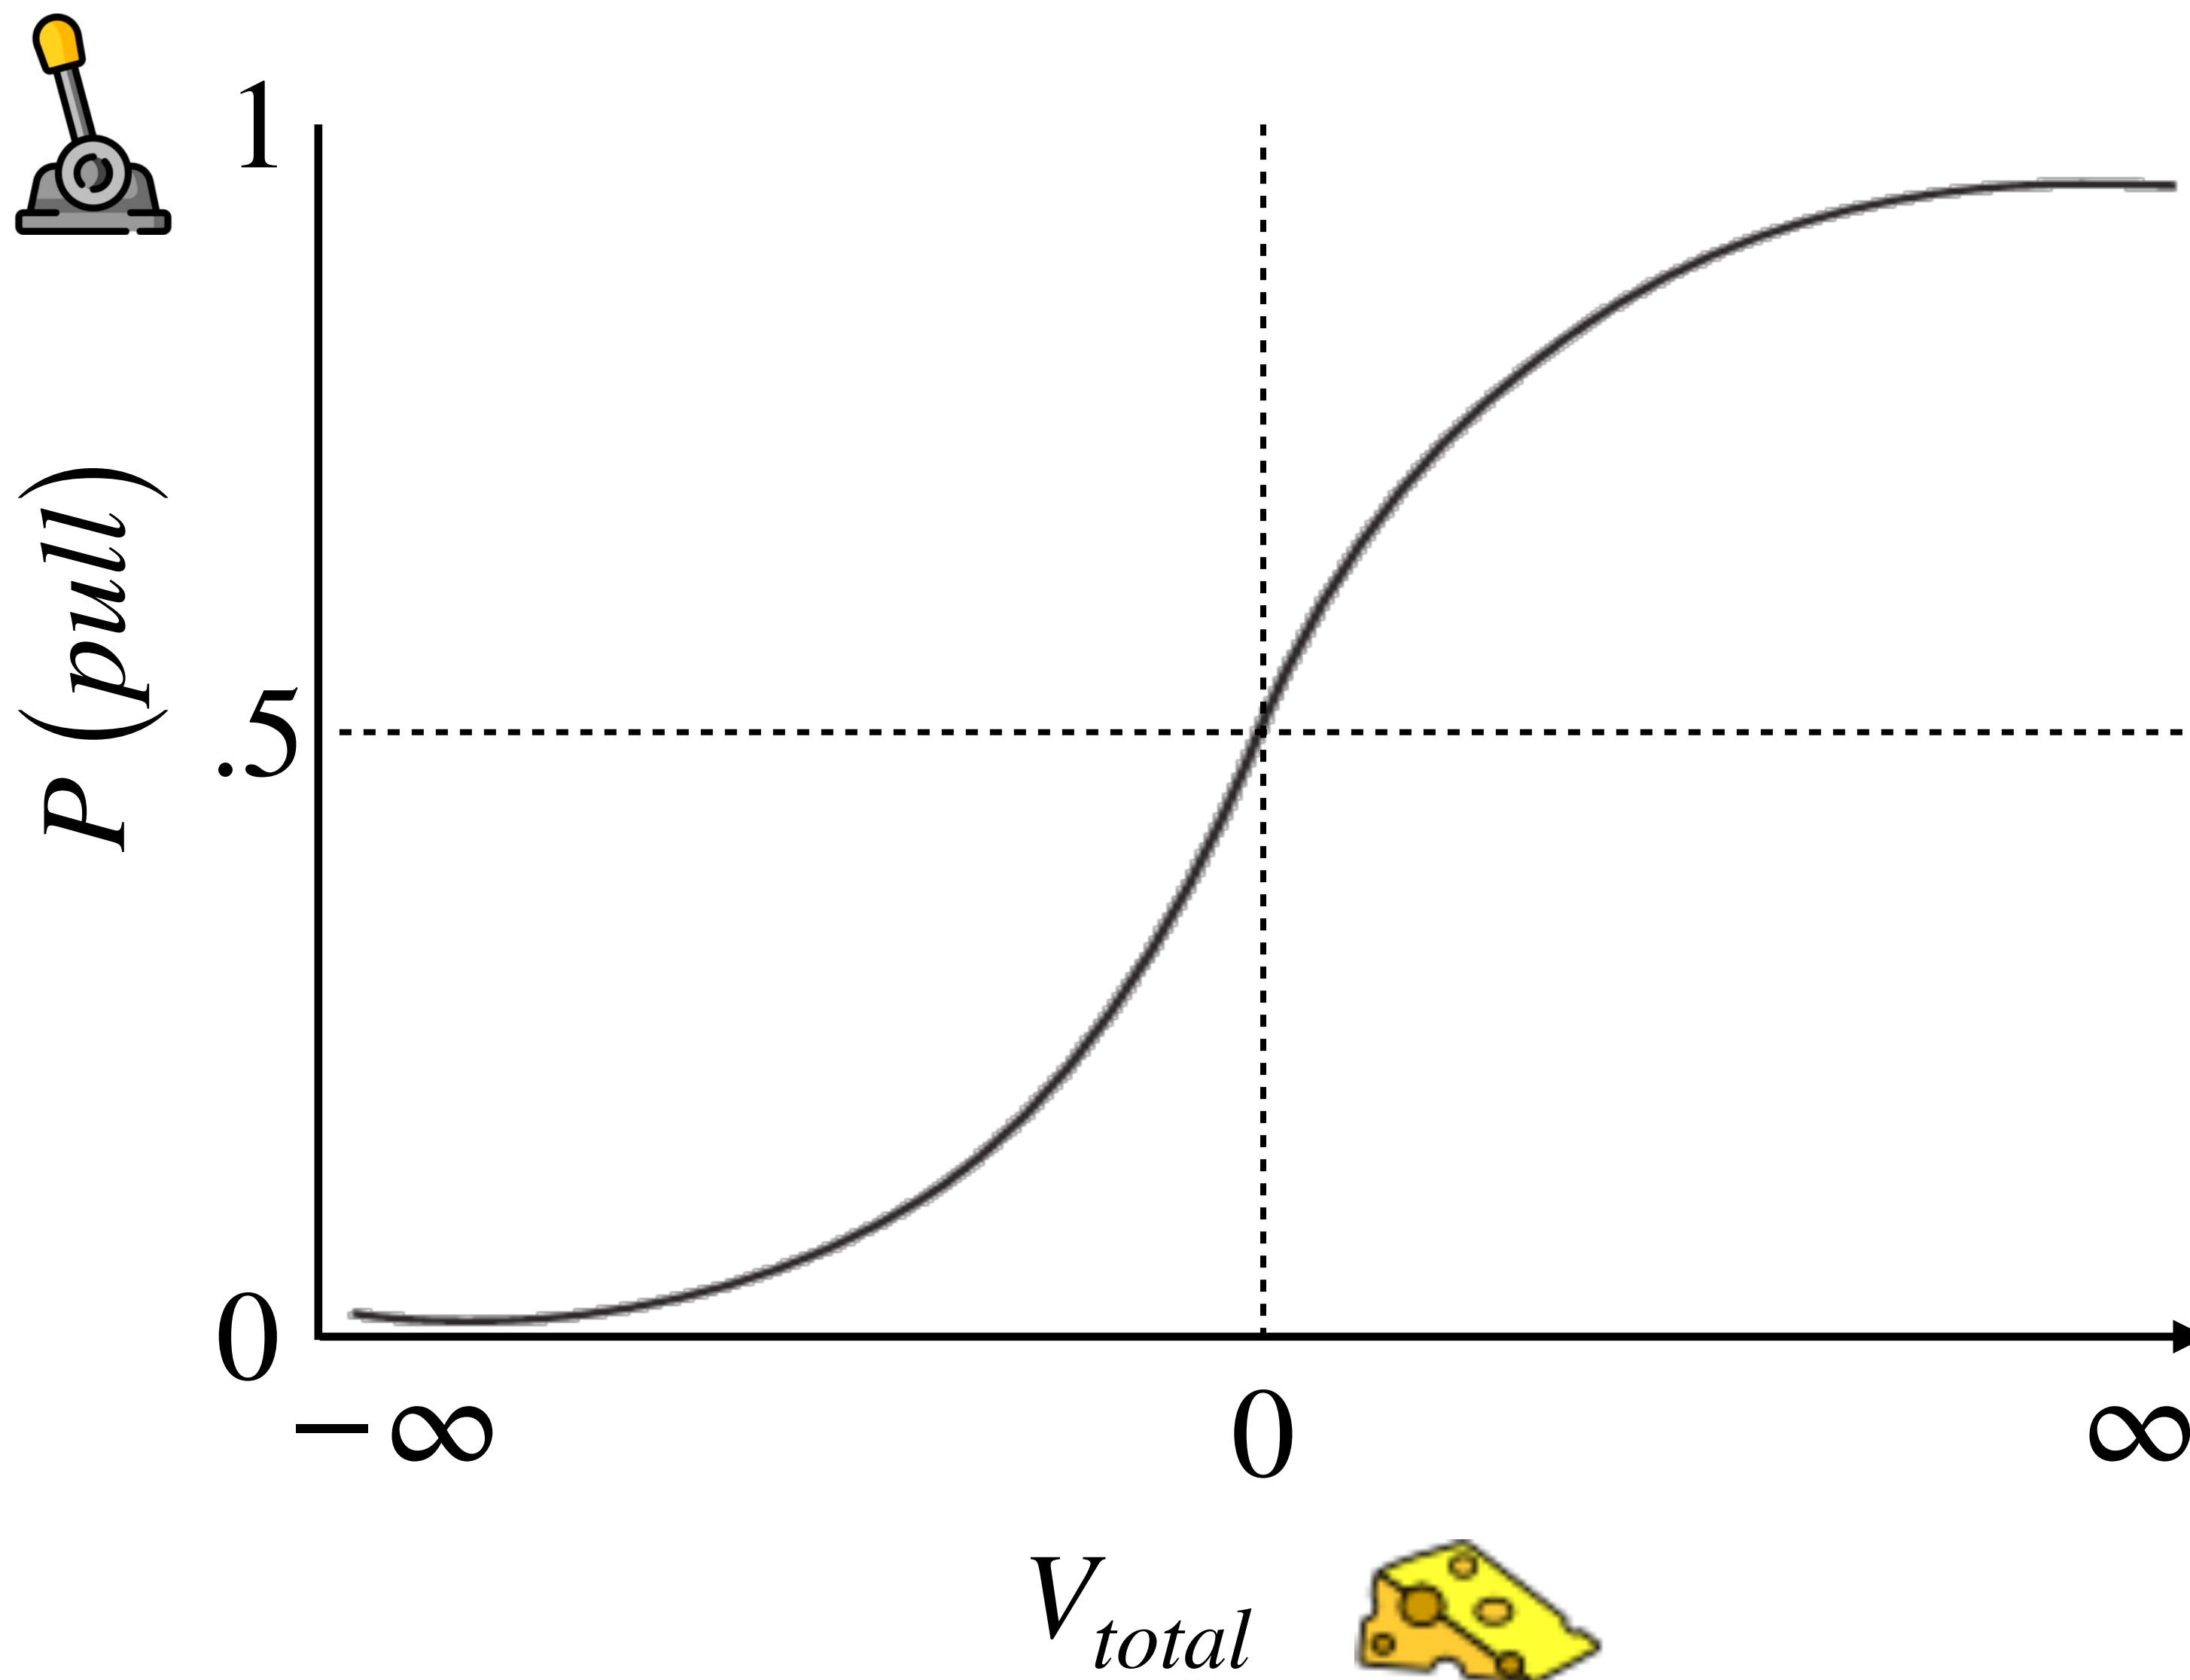


Like Ramscar et al. (2010)

Modification 1: Connecting prediction to action with a squashing function

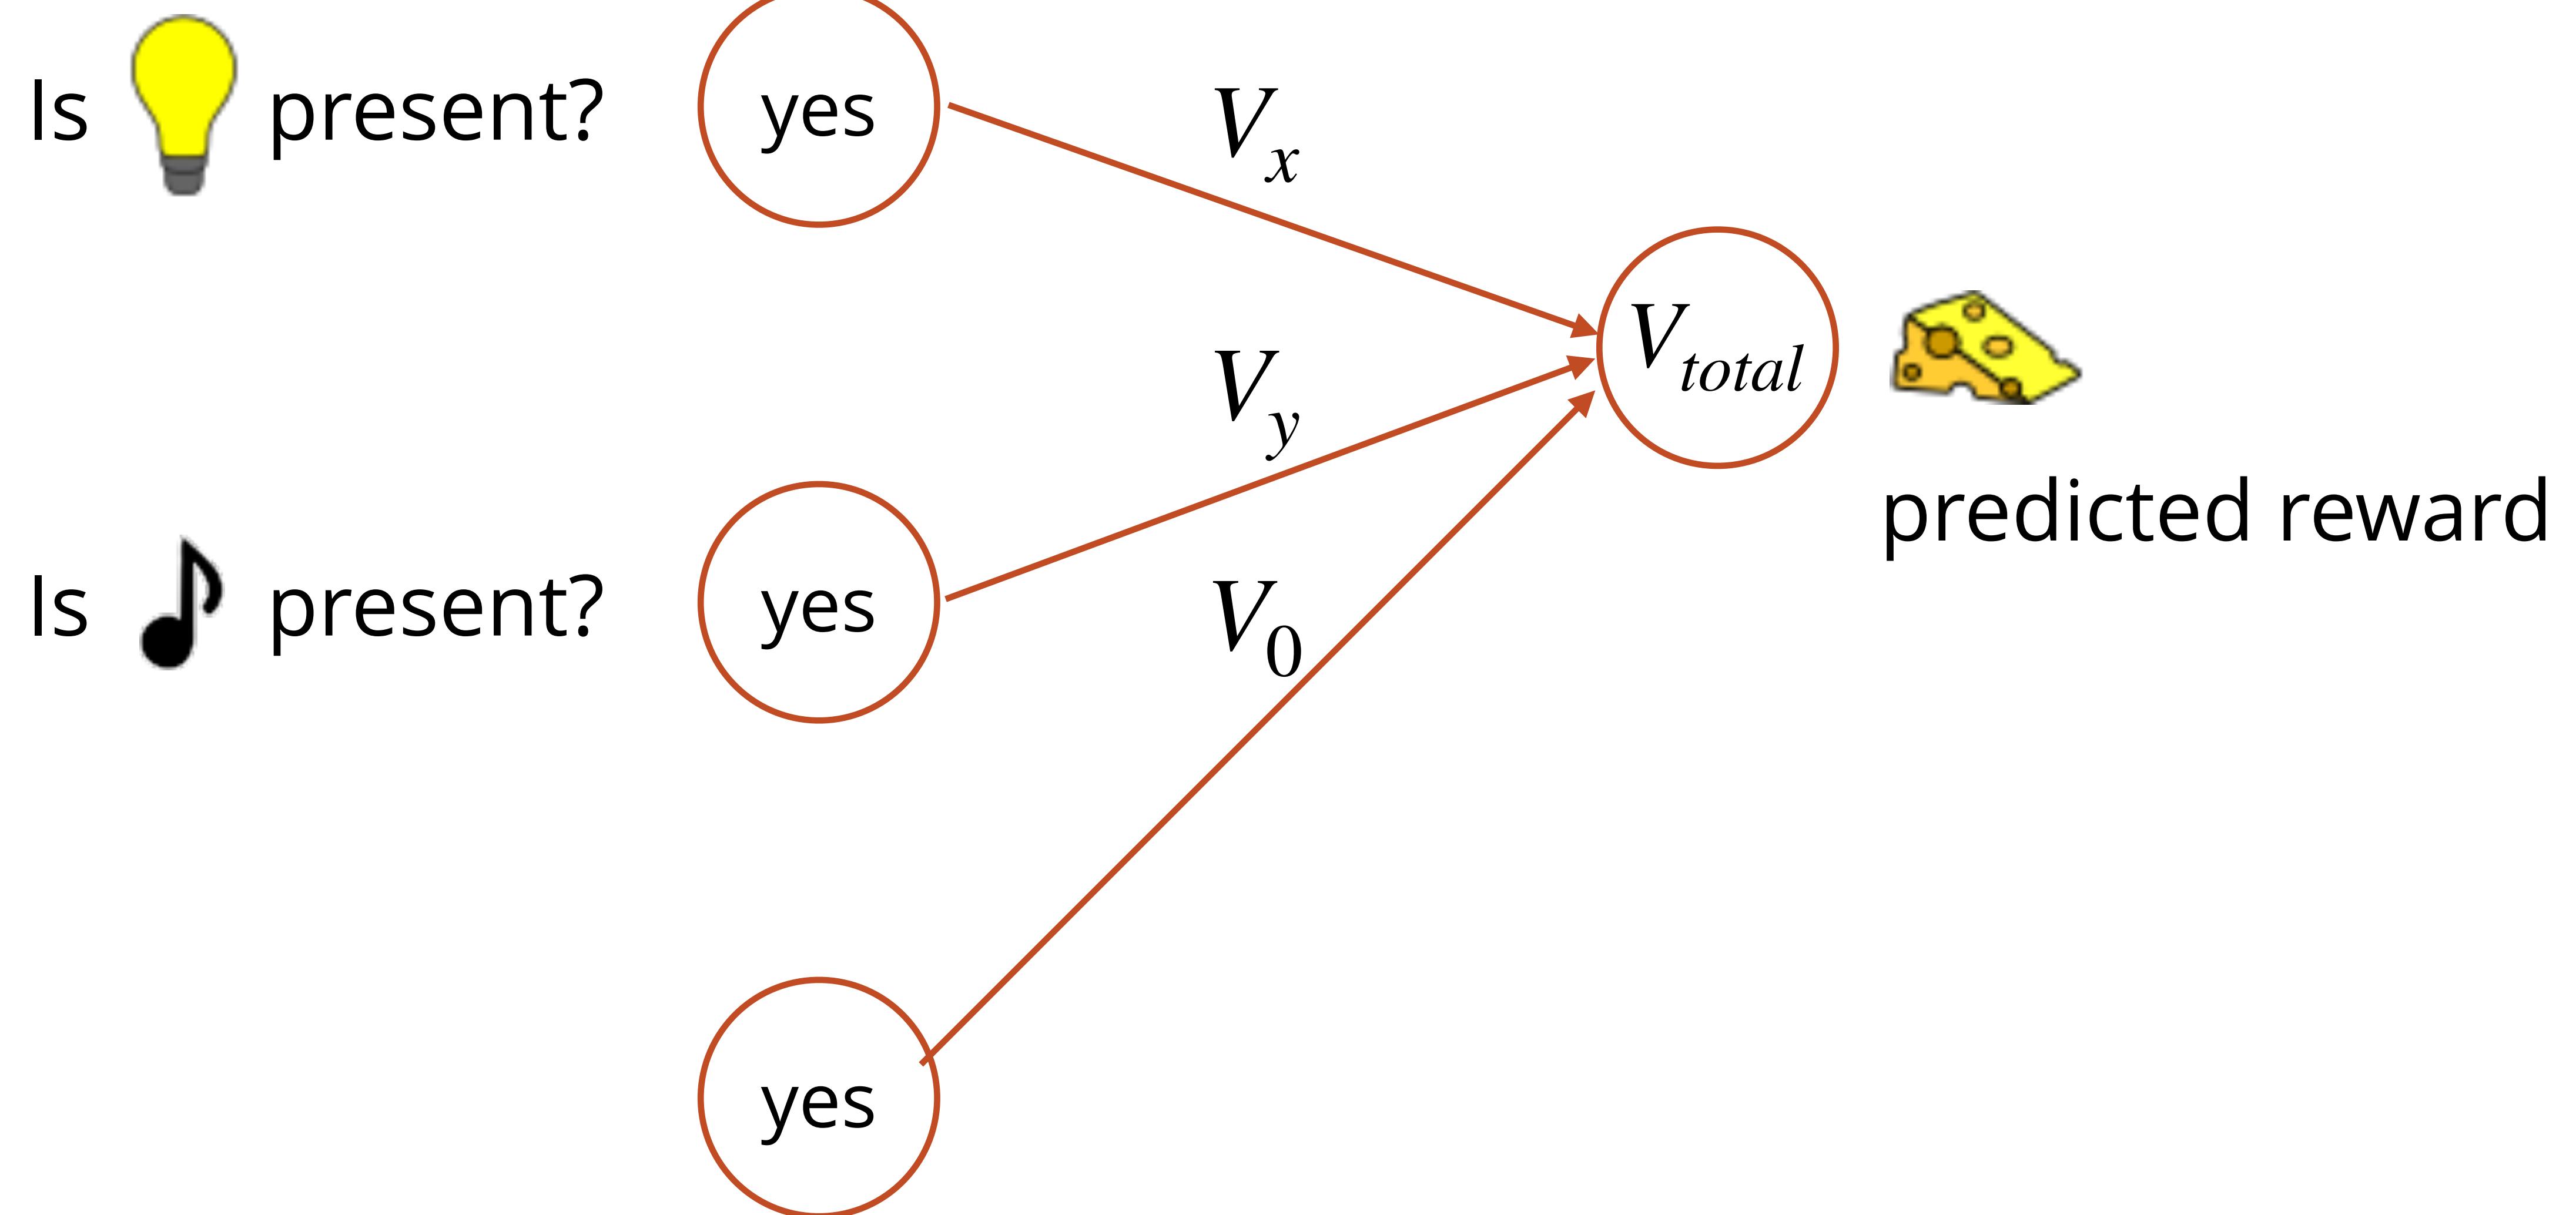


The sigmoid (logistic) function

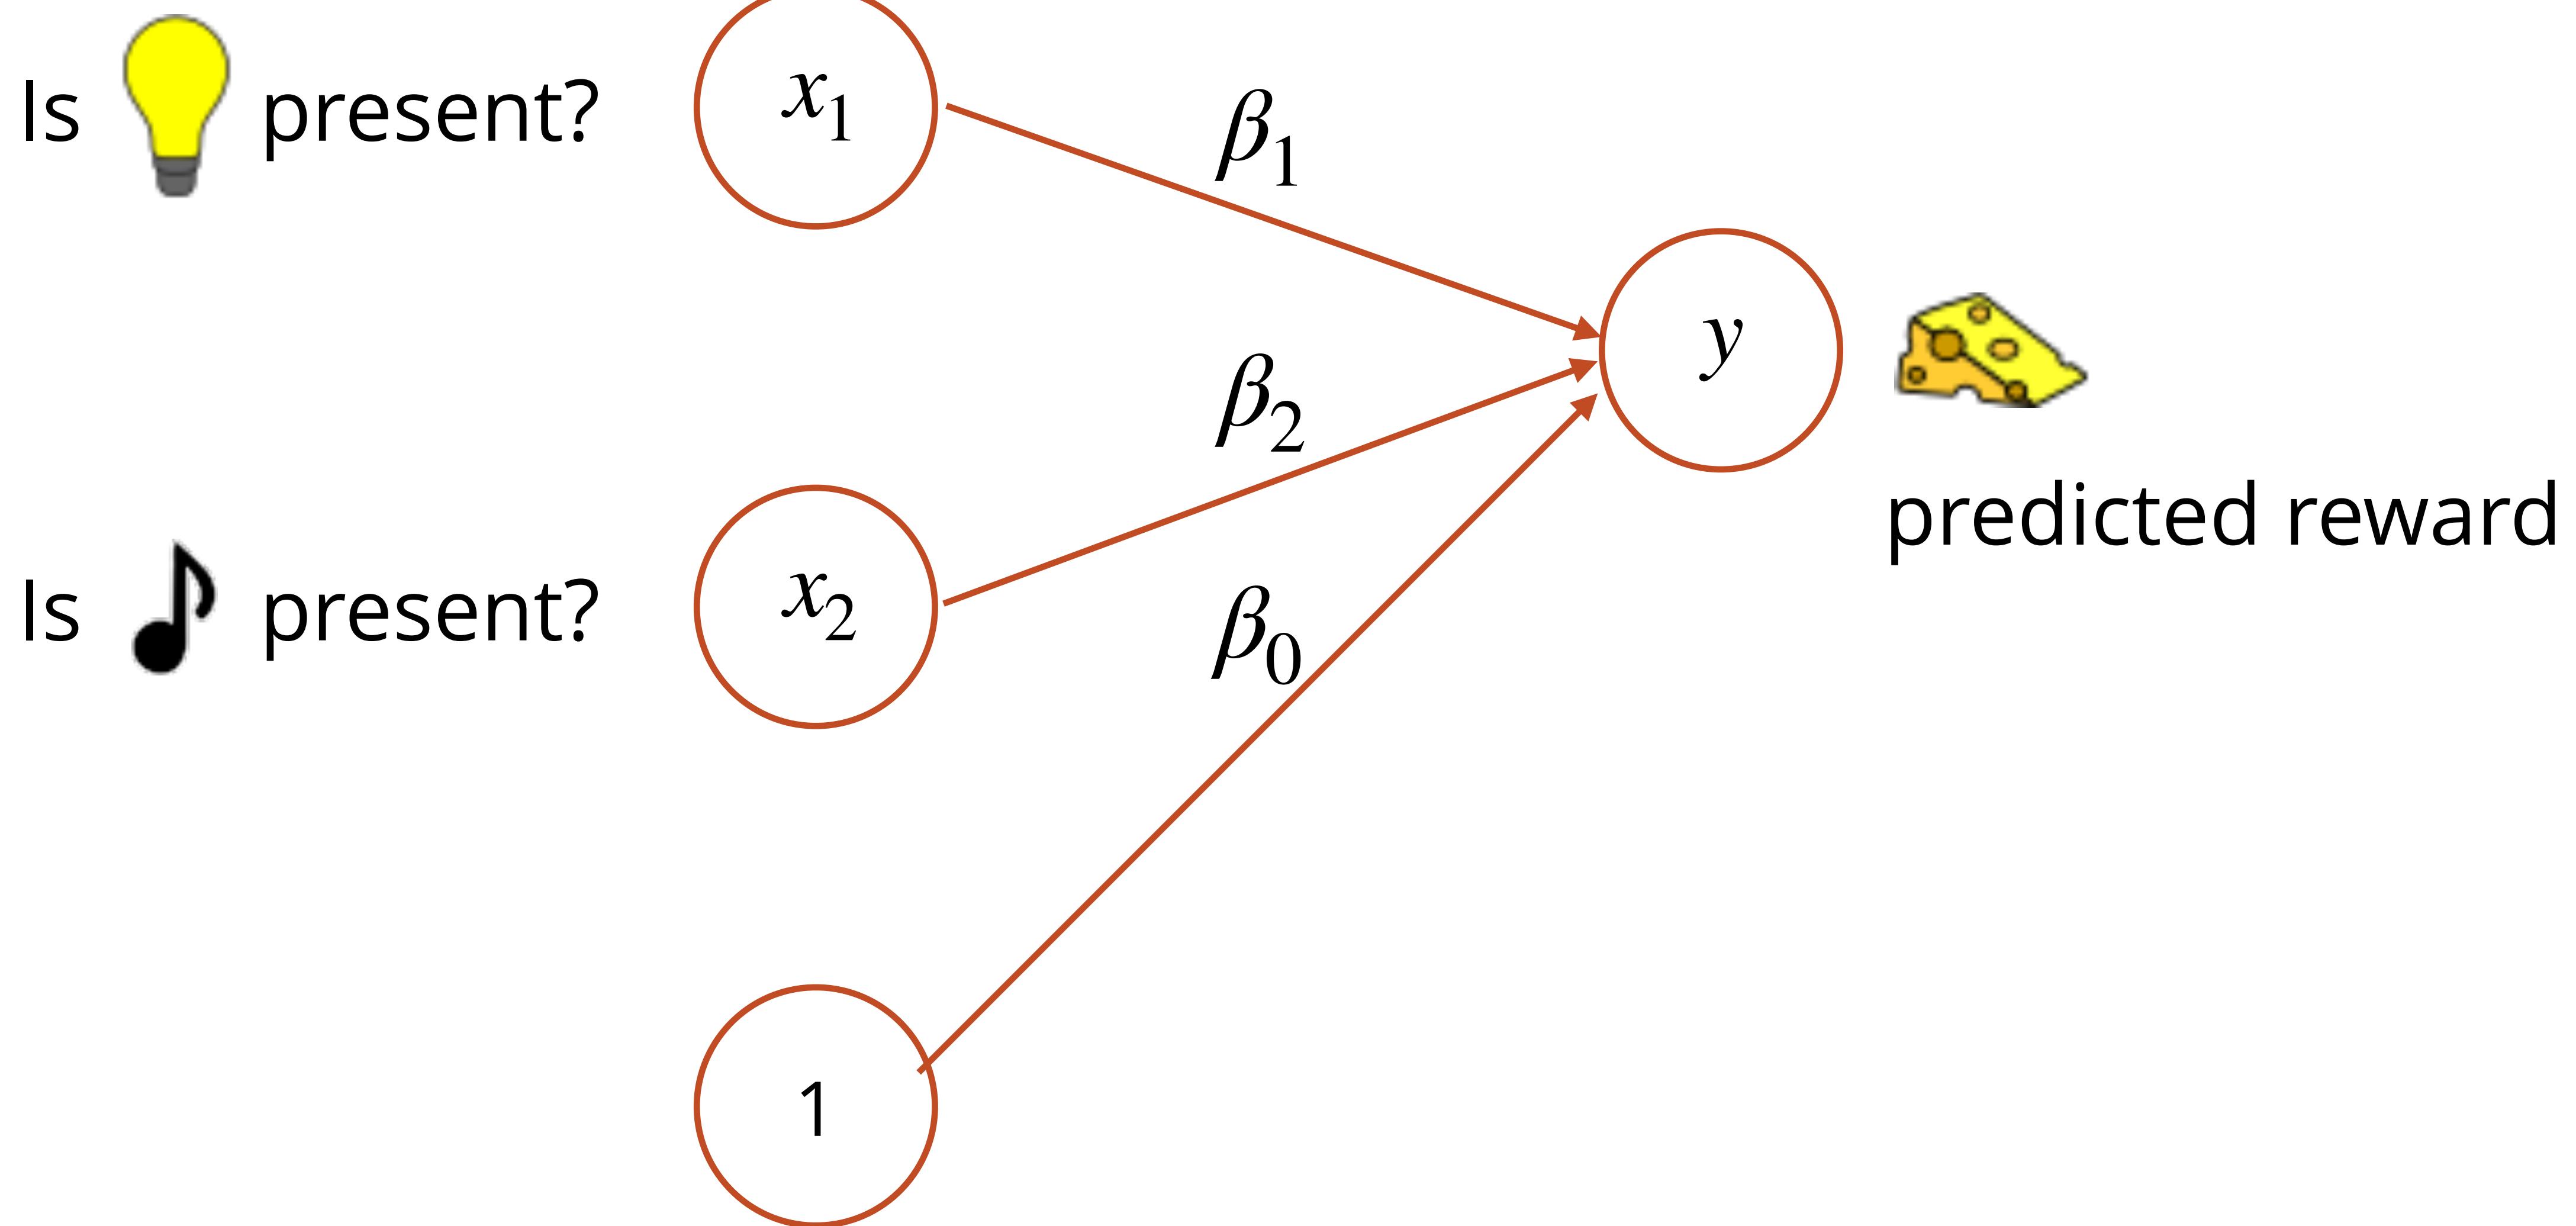


$$y = \frac{1}{1 - e^x}$$

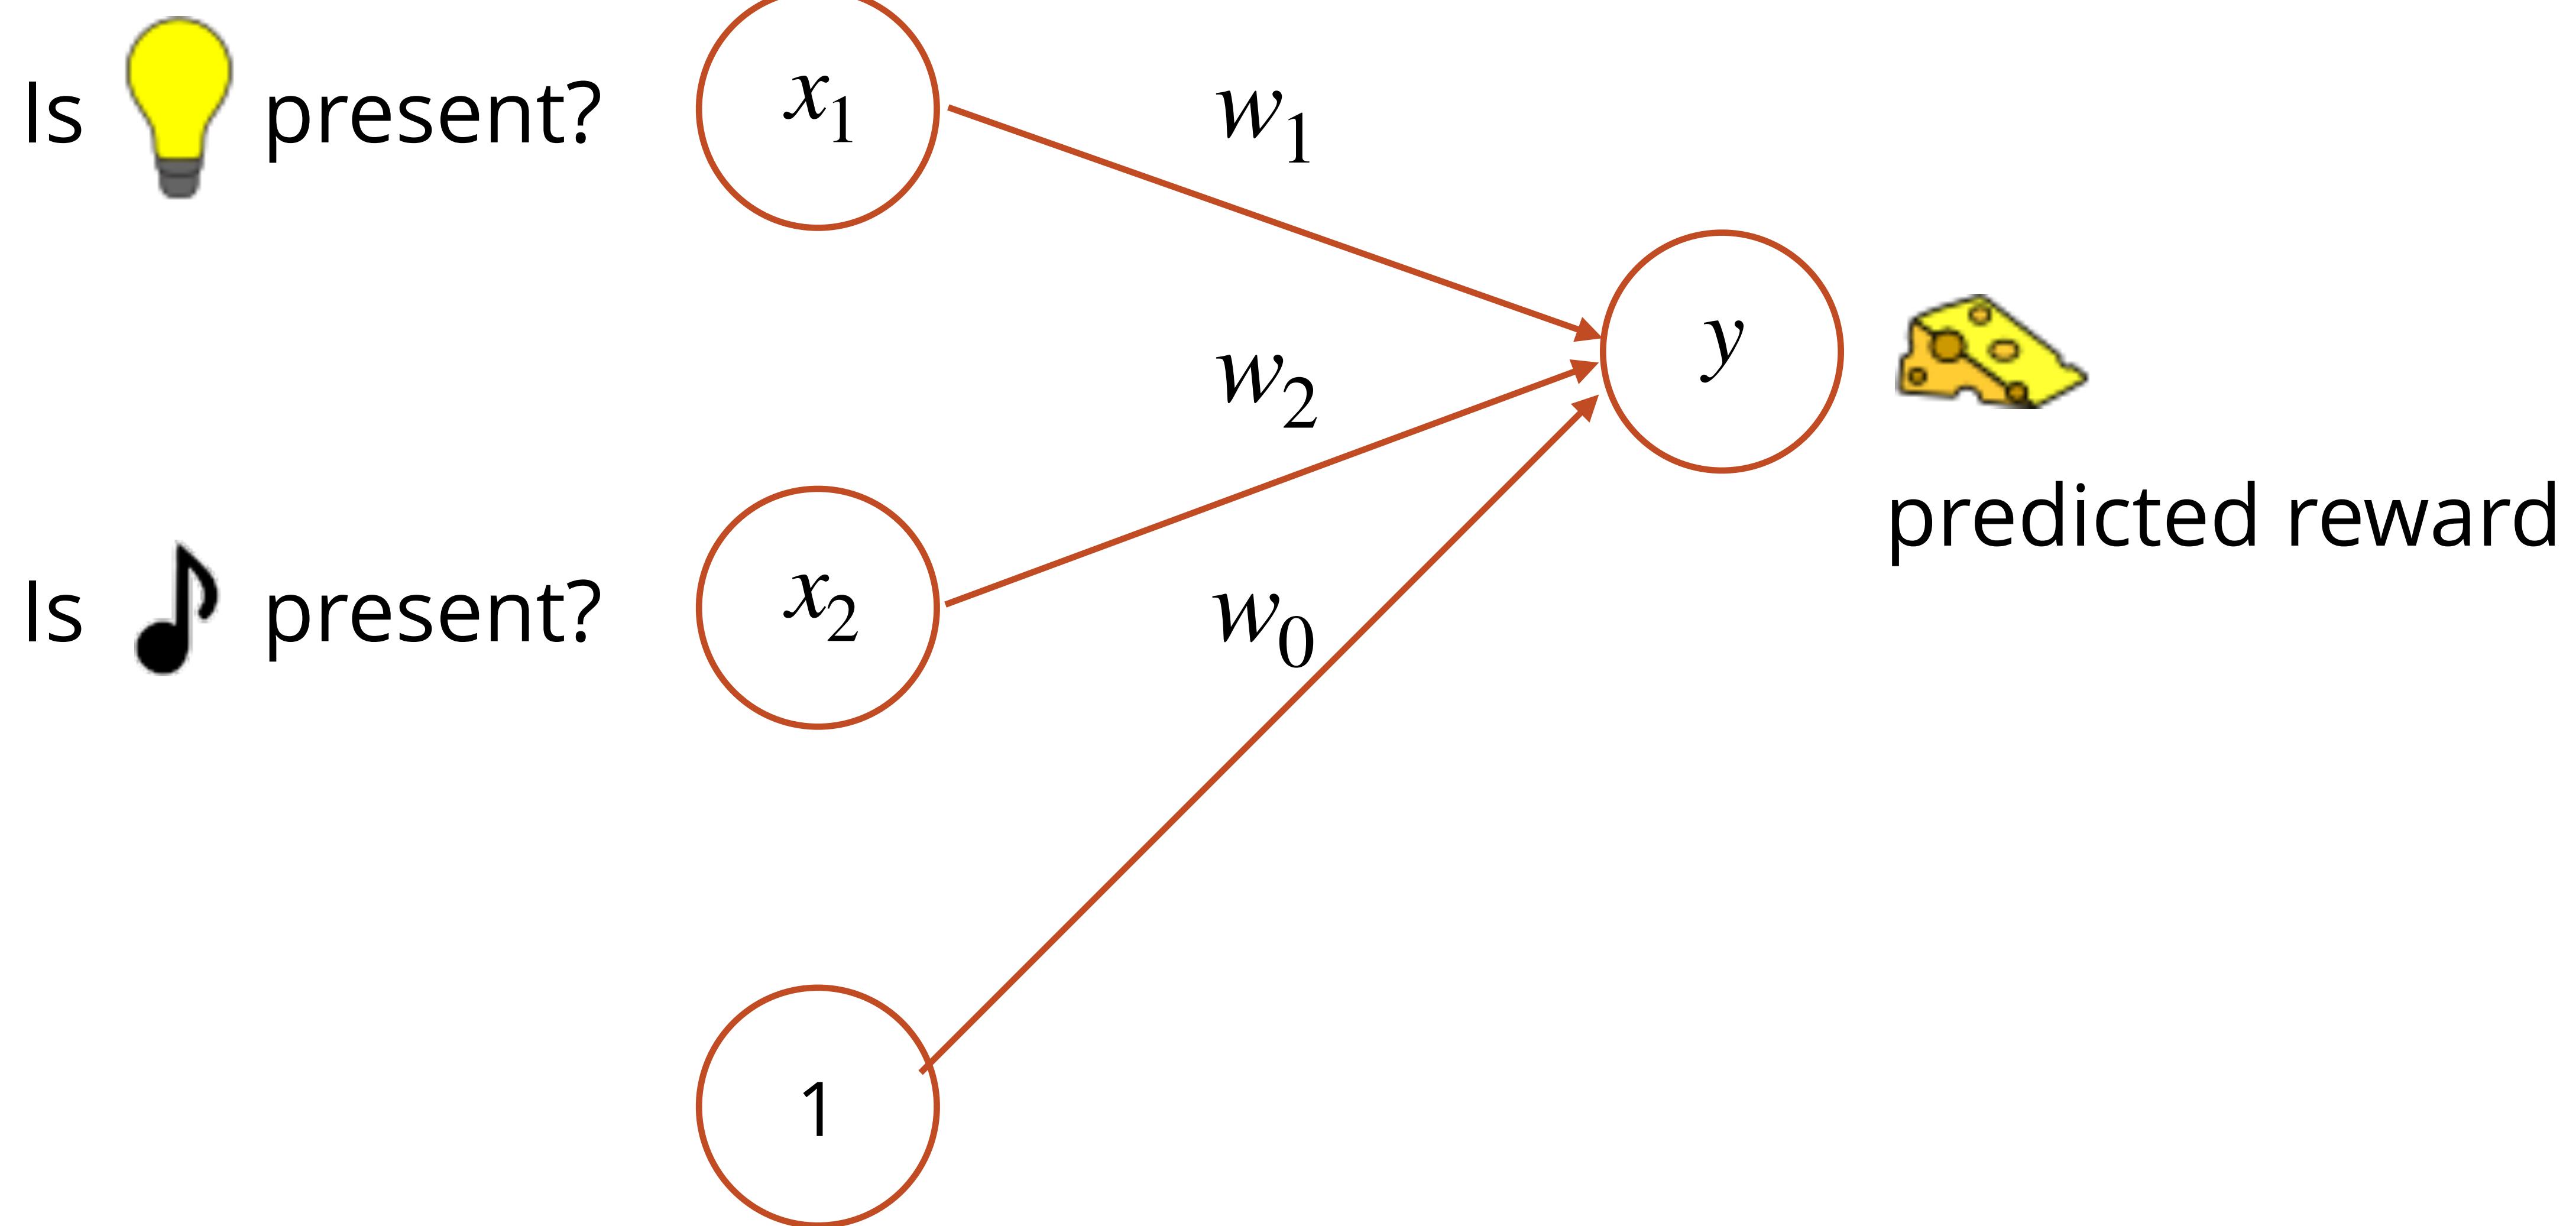
Modification 2: a bias term



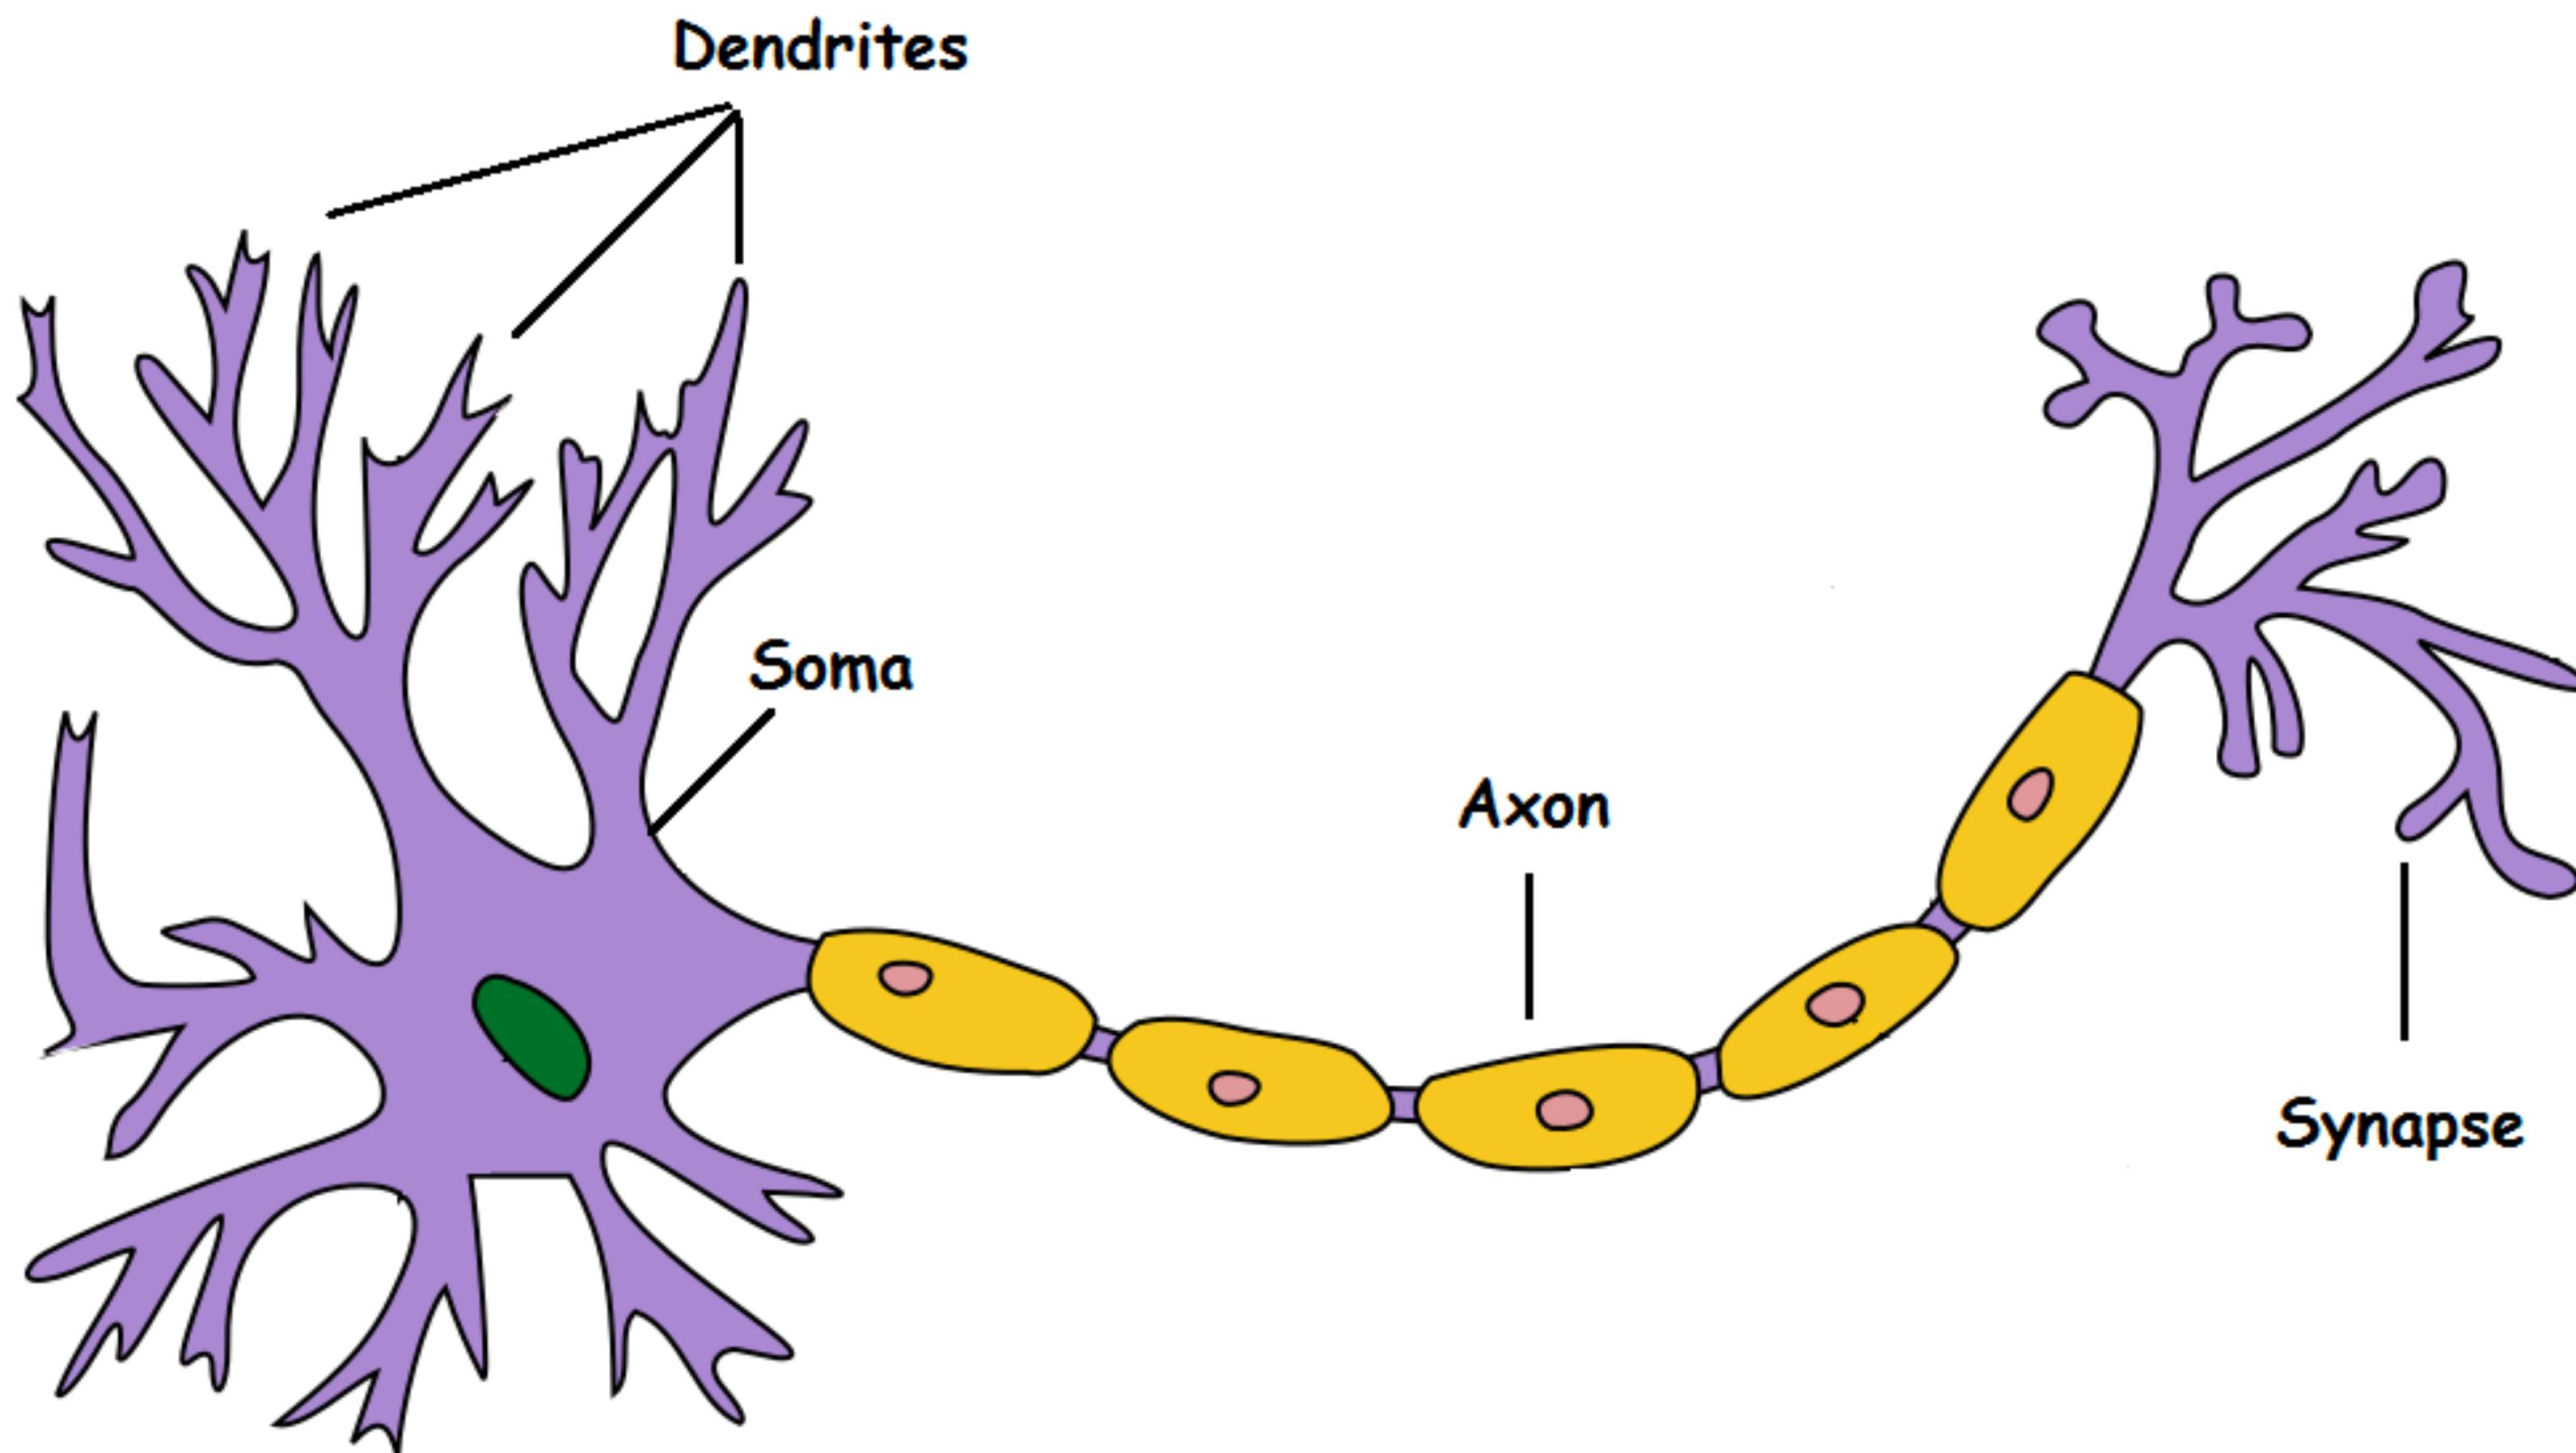
Aside: This is exactly logistic regression!



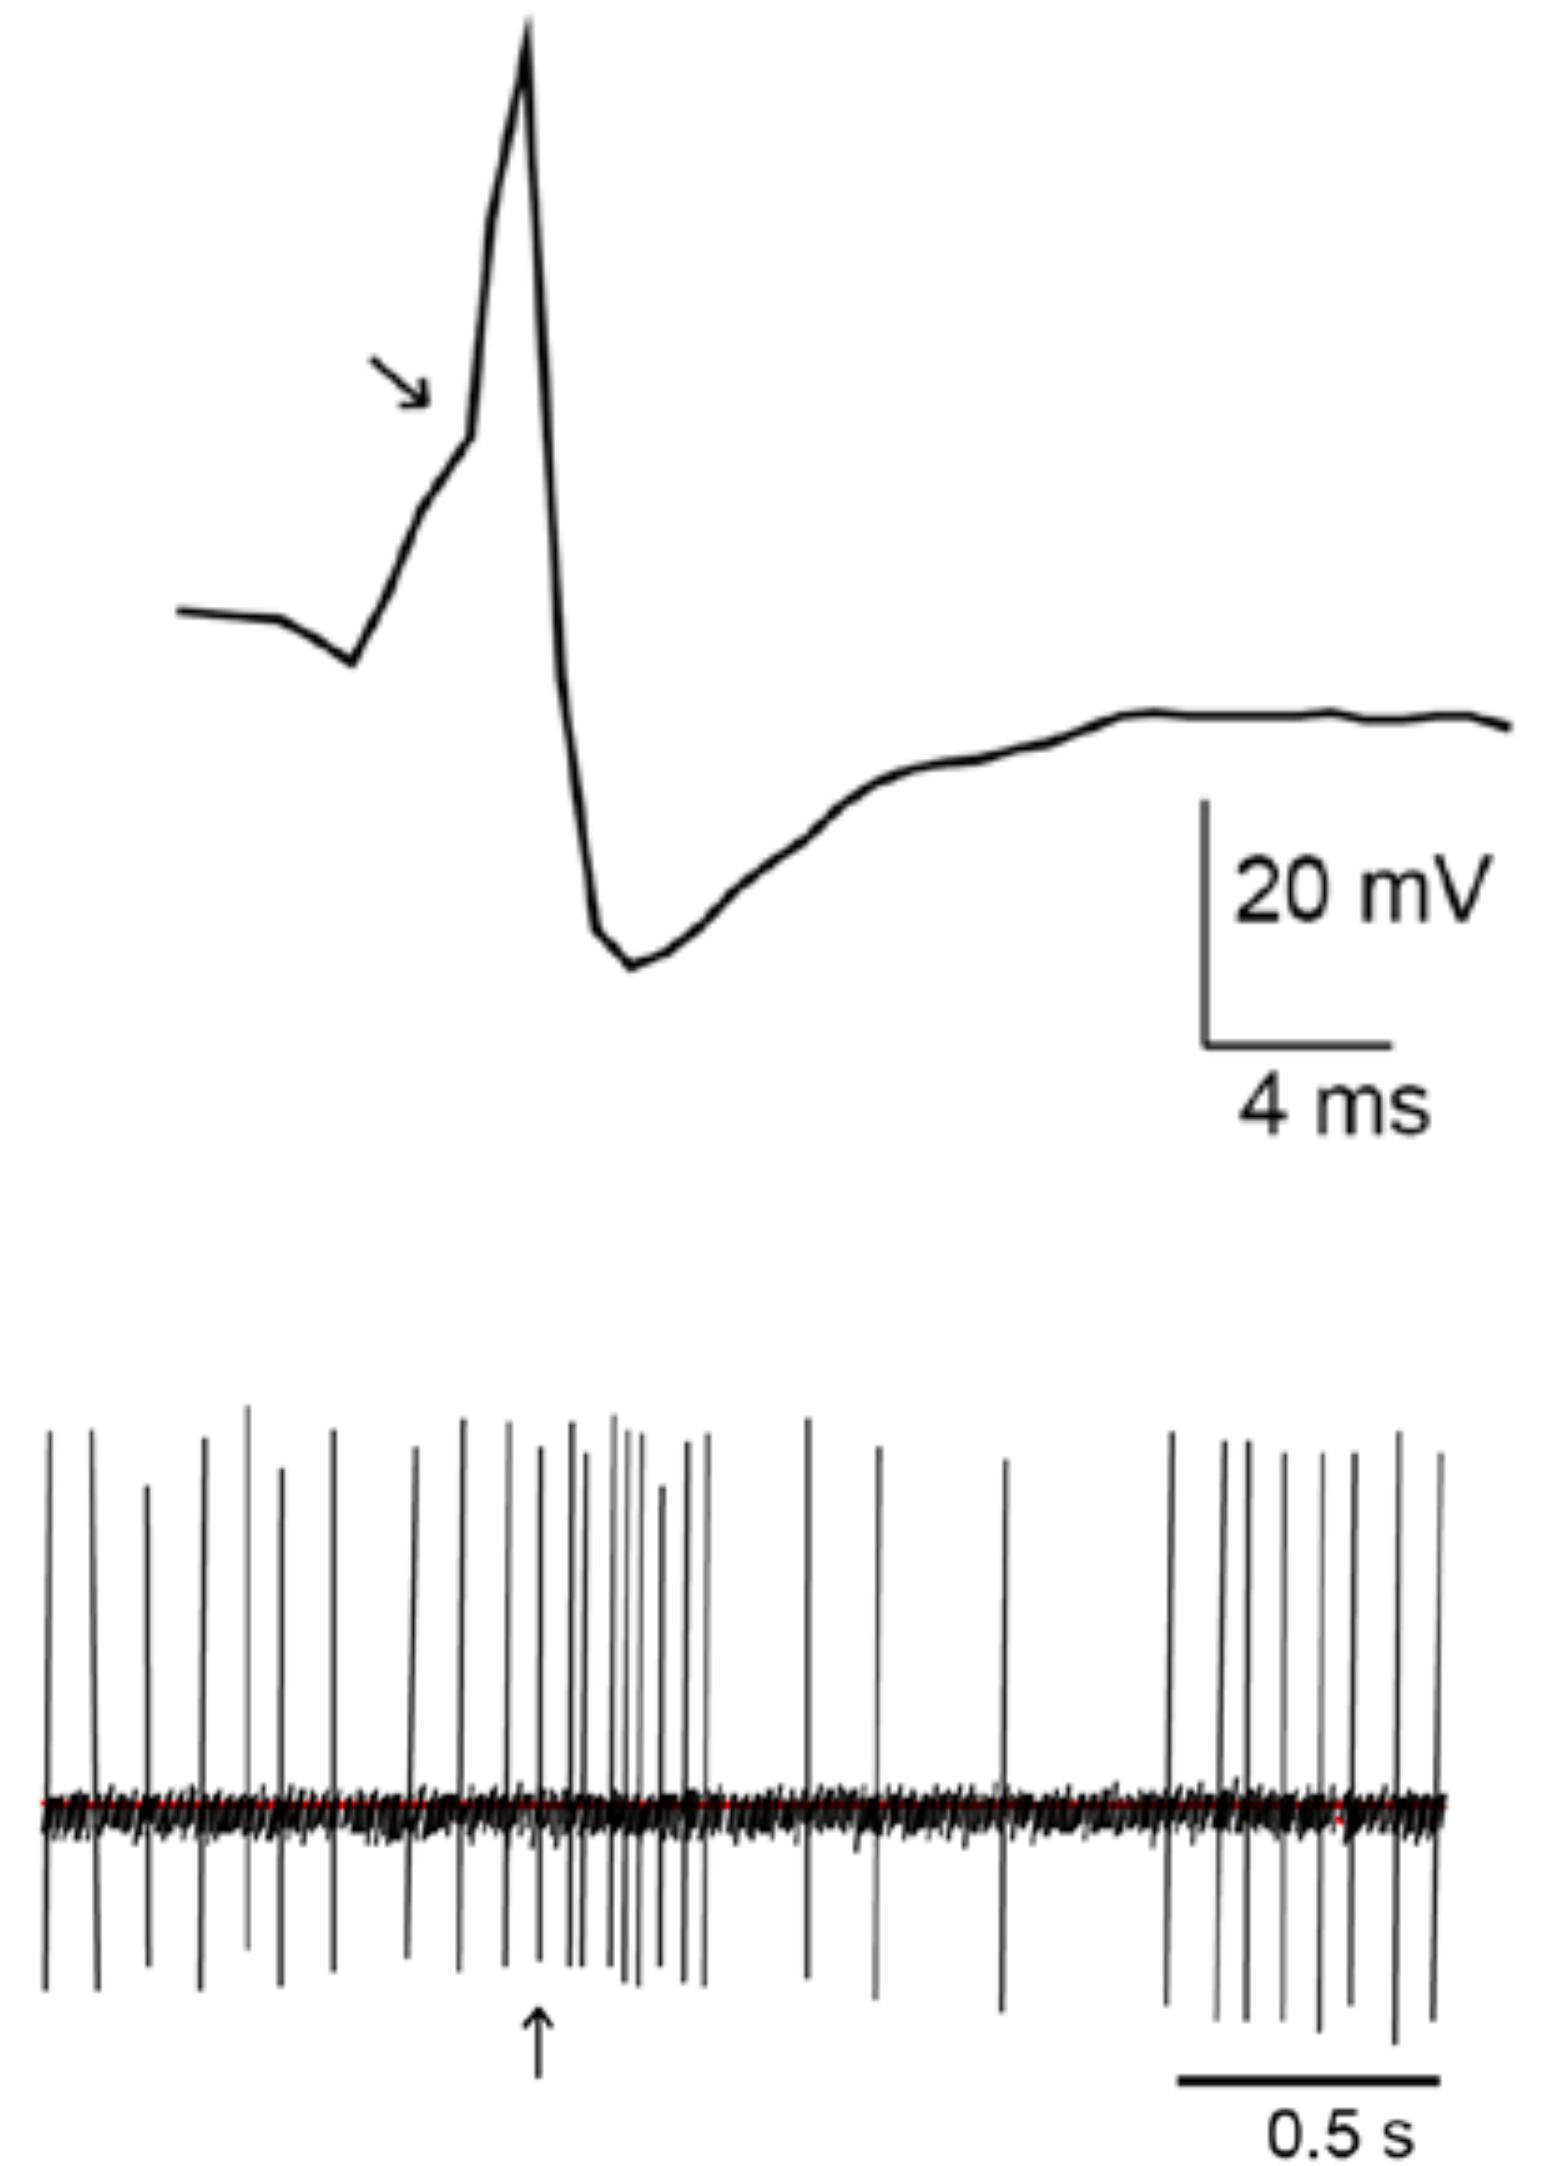
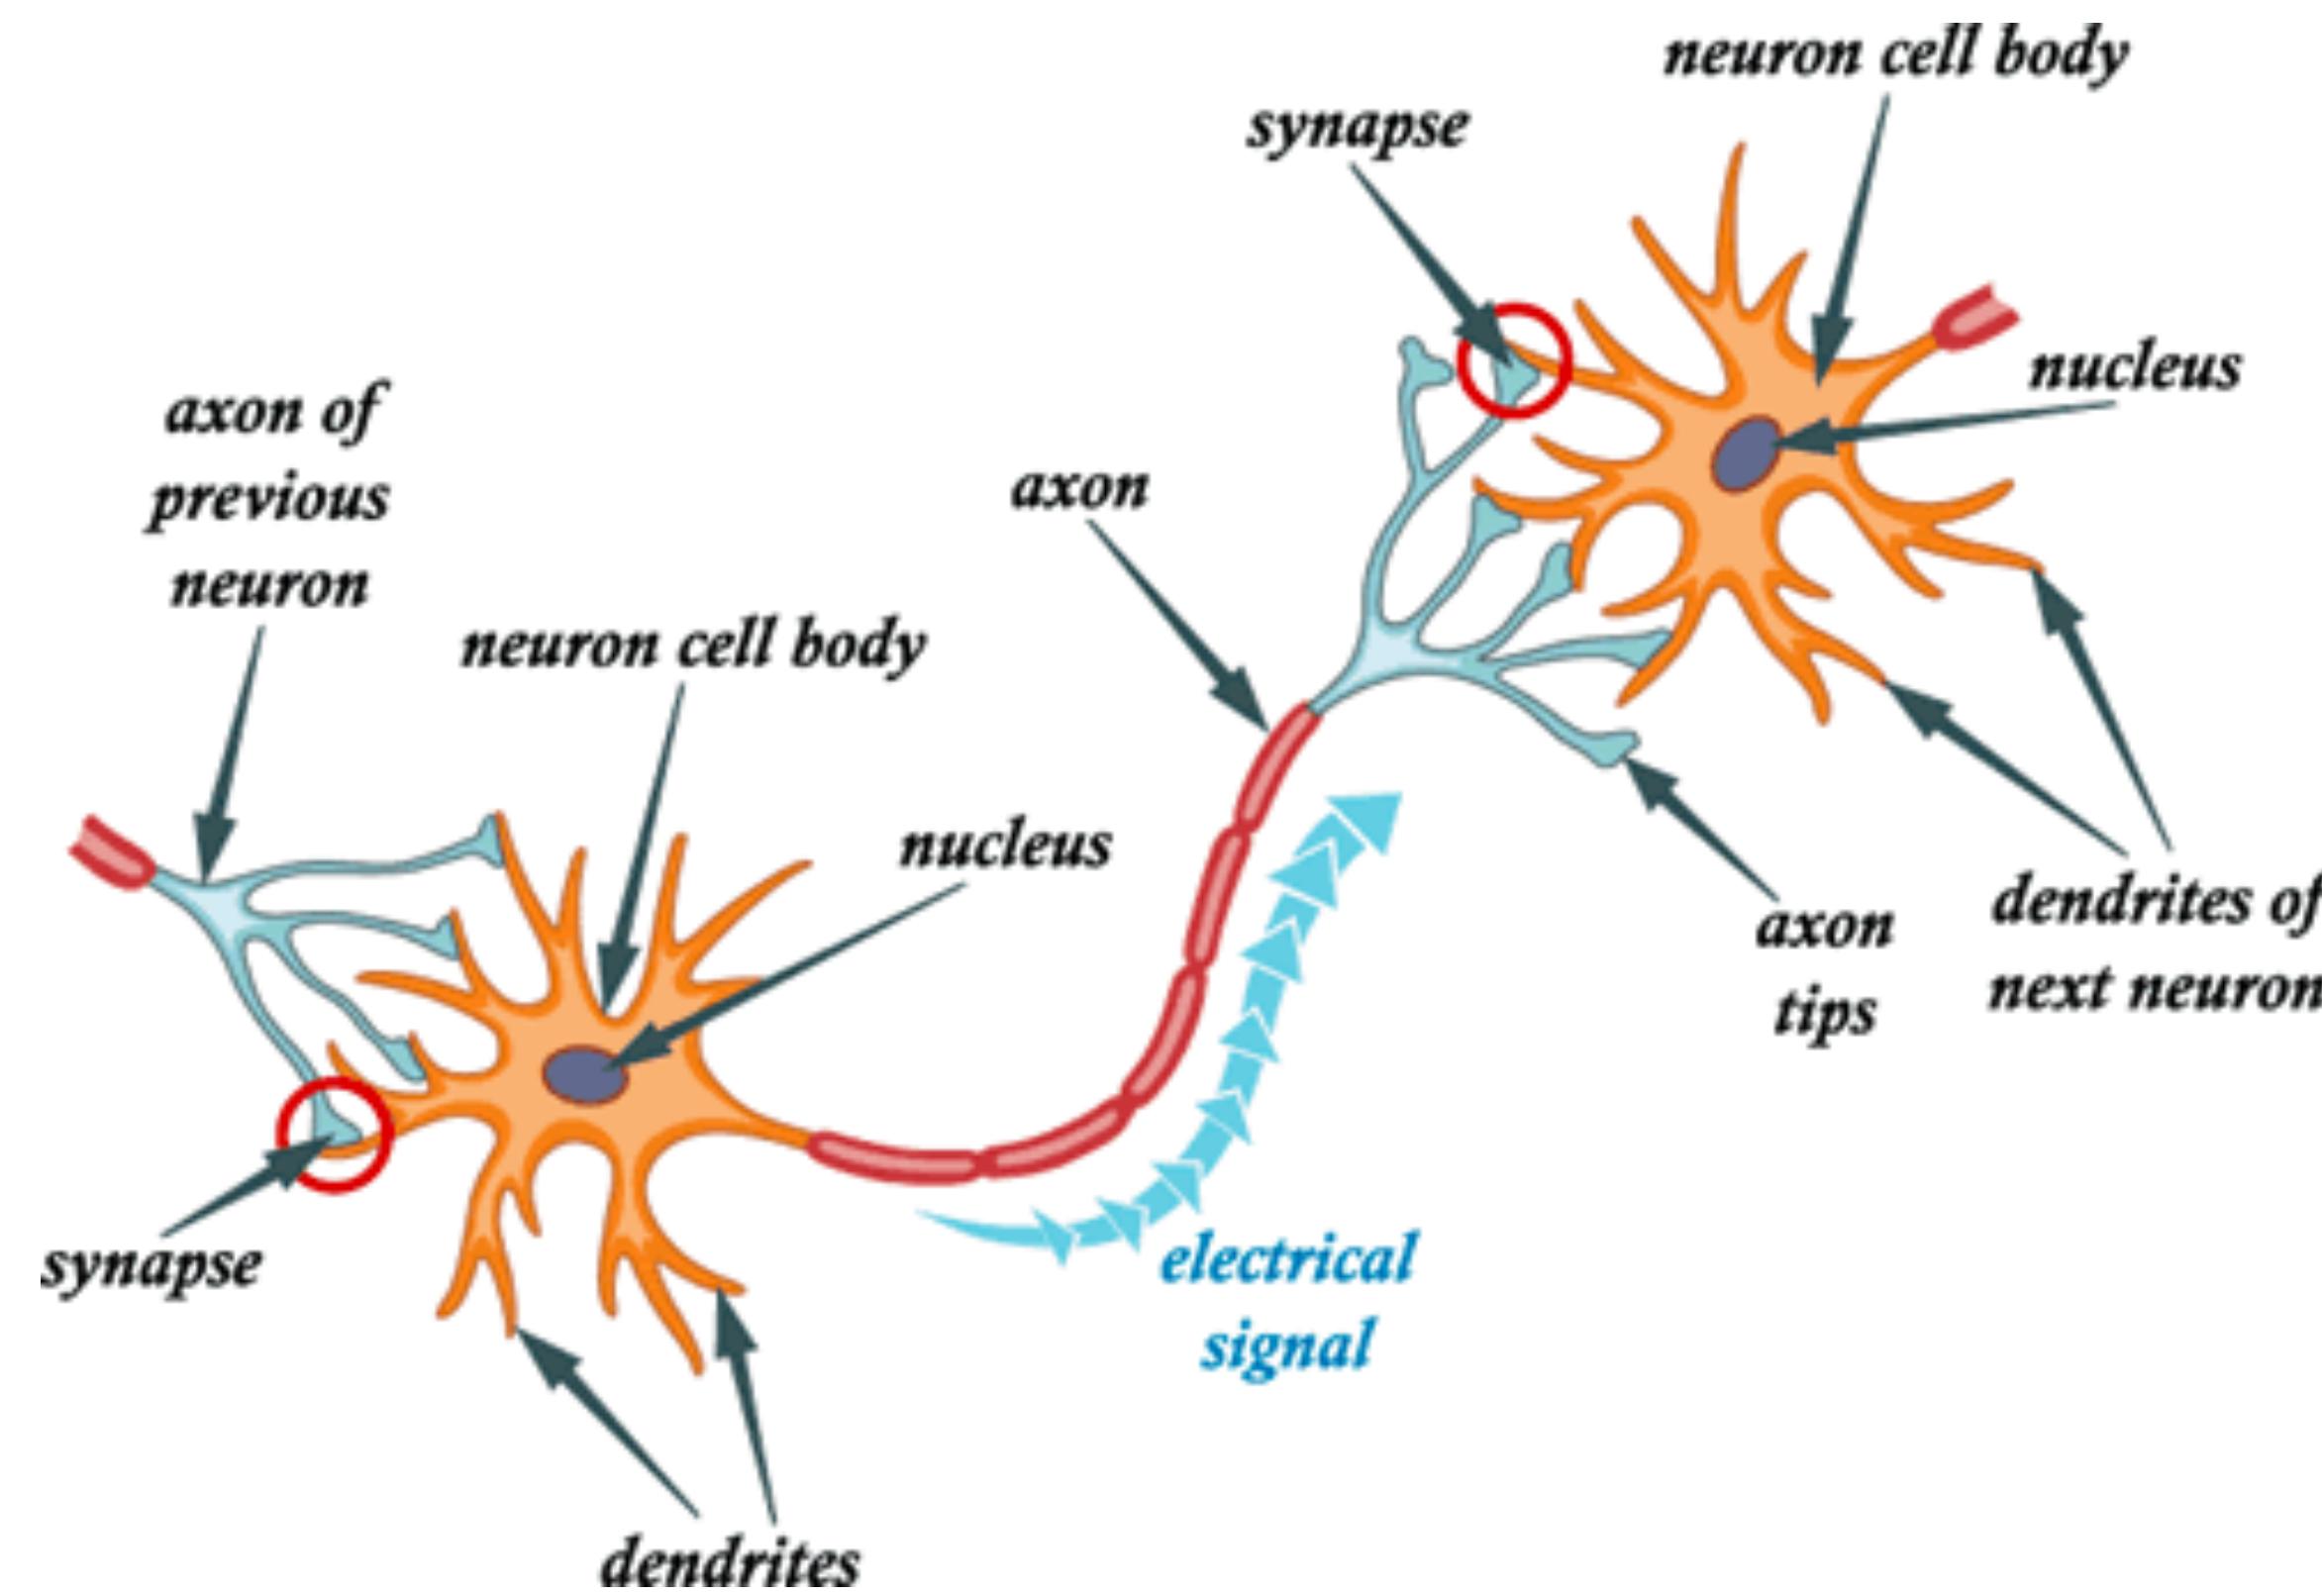
It's also a very simple model of a neuron



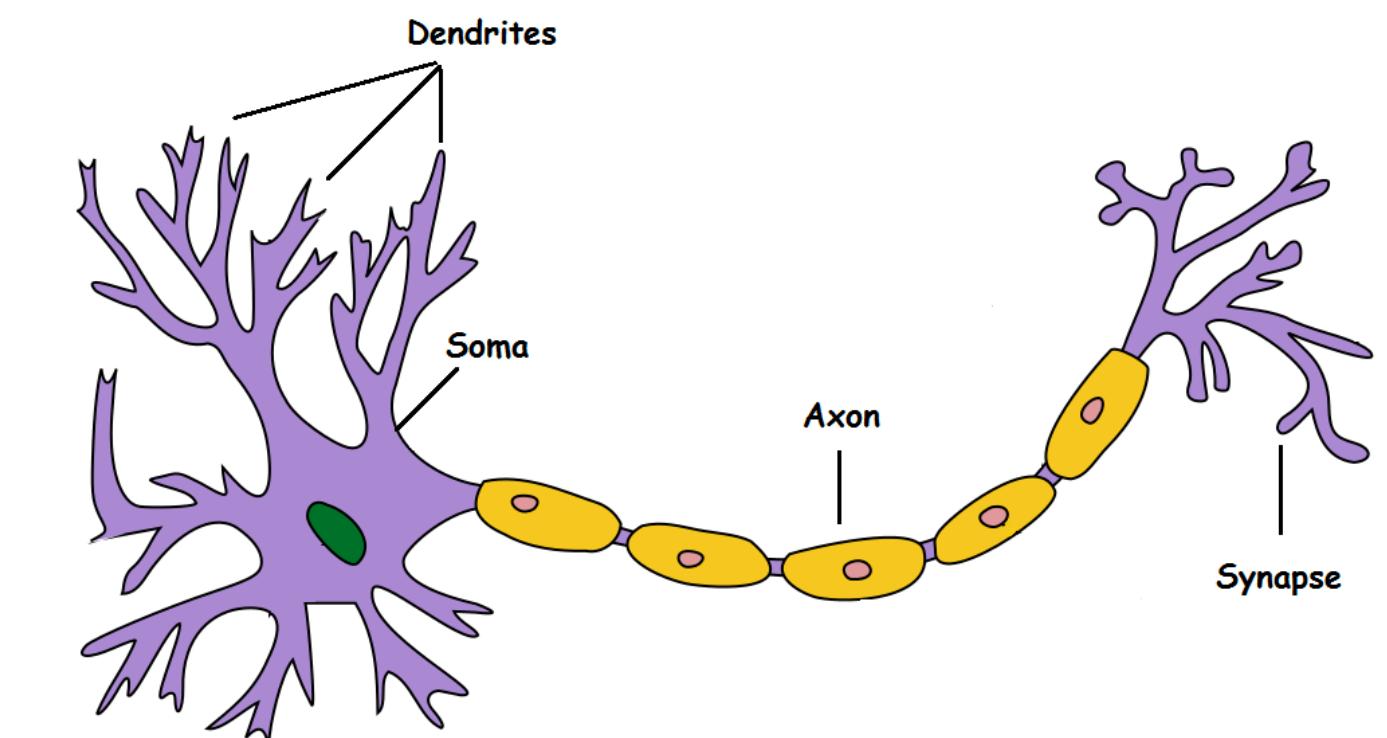
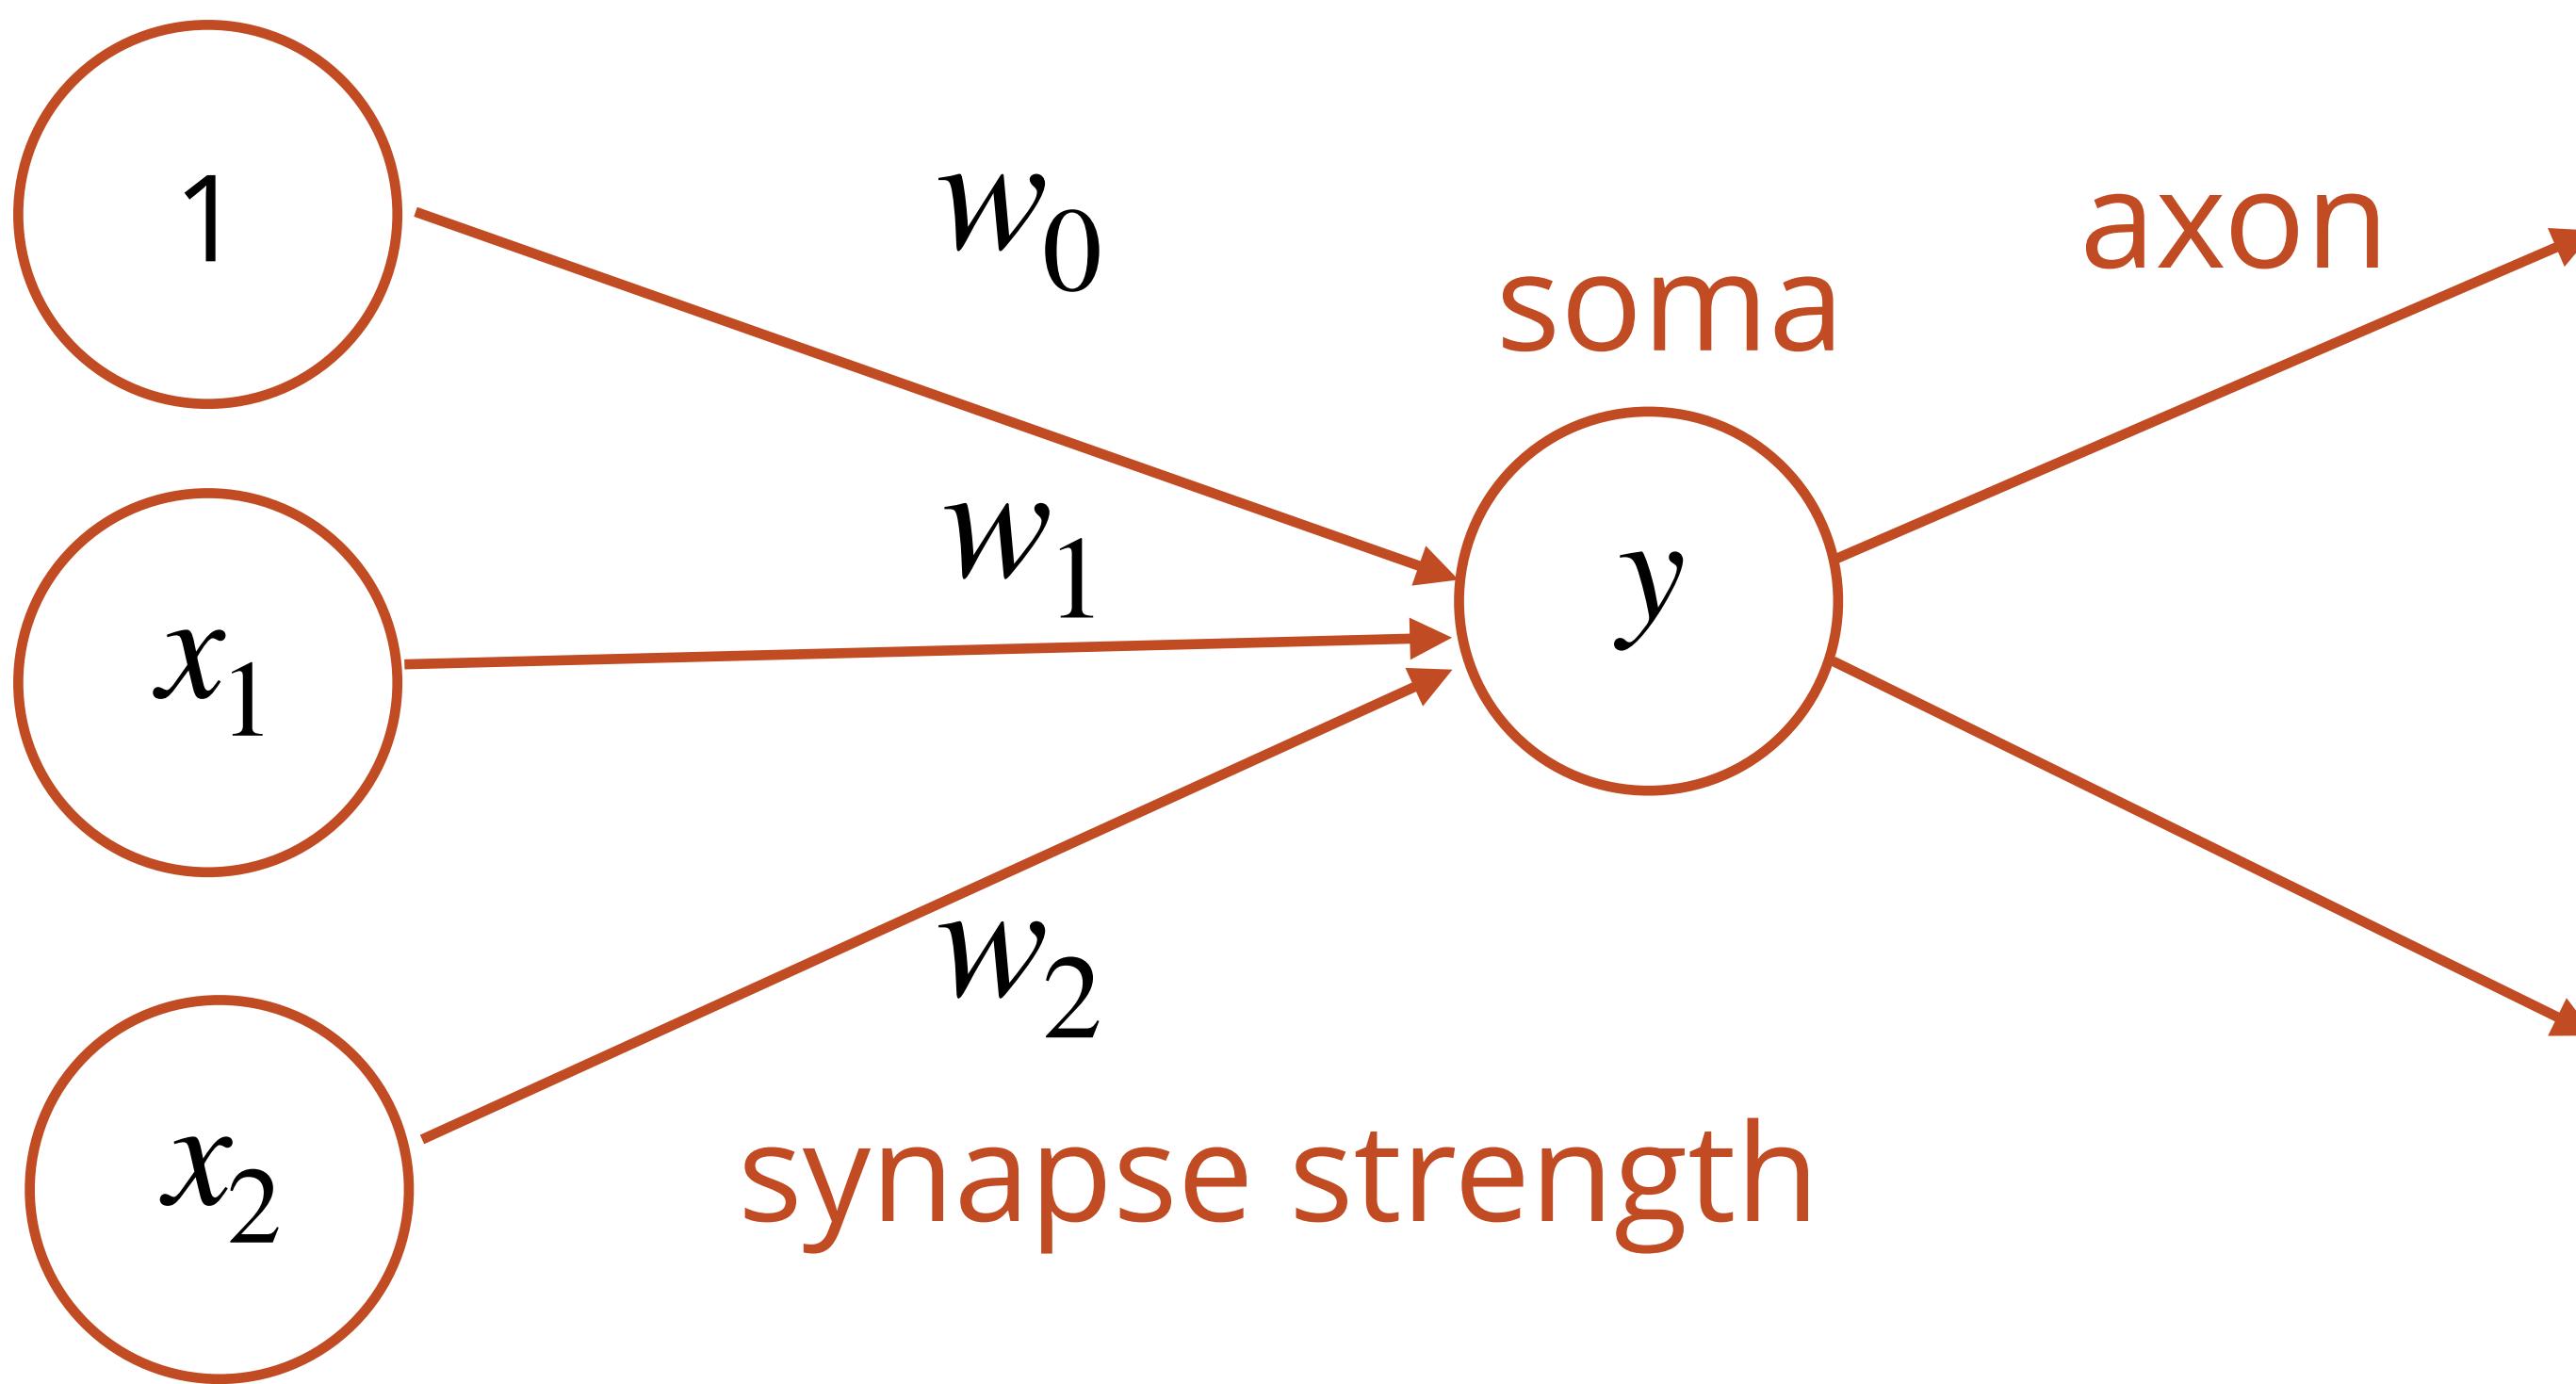
A biological neuron



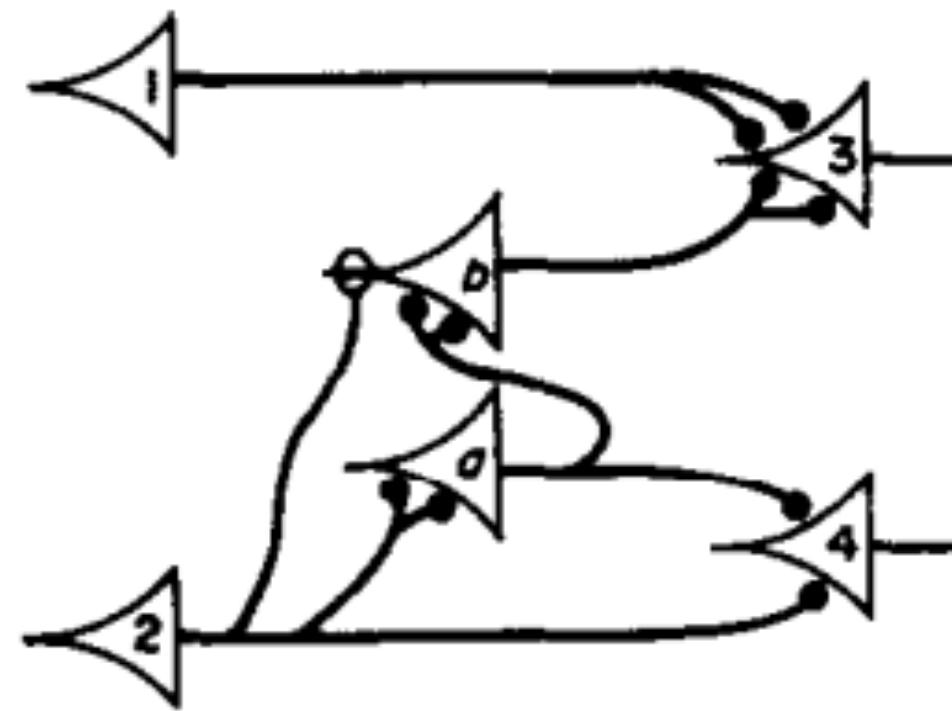
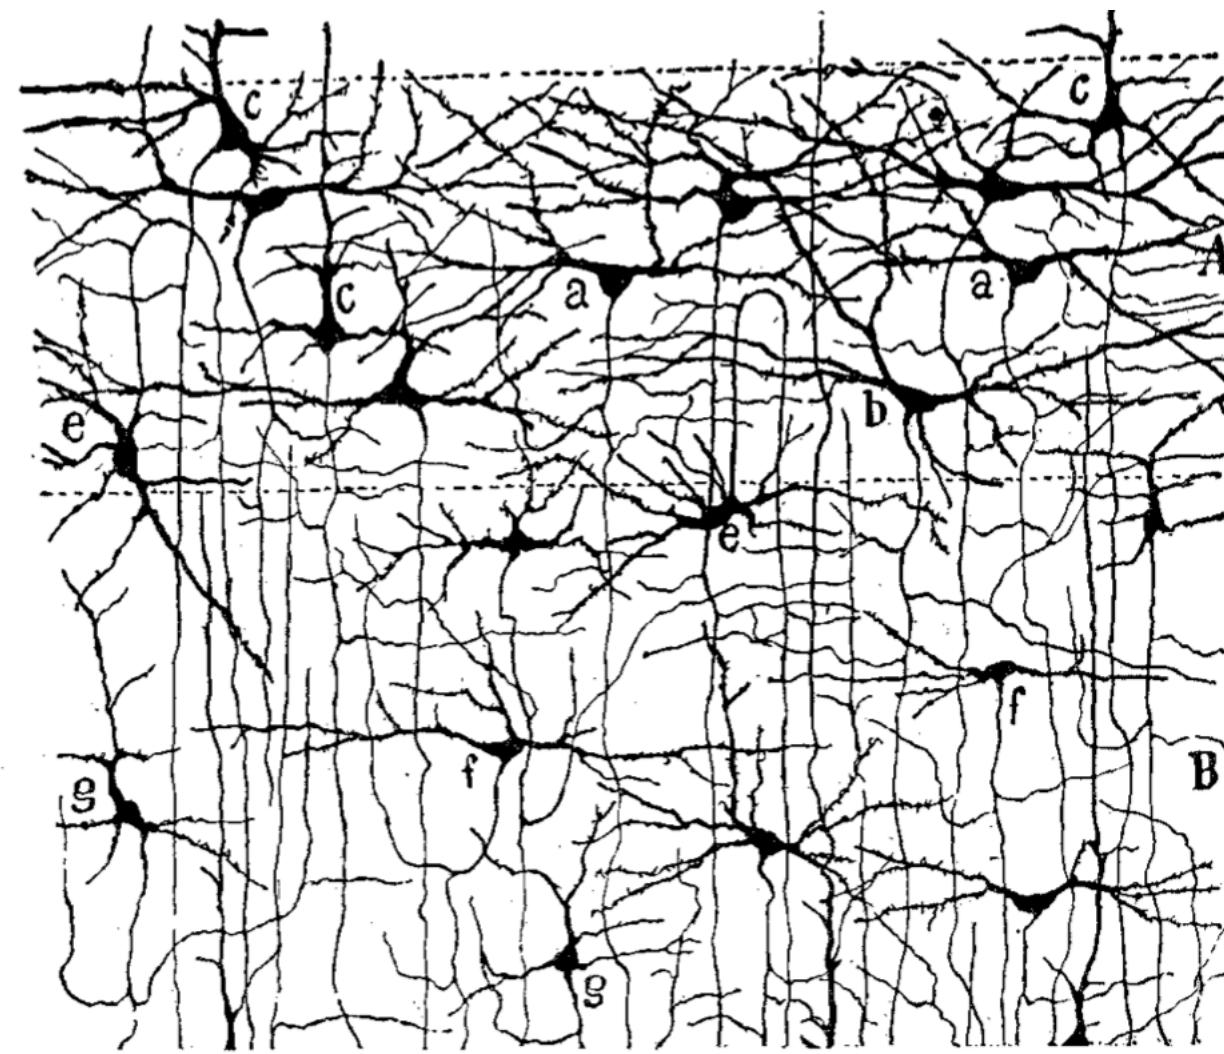
A biological neuron



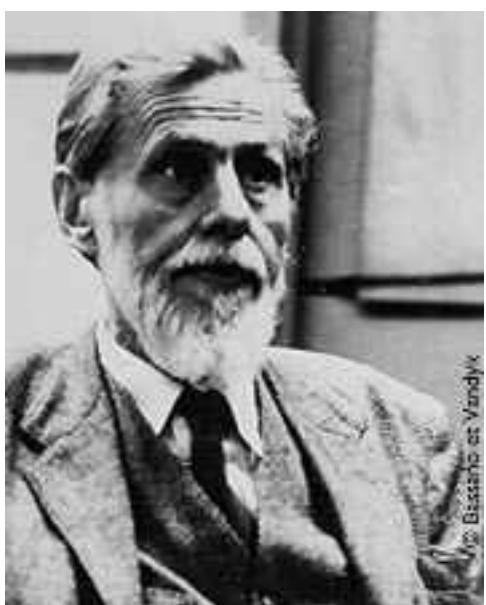
An artificial neuron



The first artificial neural networks



Because neuronal firing is discrete, neural networks can approximate boolean logic!



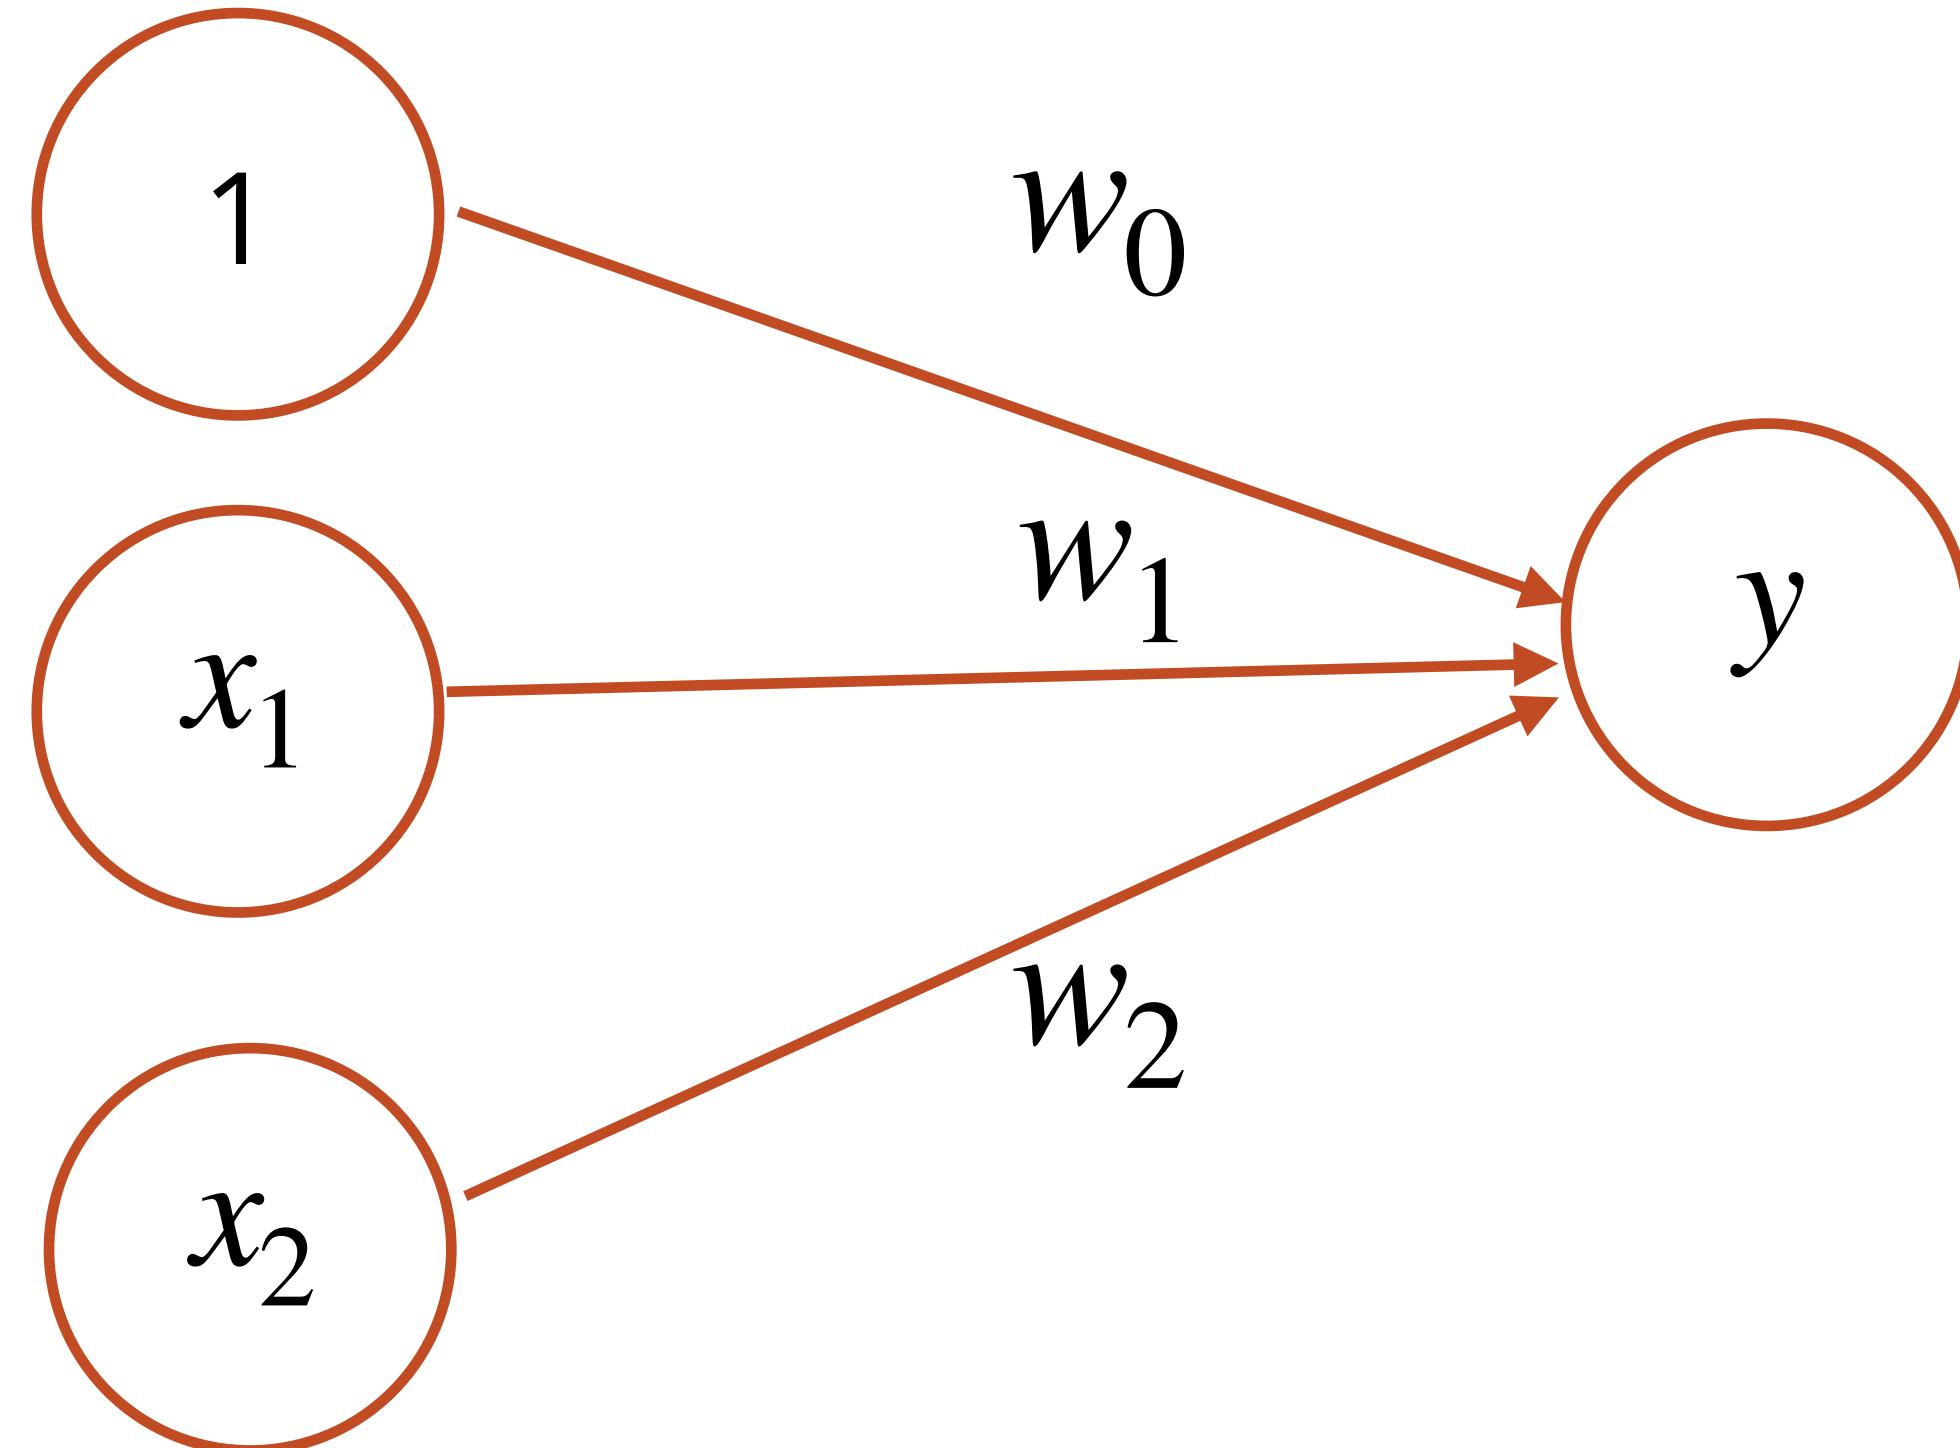
Warren McCulloch



Walter Pitts

What's more, they can *learn* to do logic

The Perceptron learning rule (Rosenblatt, 1958)



Perceptron learning rule

Rescorla-Wagner learning rule

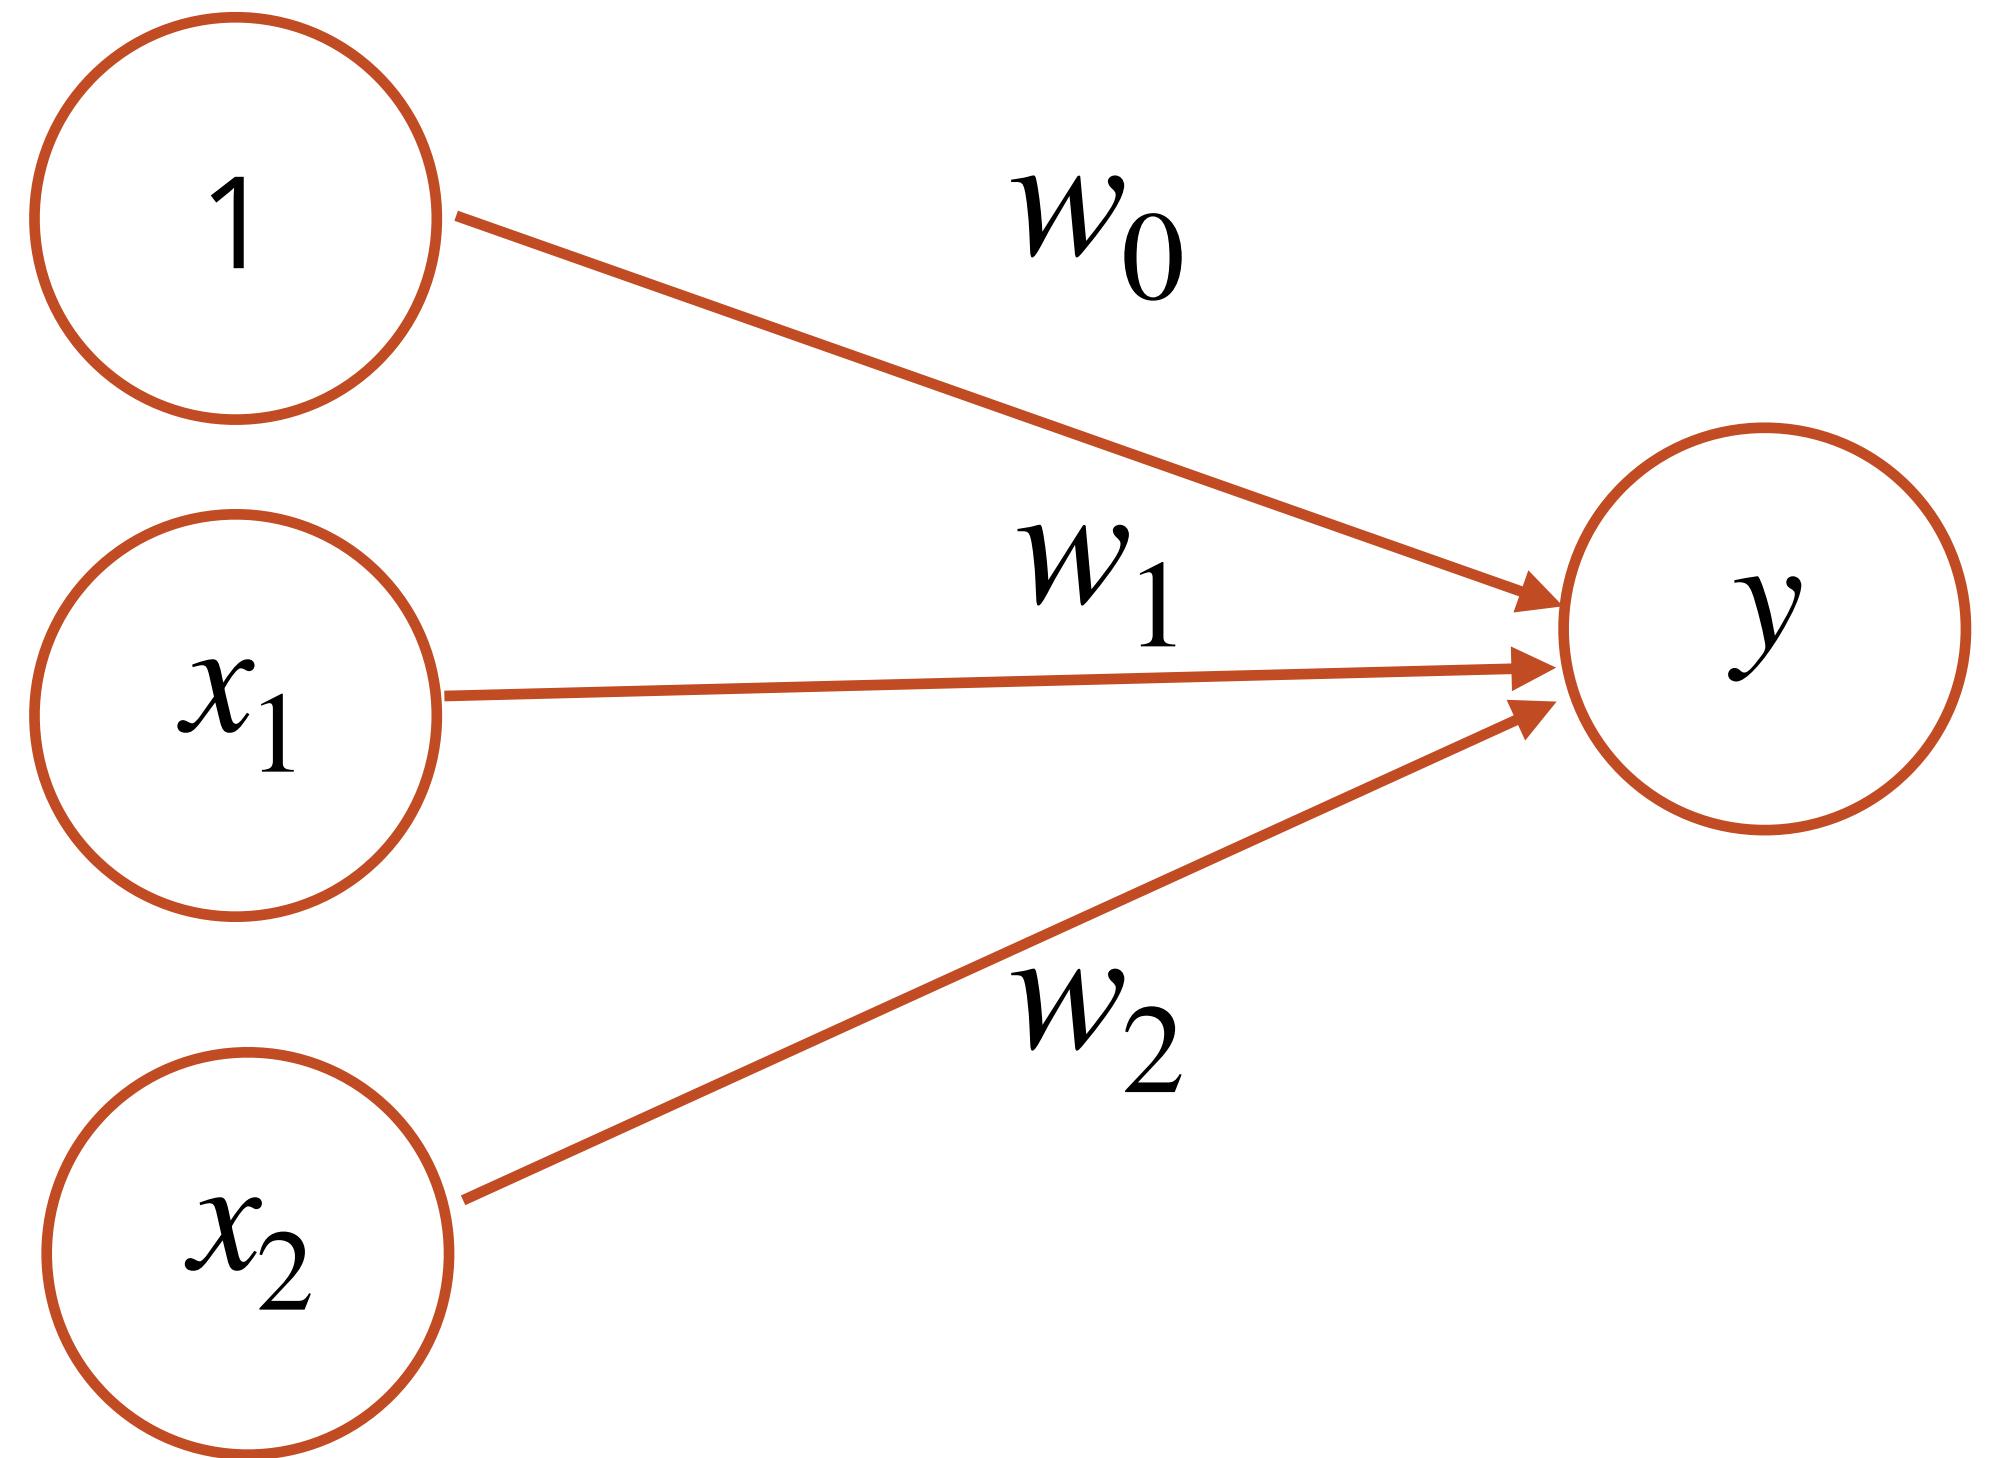
$$f(x) = \frac{1}{1 - e^x}$$

$$y = f \sum_i w_i \cdot x_i$$

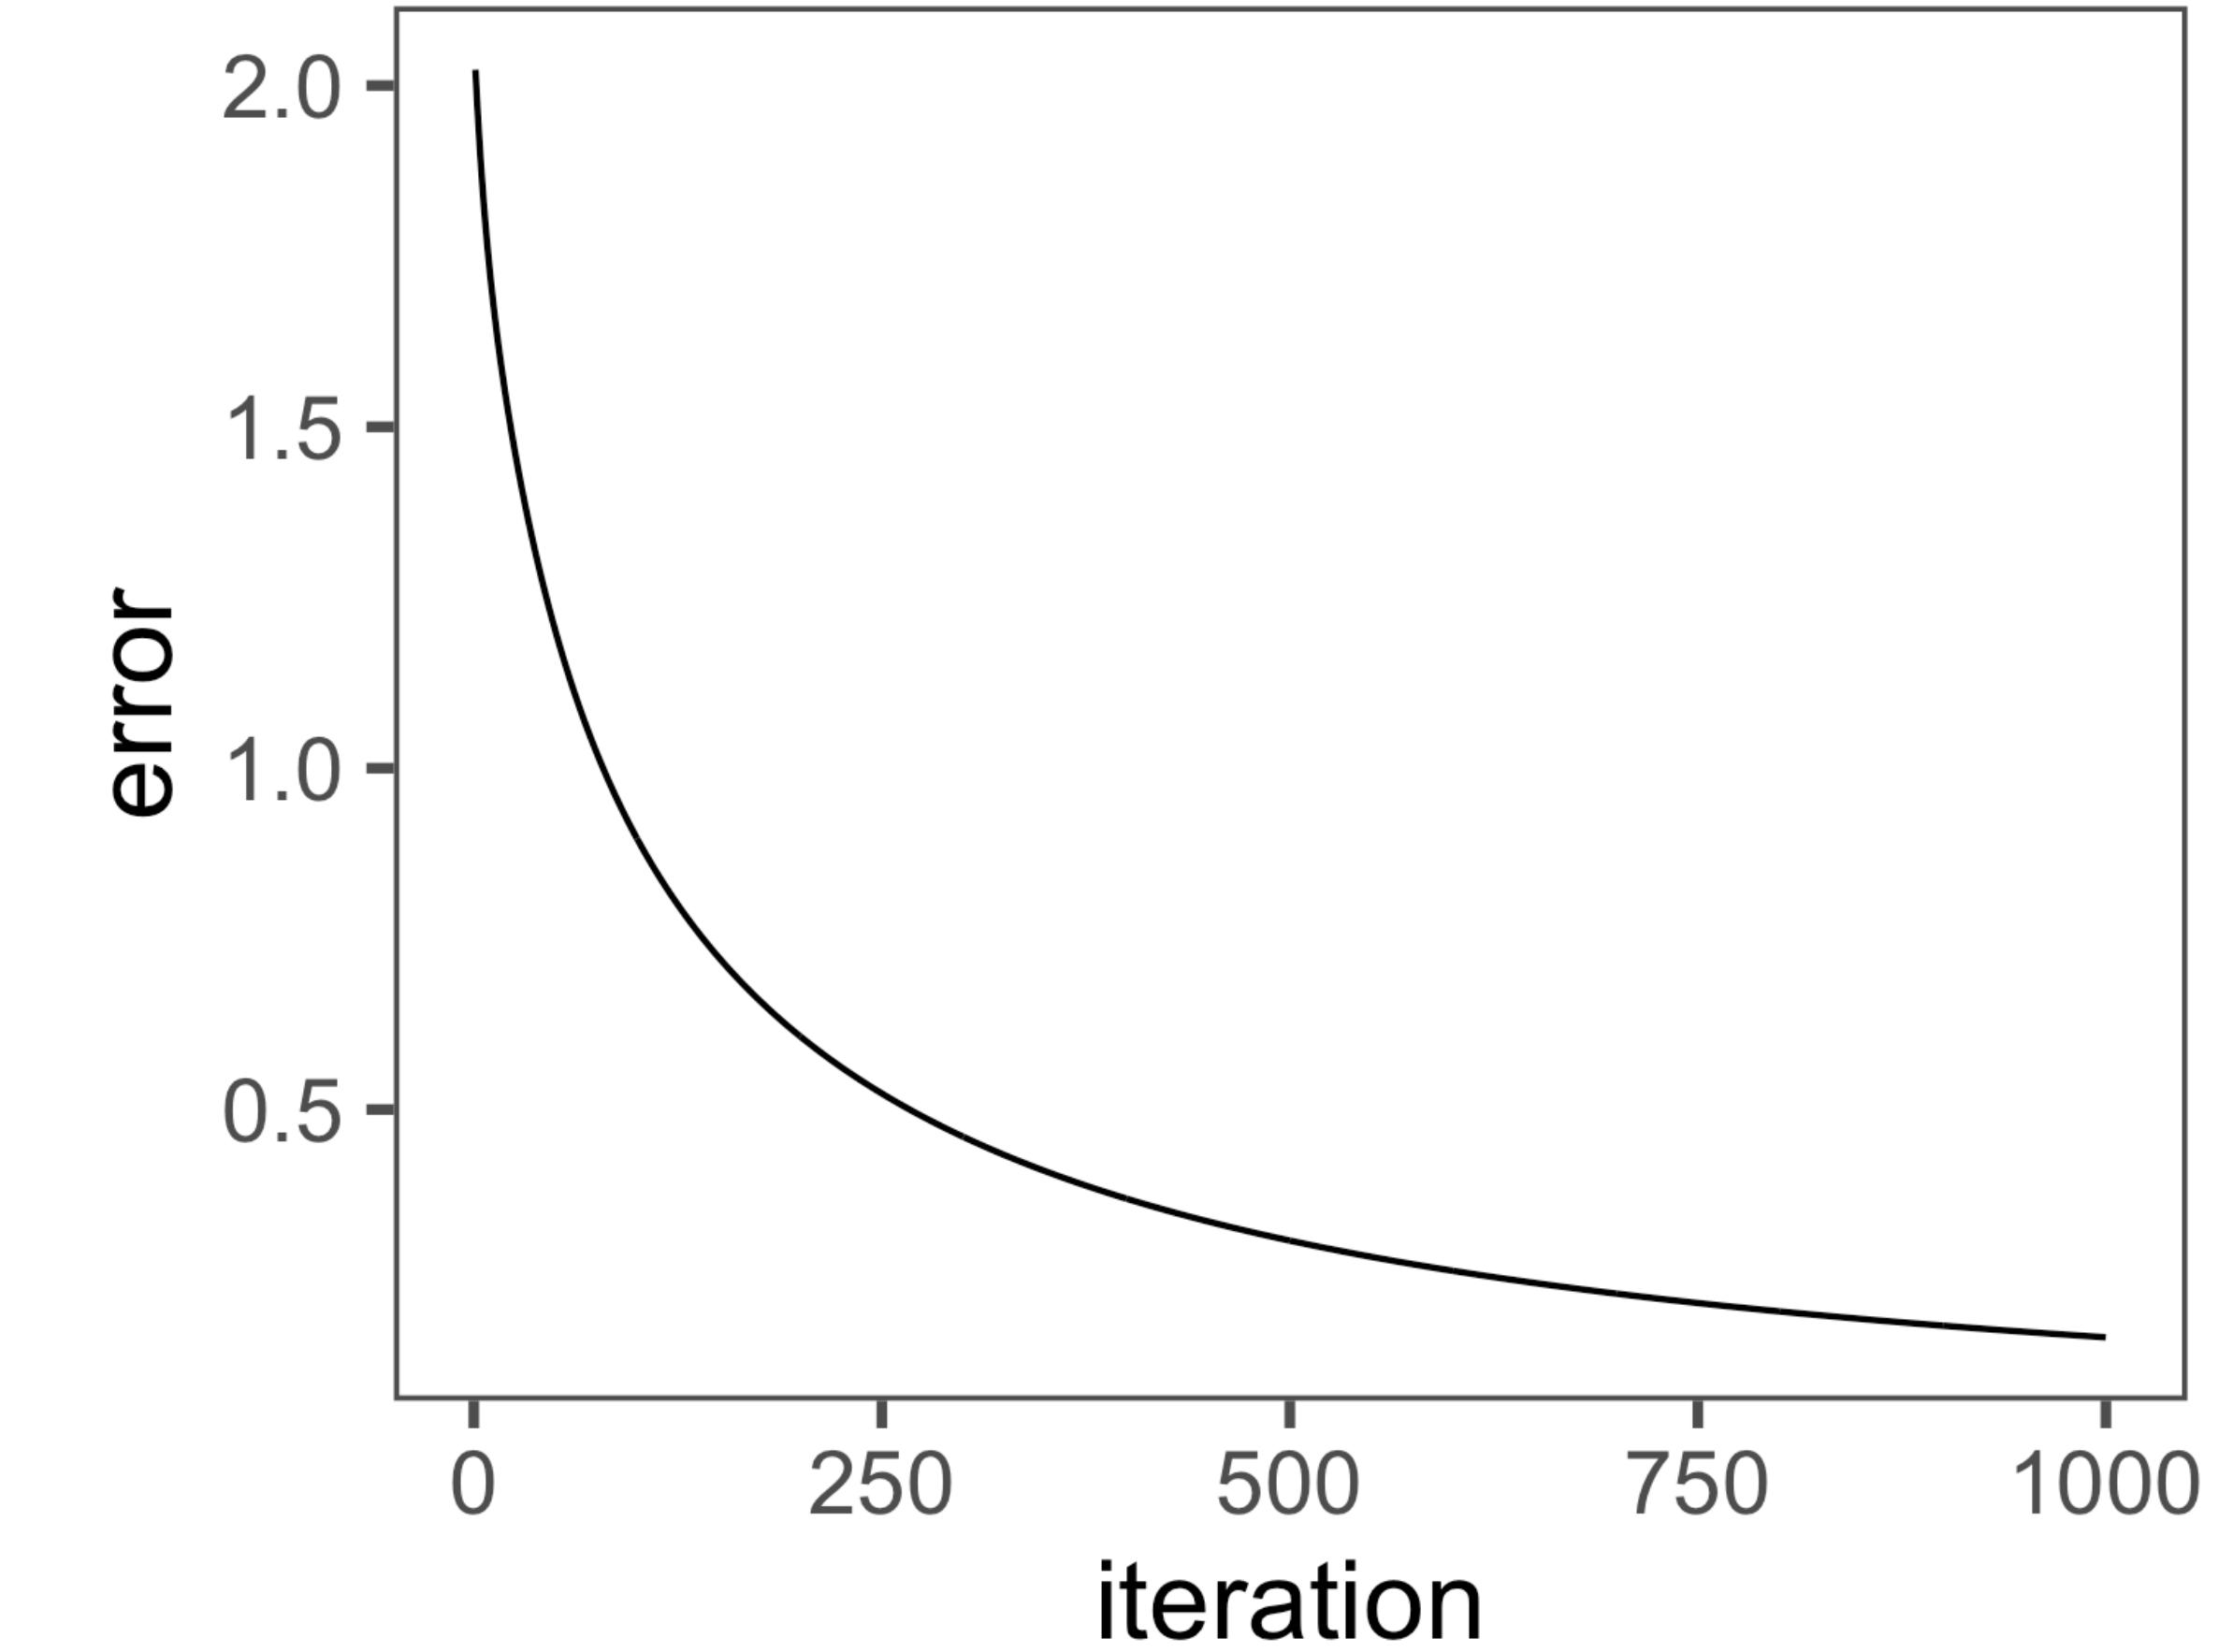
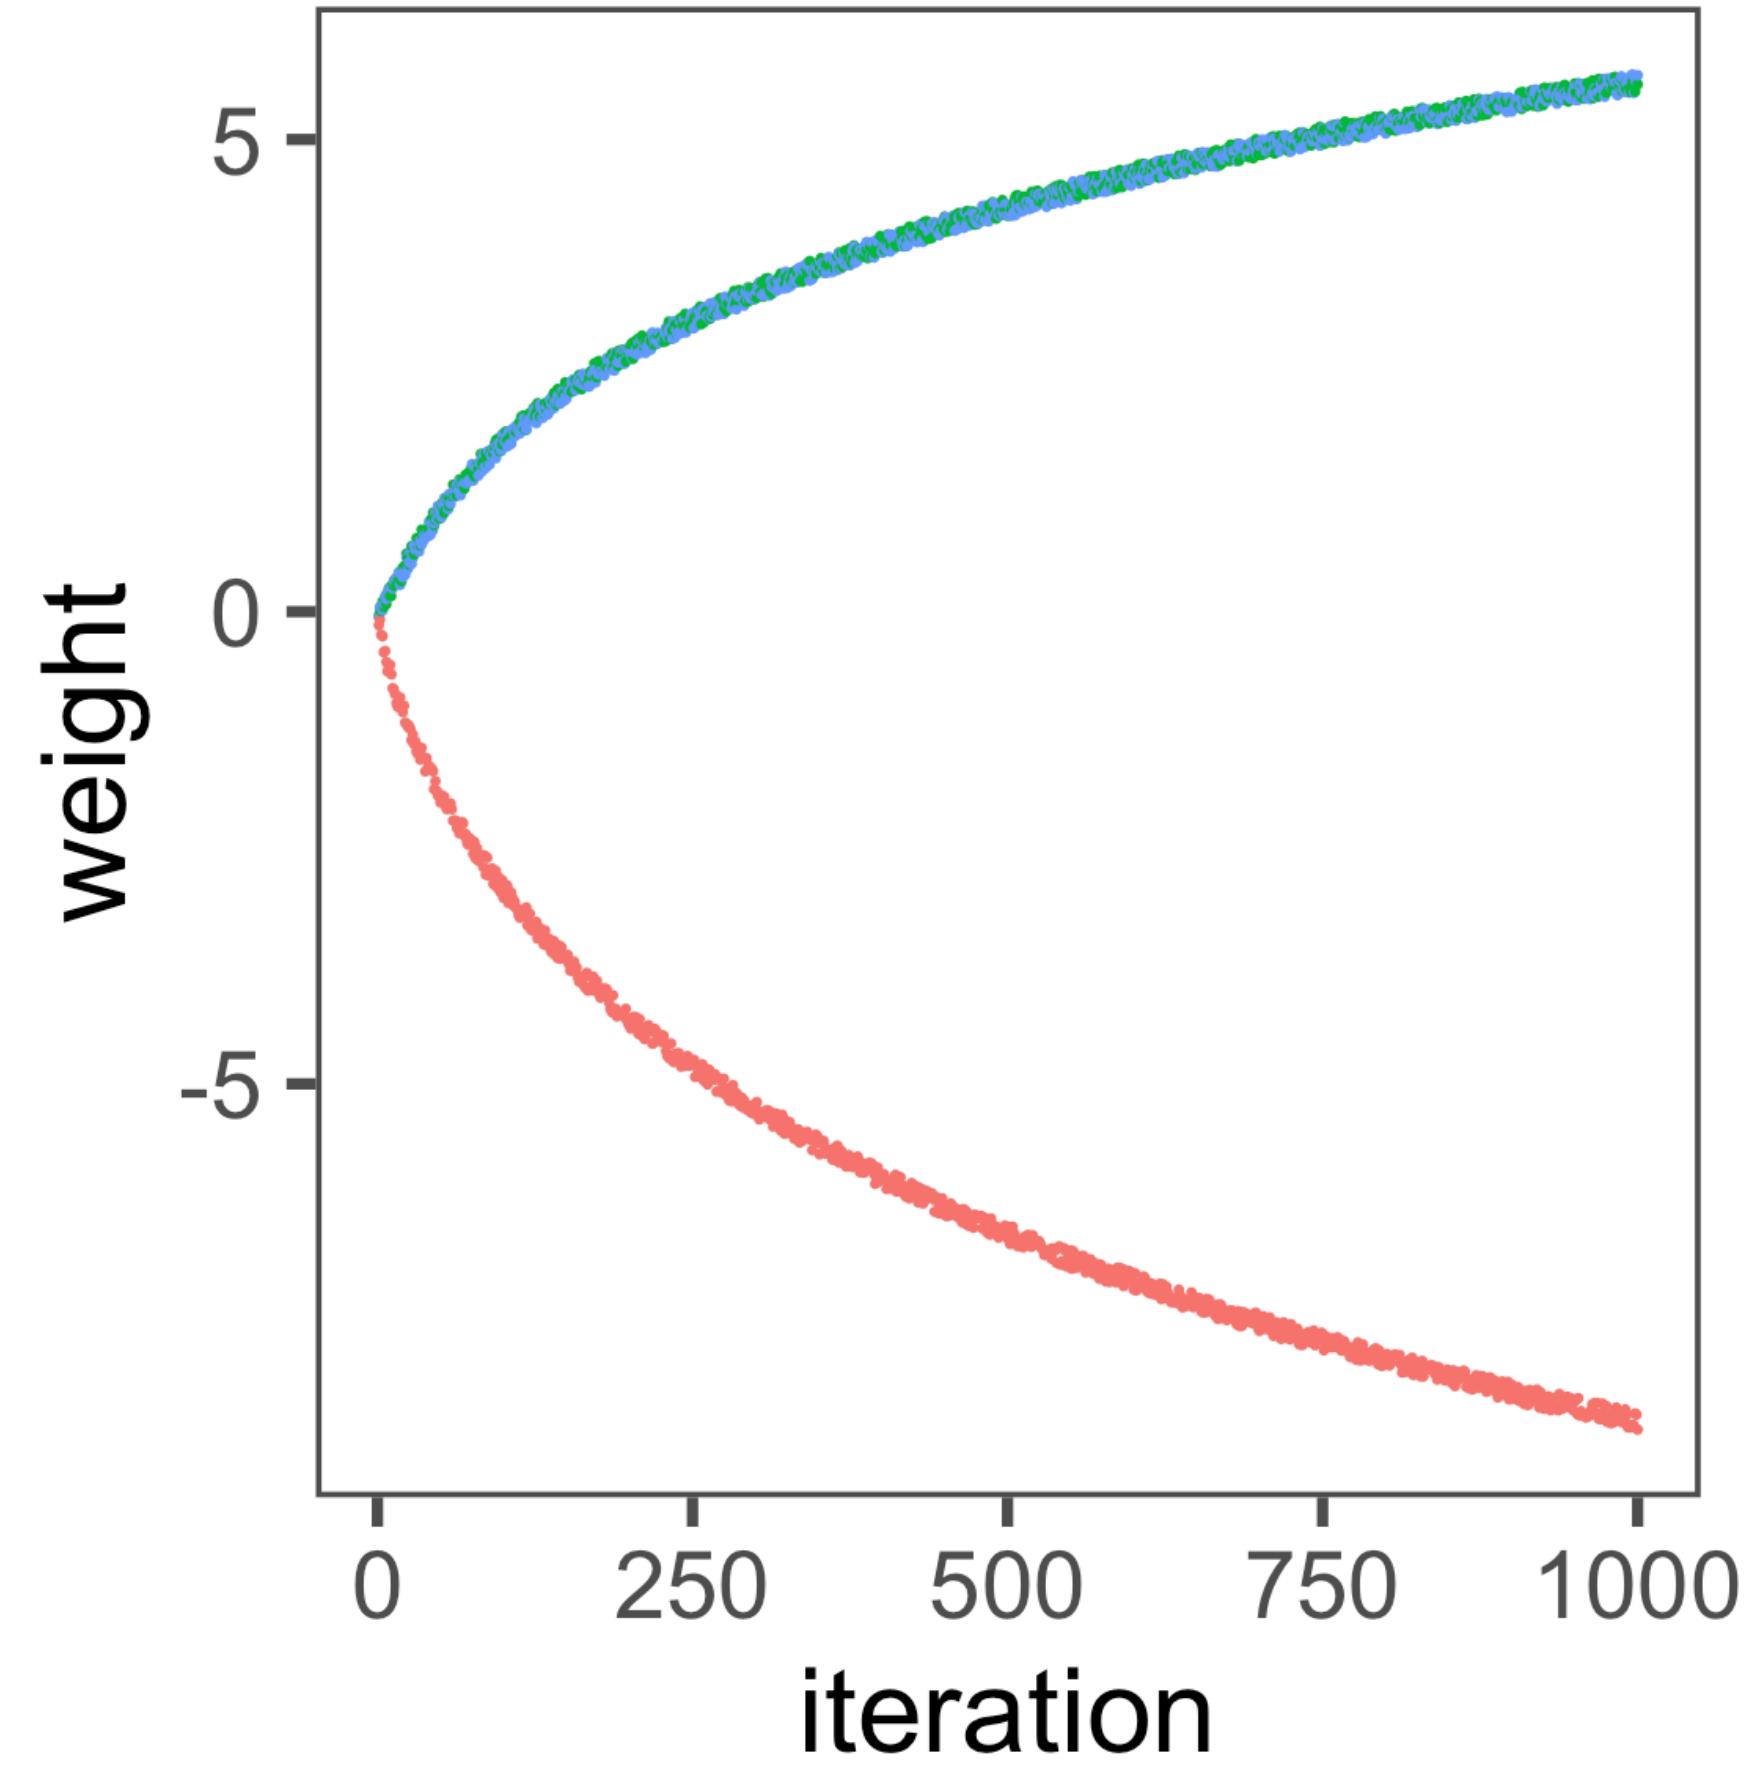
$$\Delta w_i = \alpha \cdot (y - \hat{y}) x_i$$

$$\Delta V = \alpha \cdot (\lambda - V_{total})$$

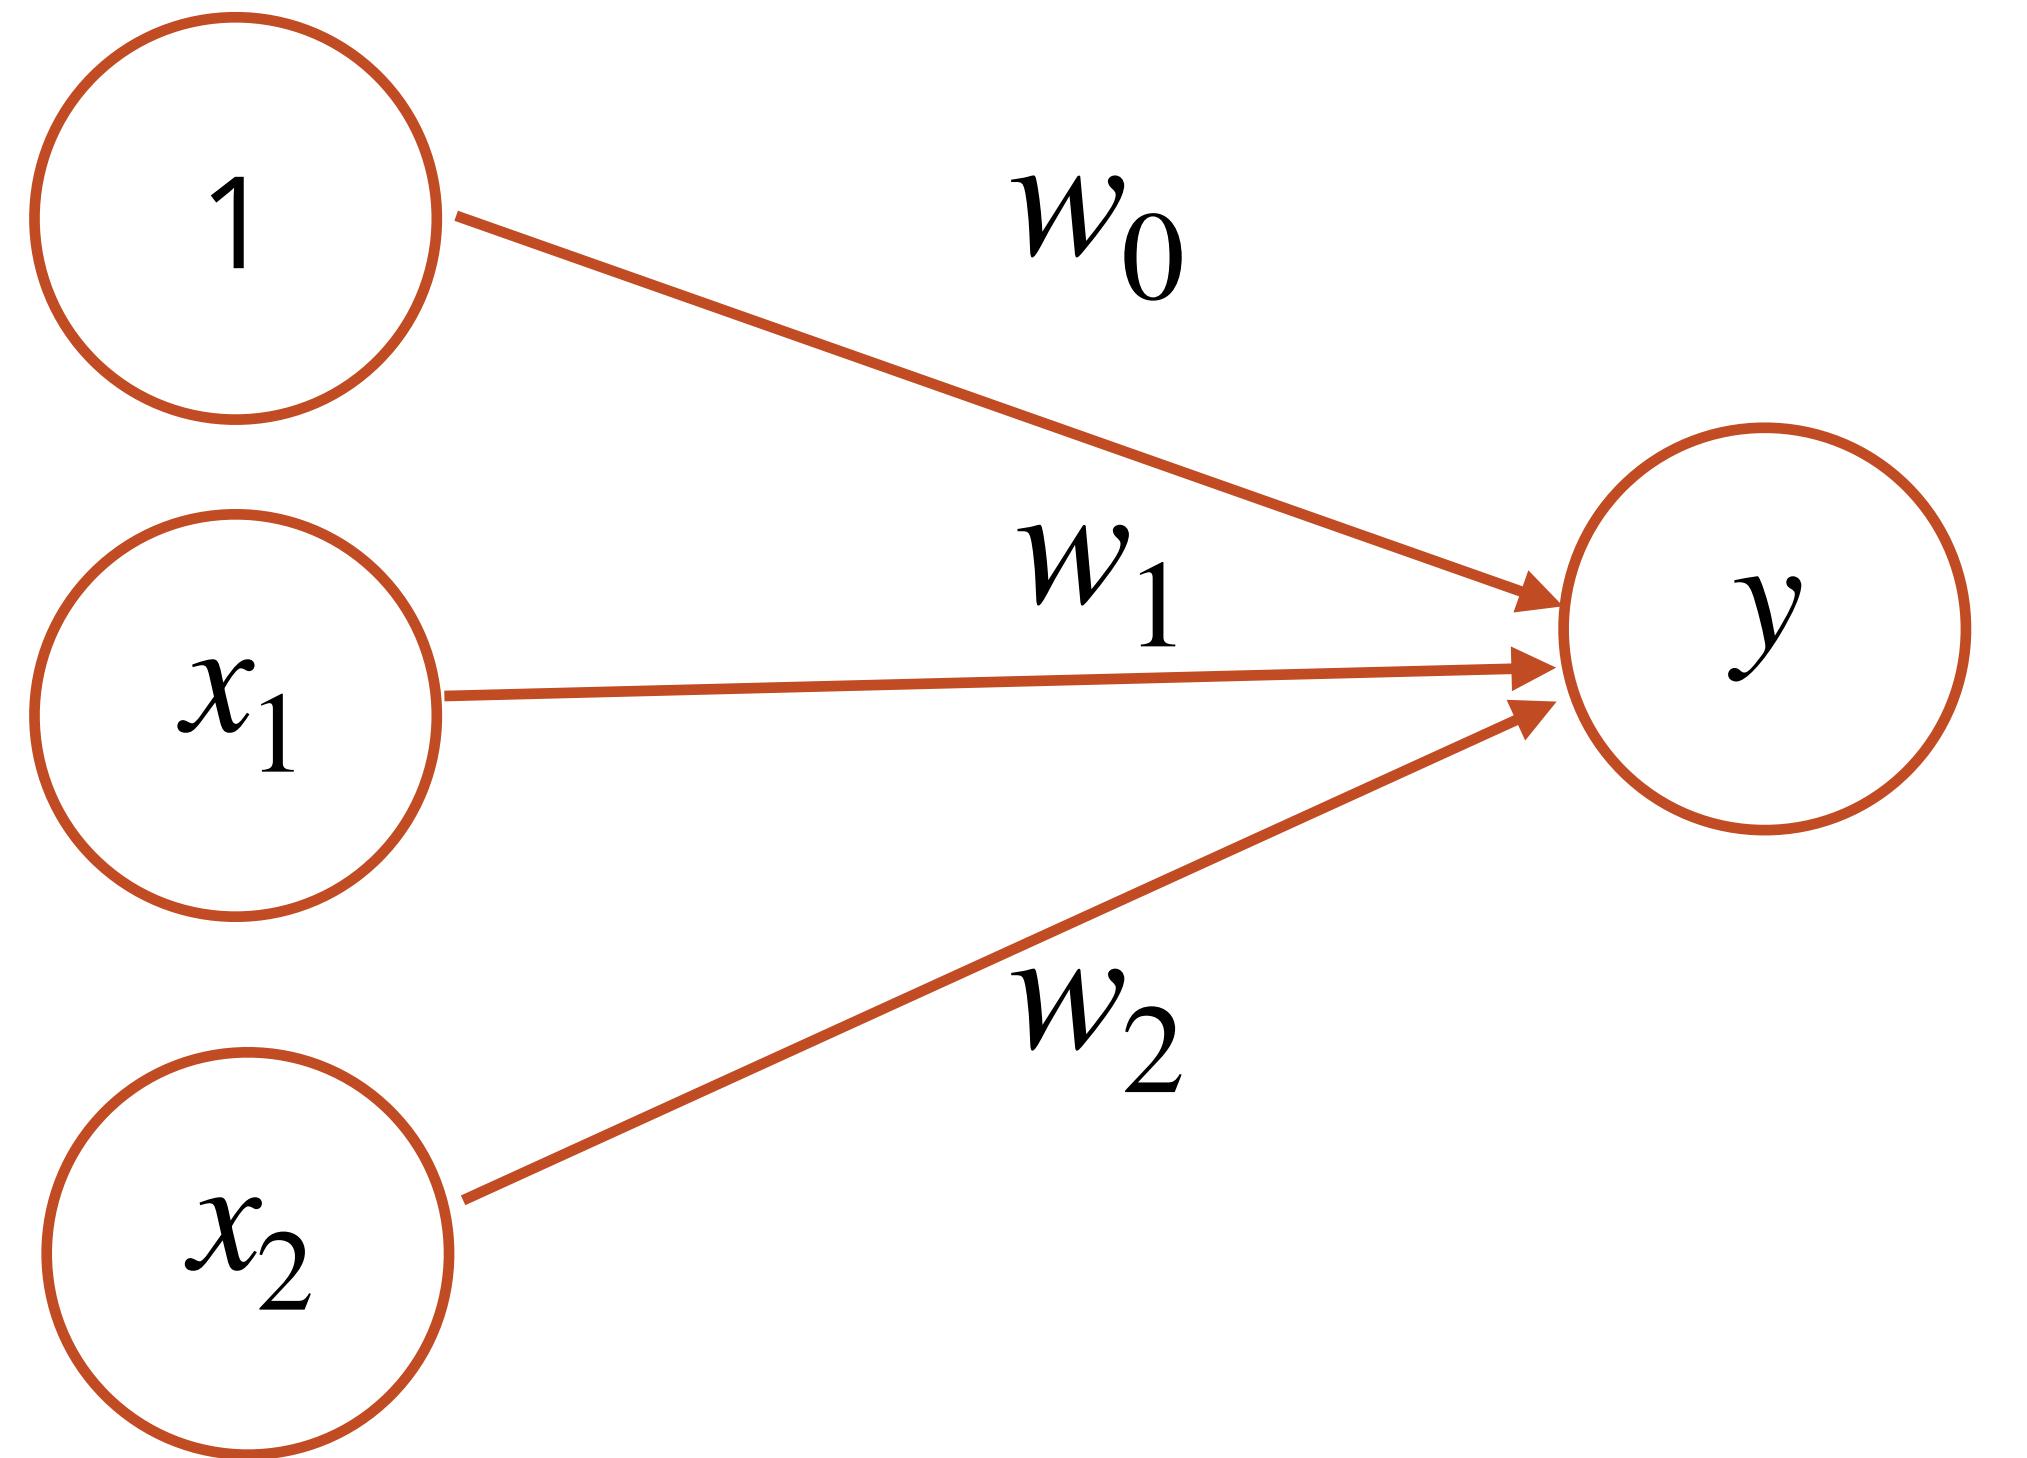
Building an AND network



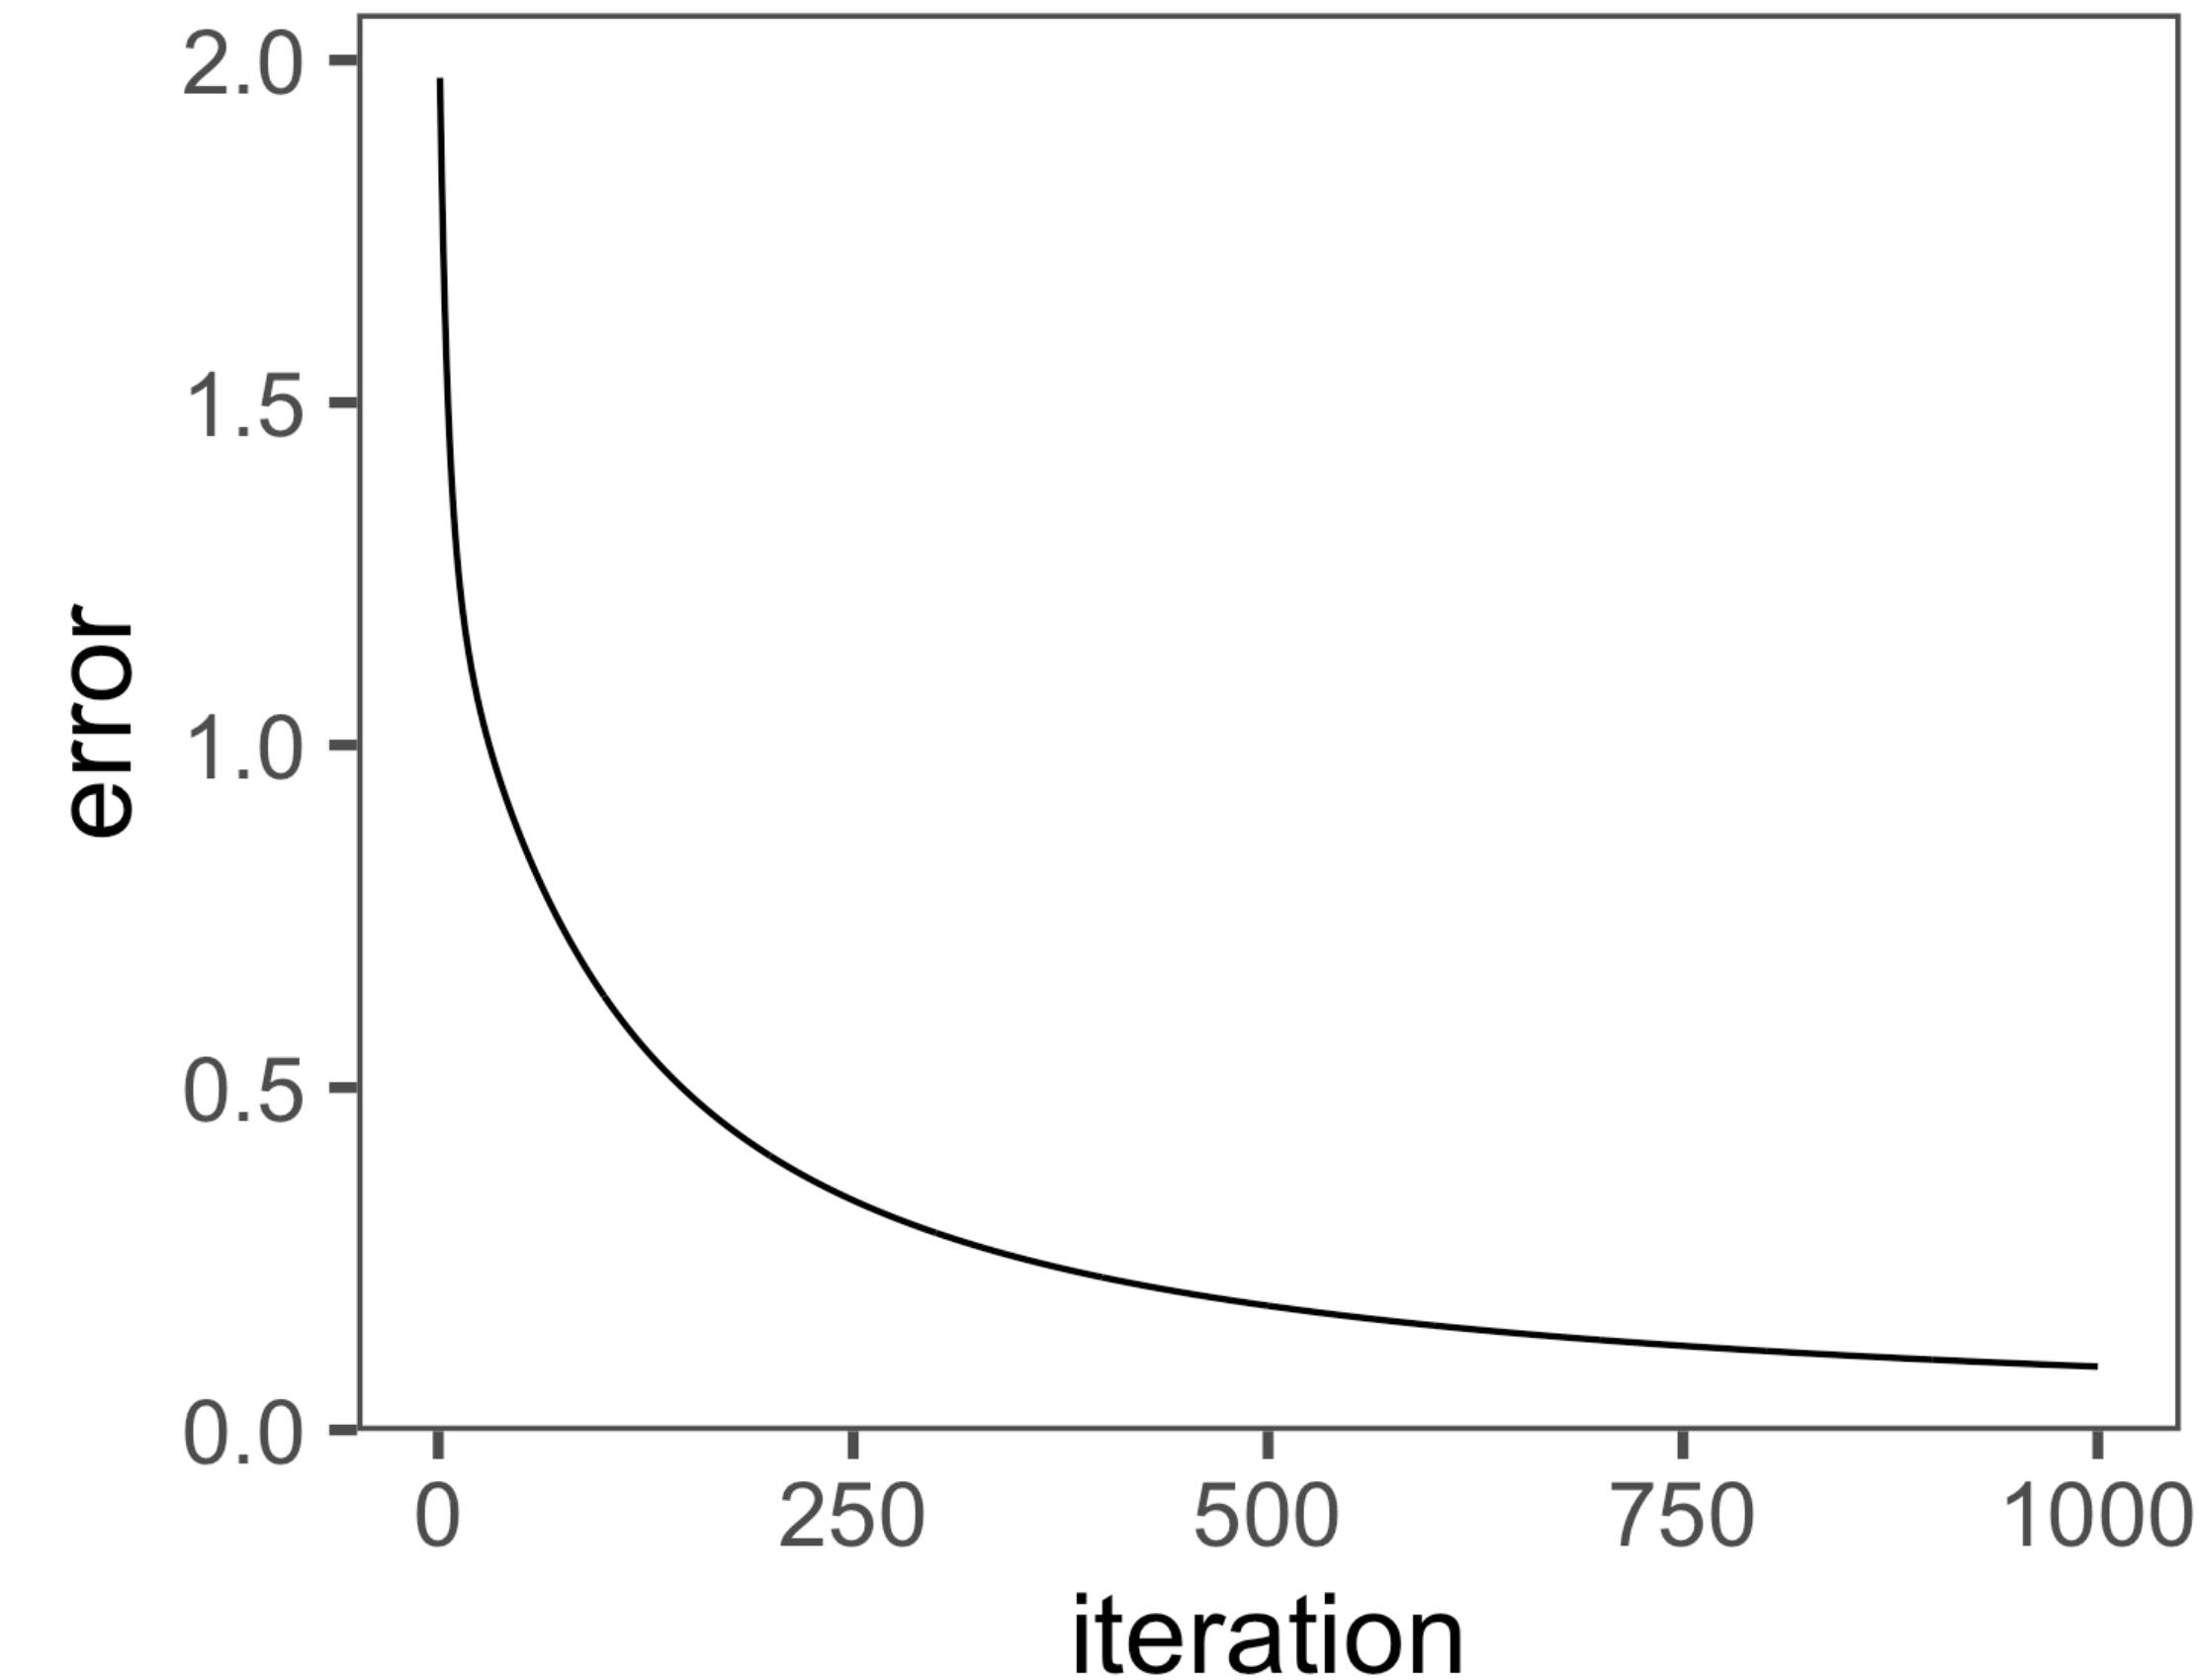
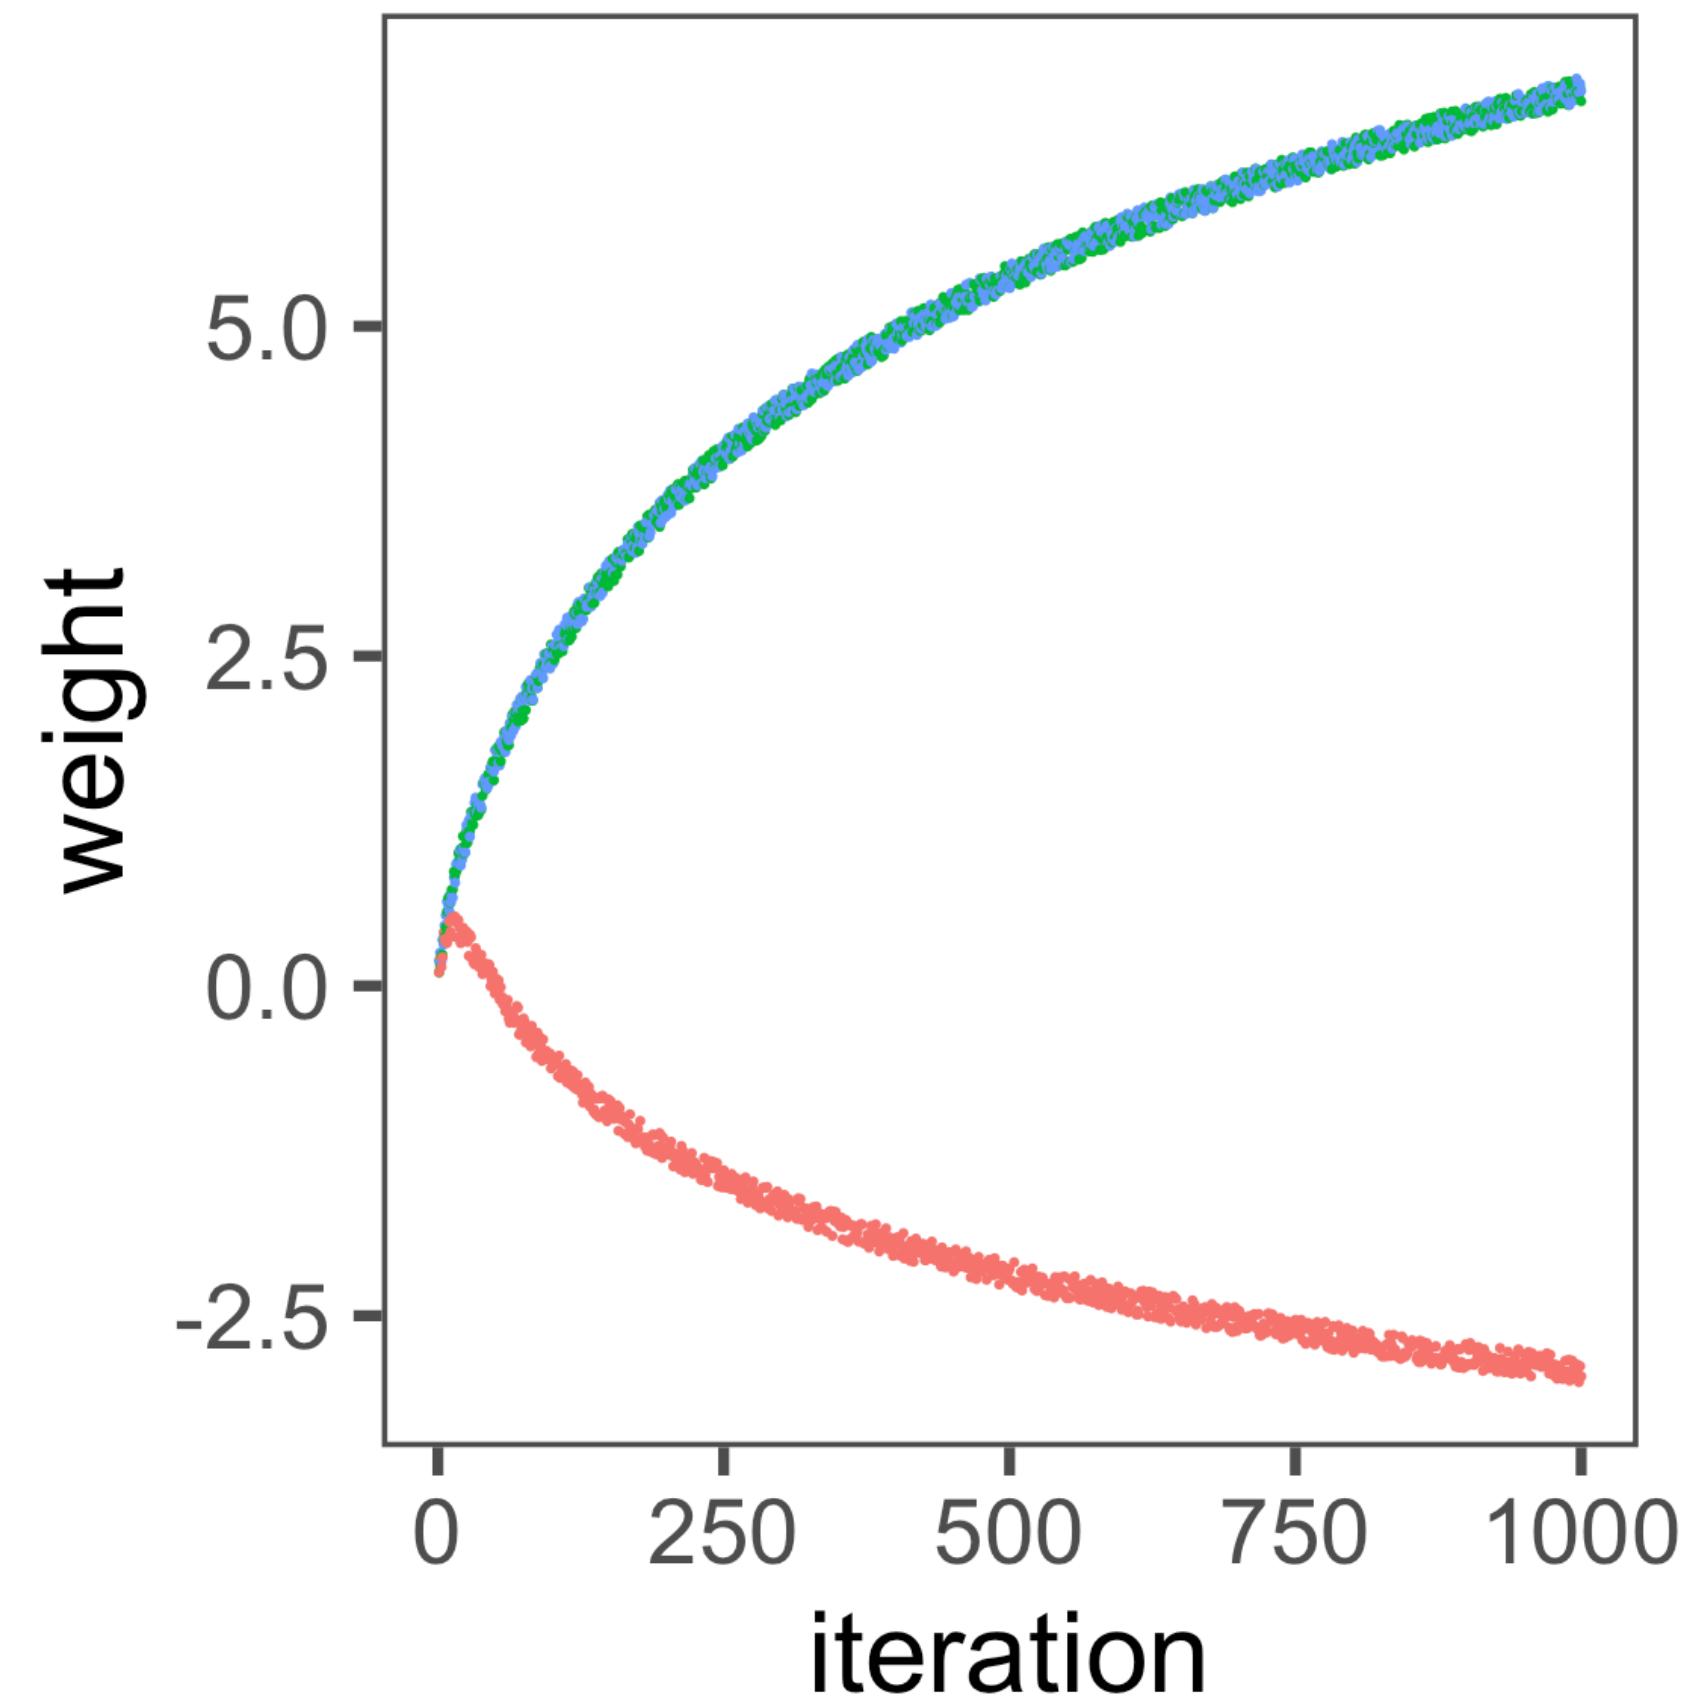
Learning the AND function



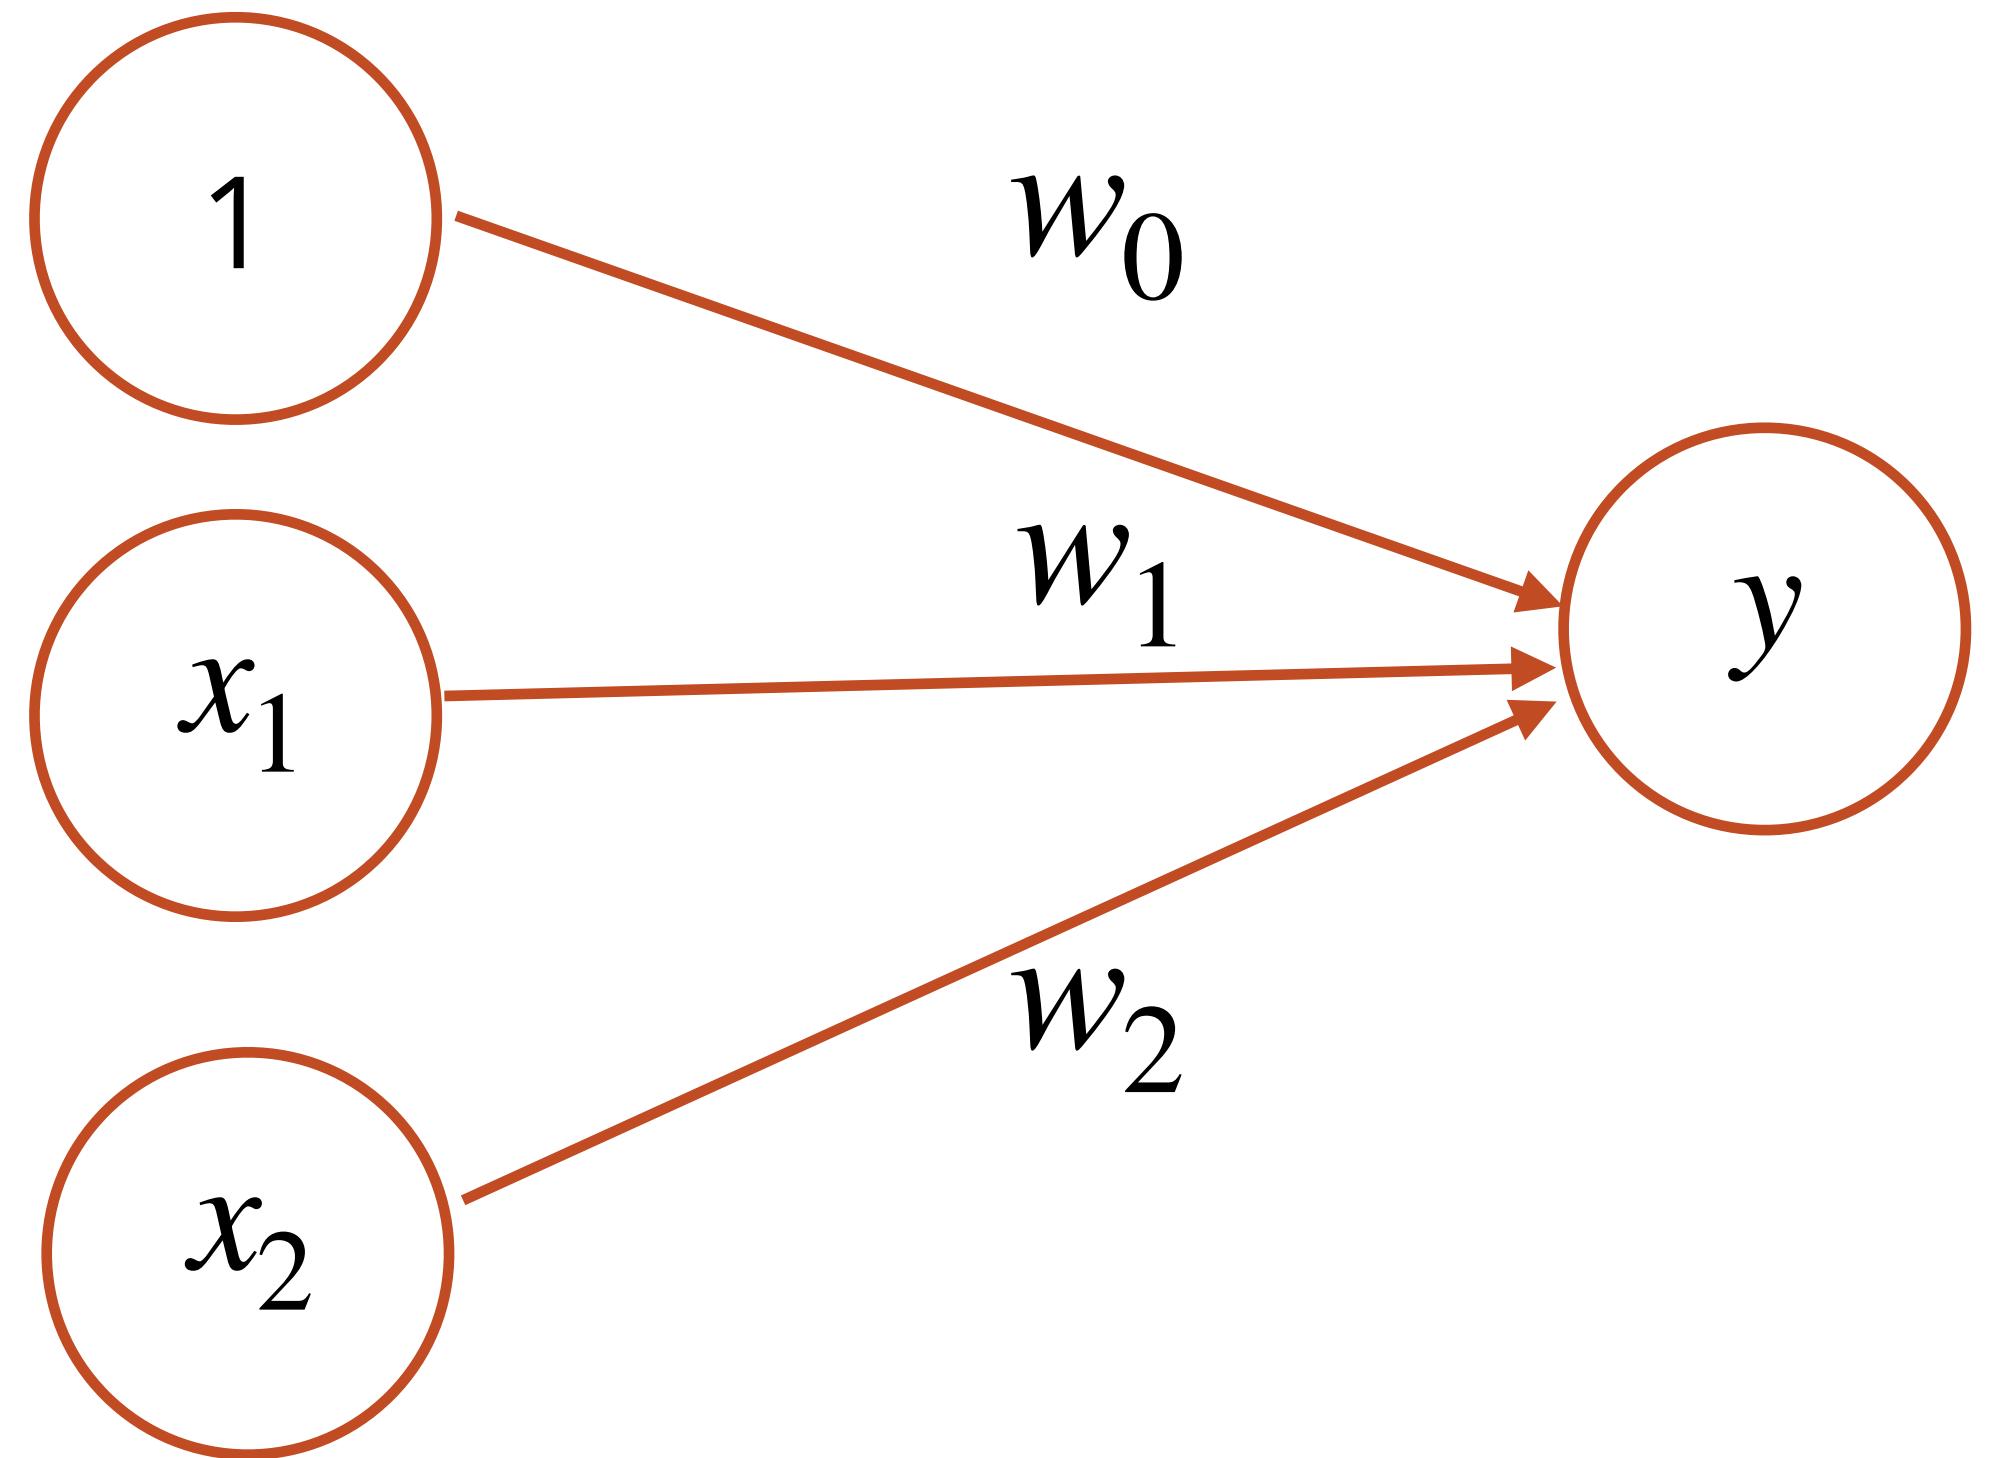
Building an OR network



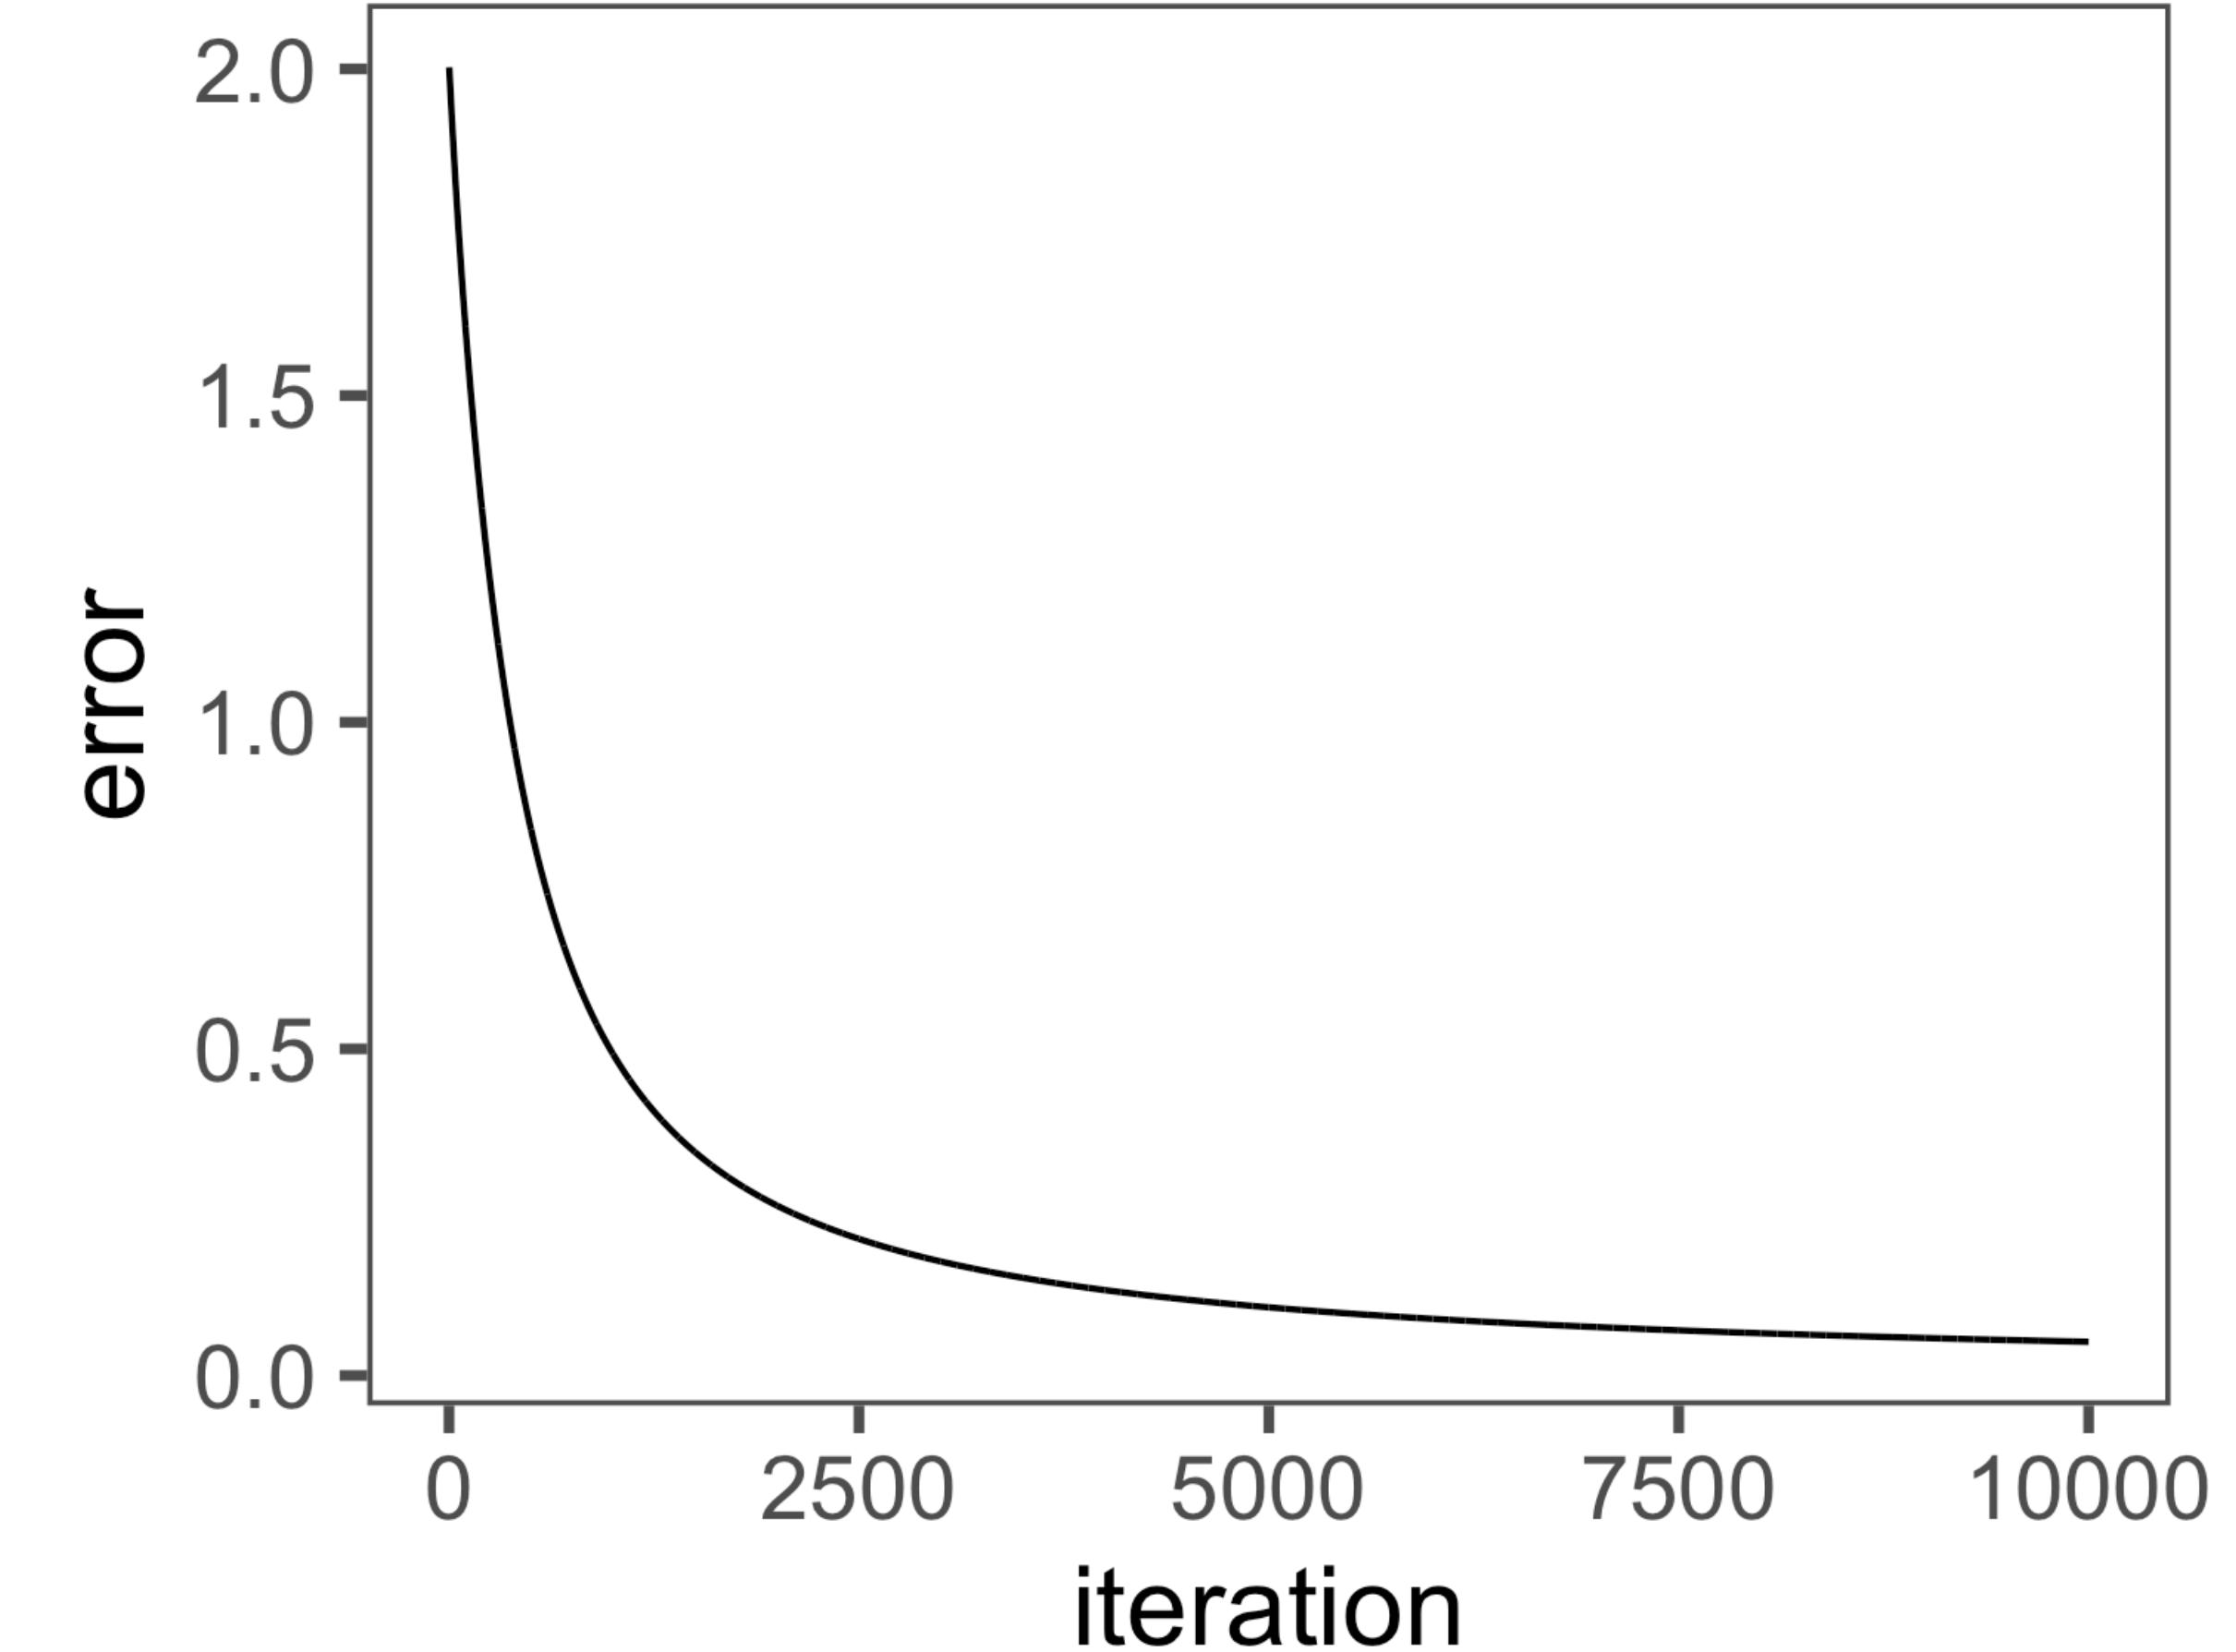
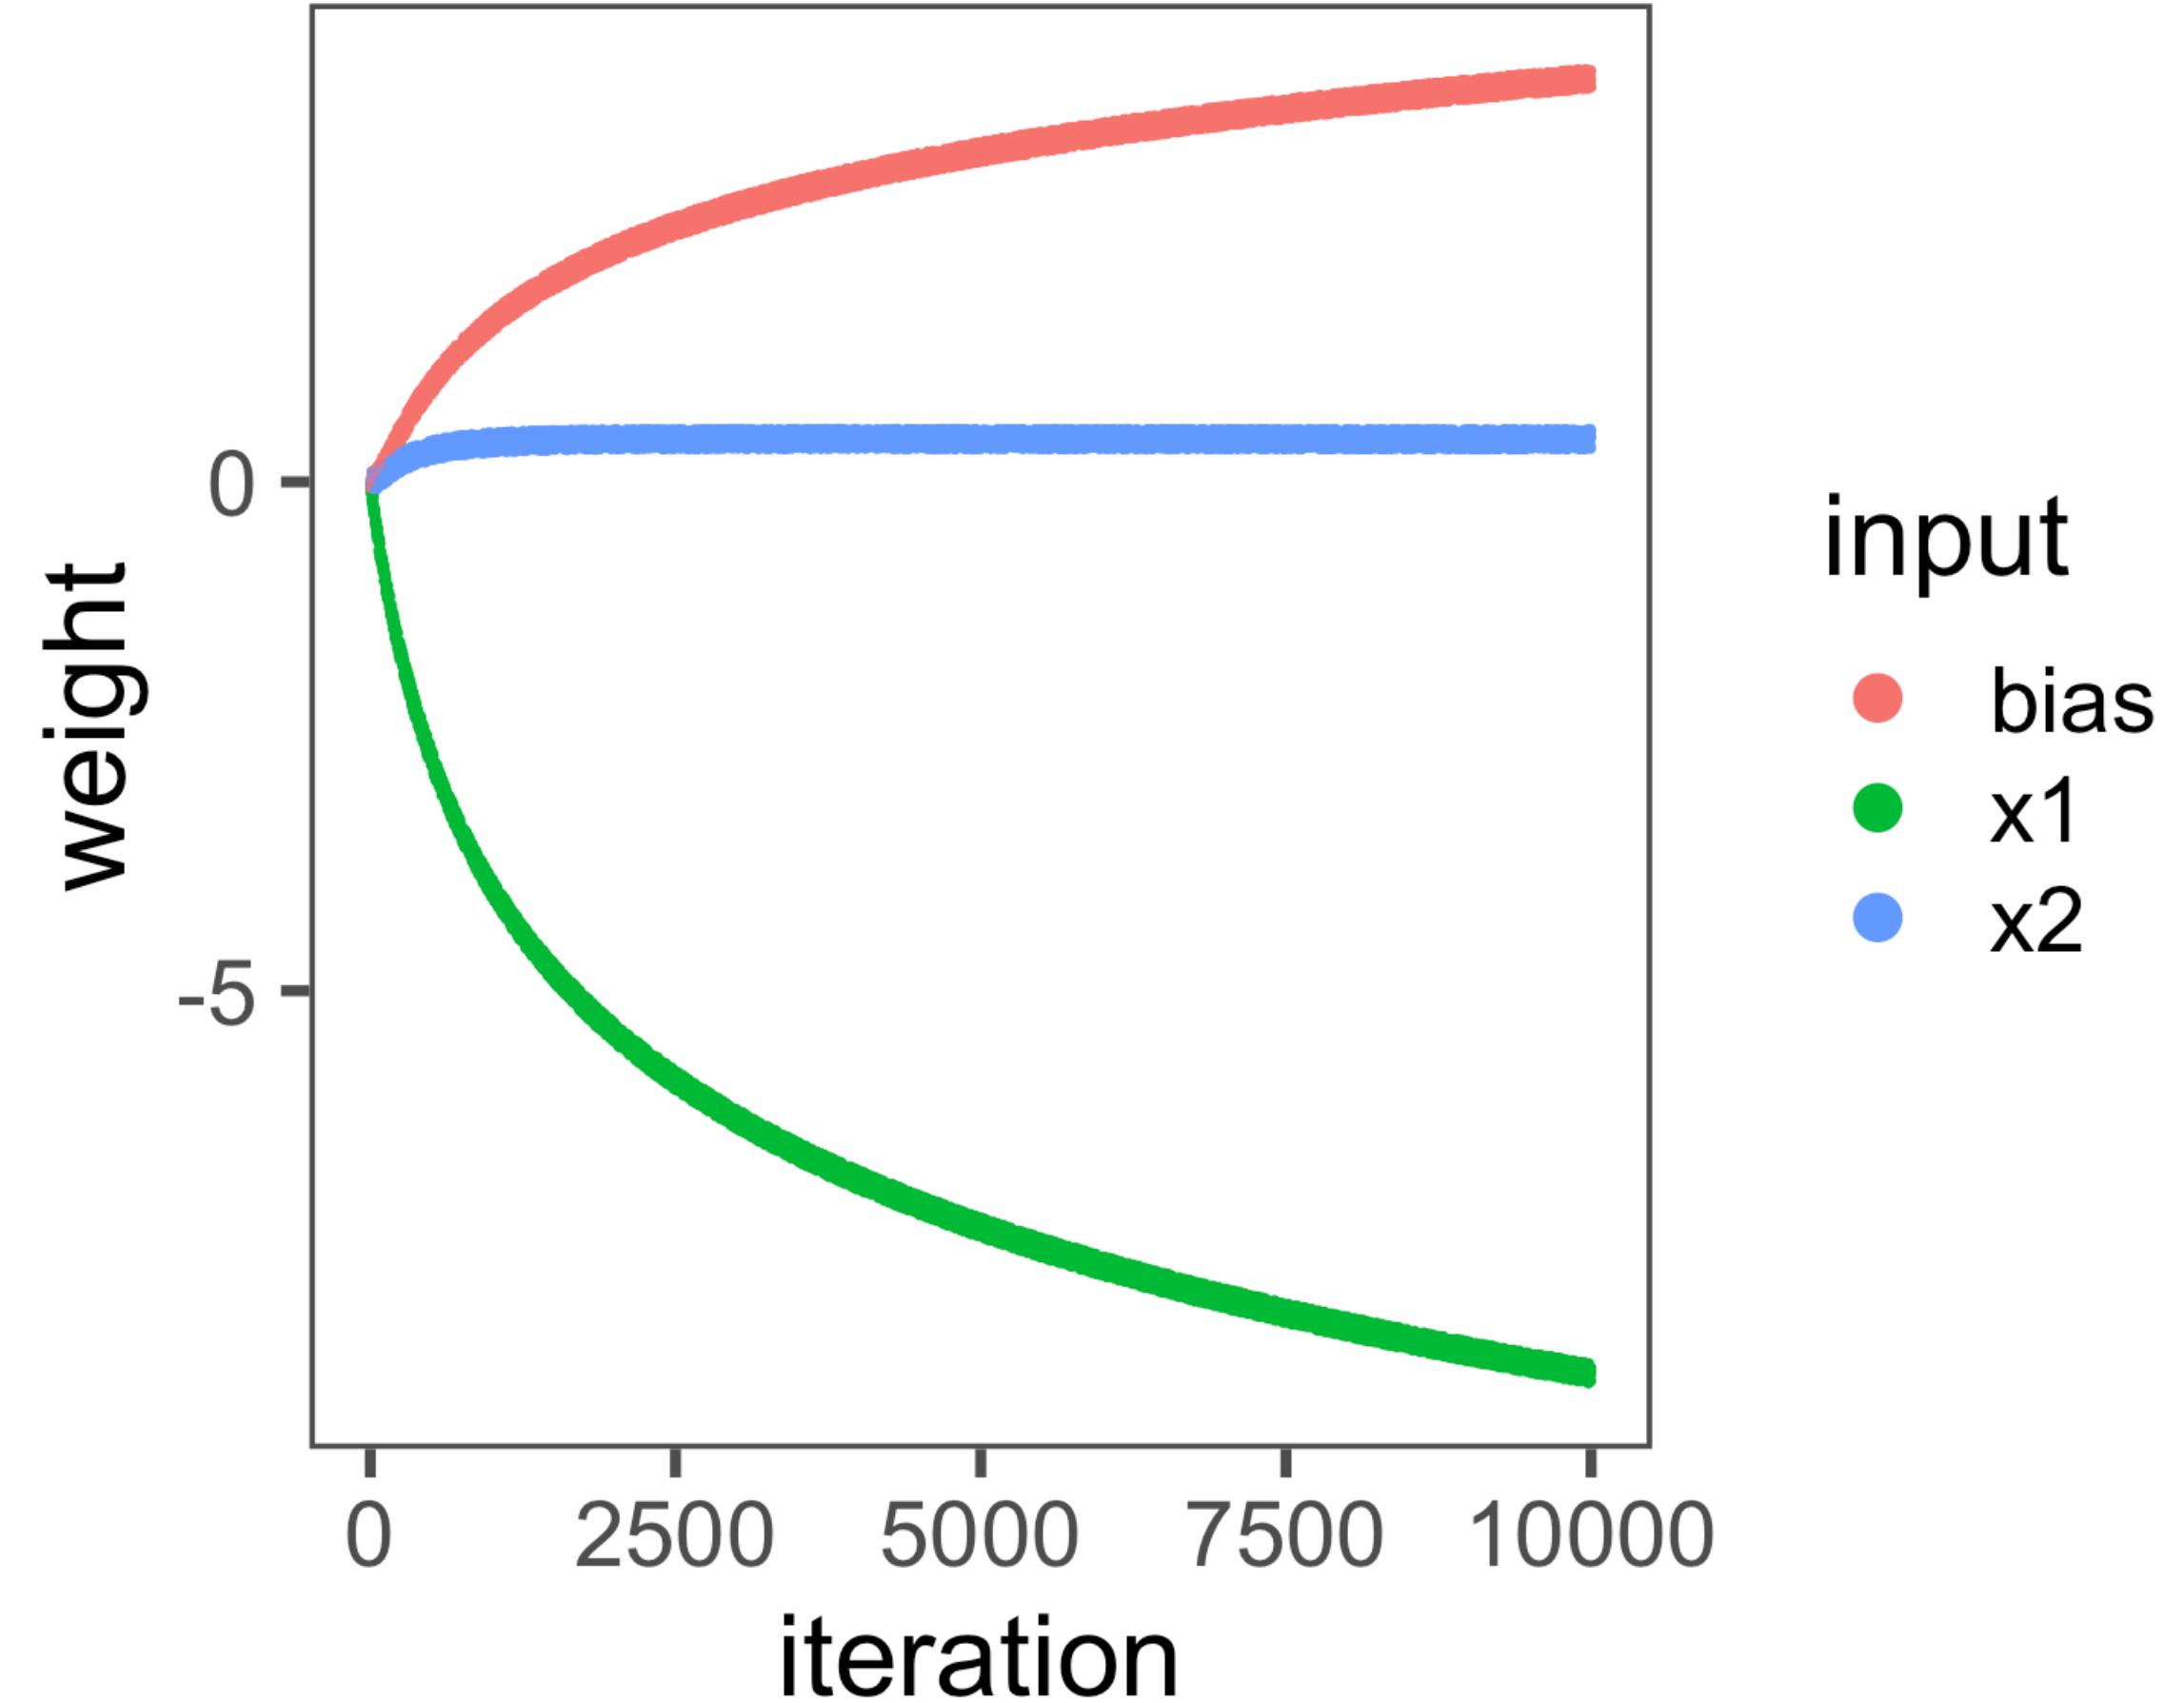
Learning the OR function



Building a NOT(x_1) network



Learning the NOT(x) function



Perceptrons as general classifiers



iris virginica

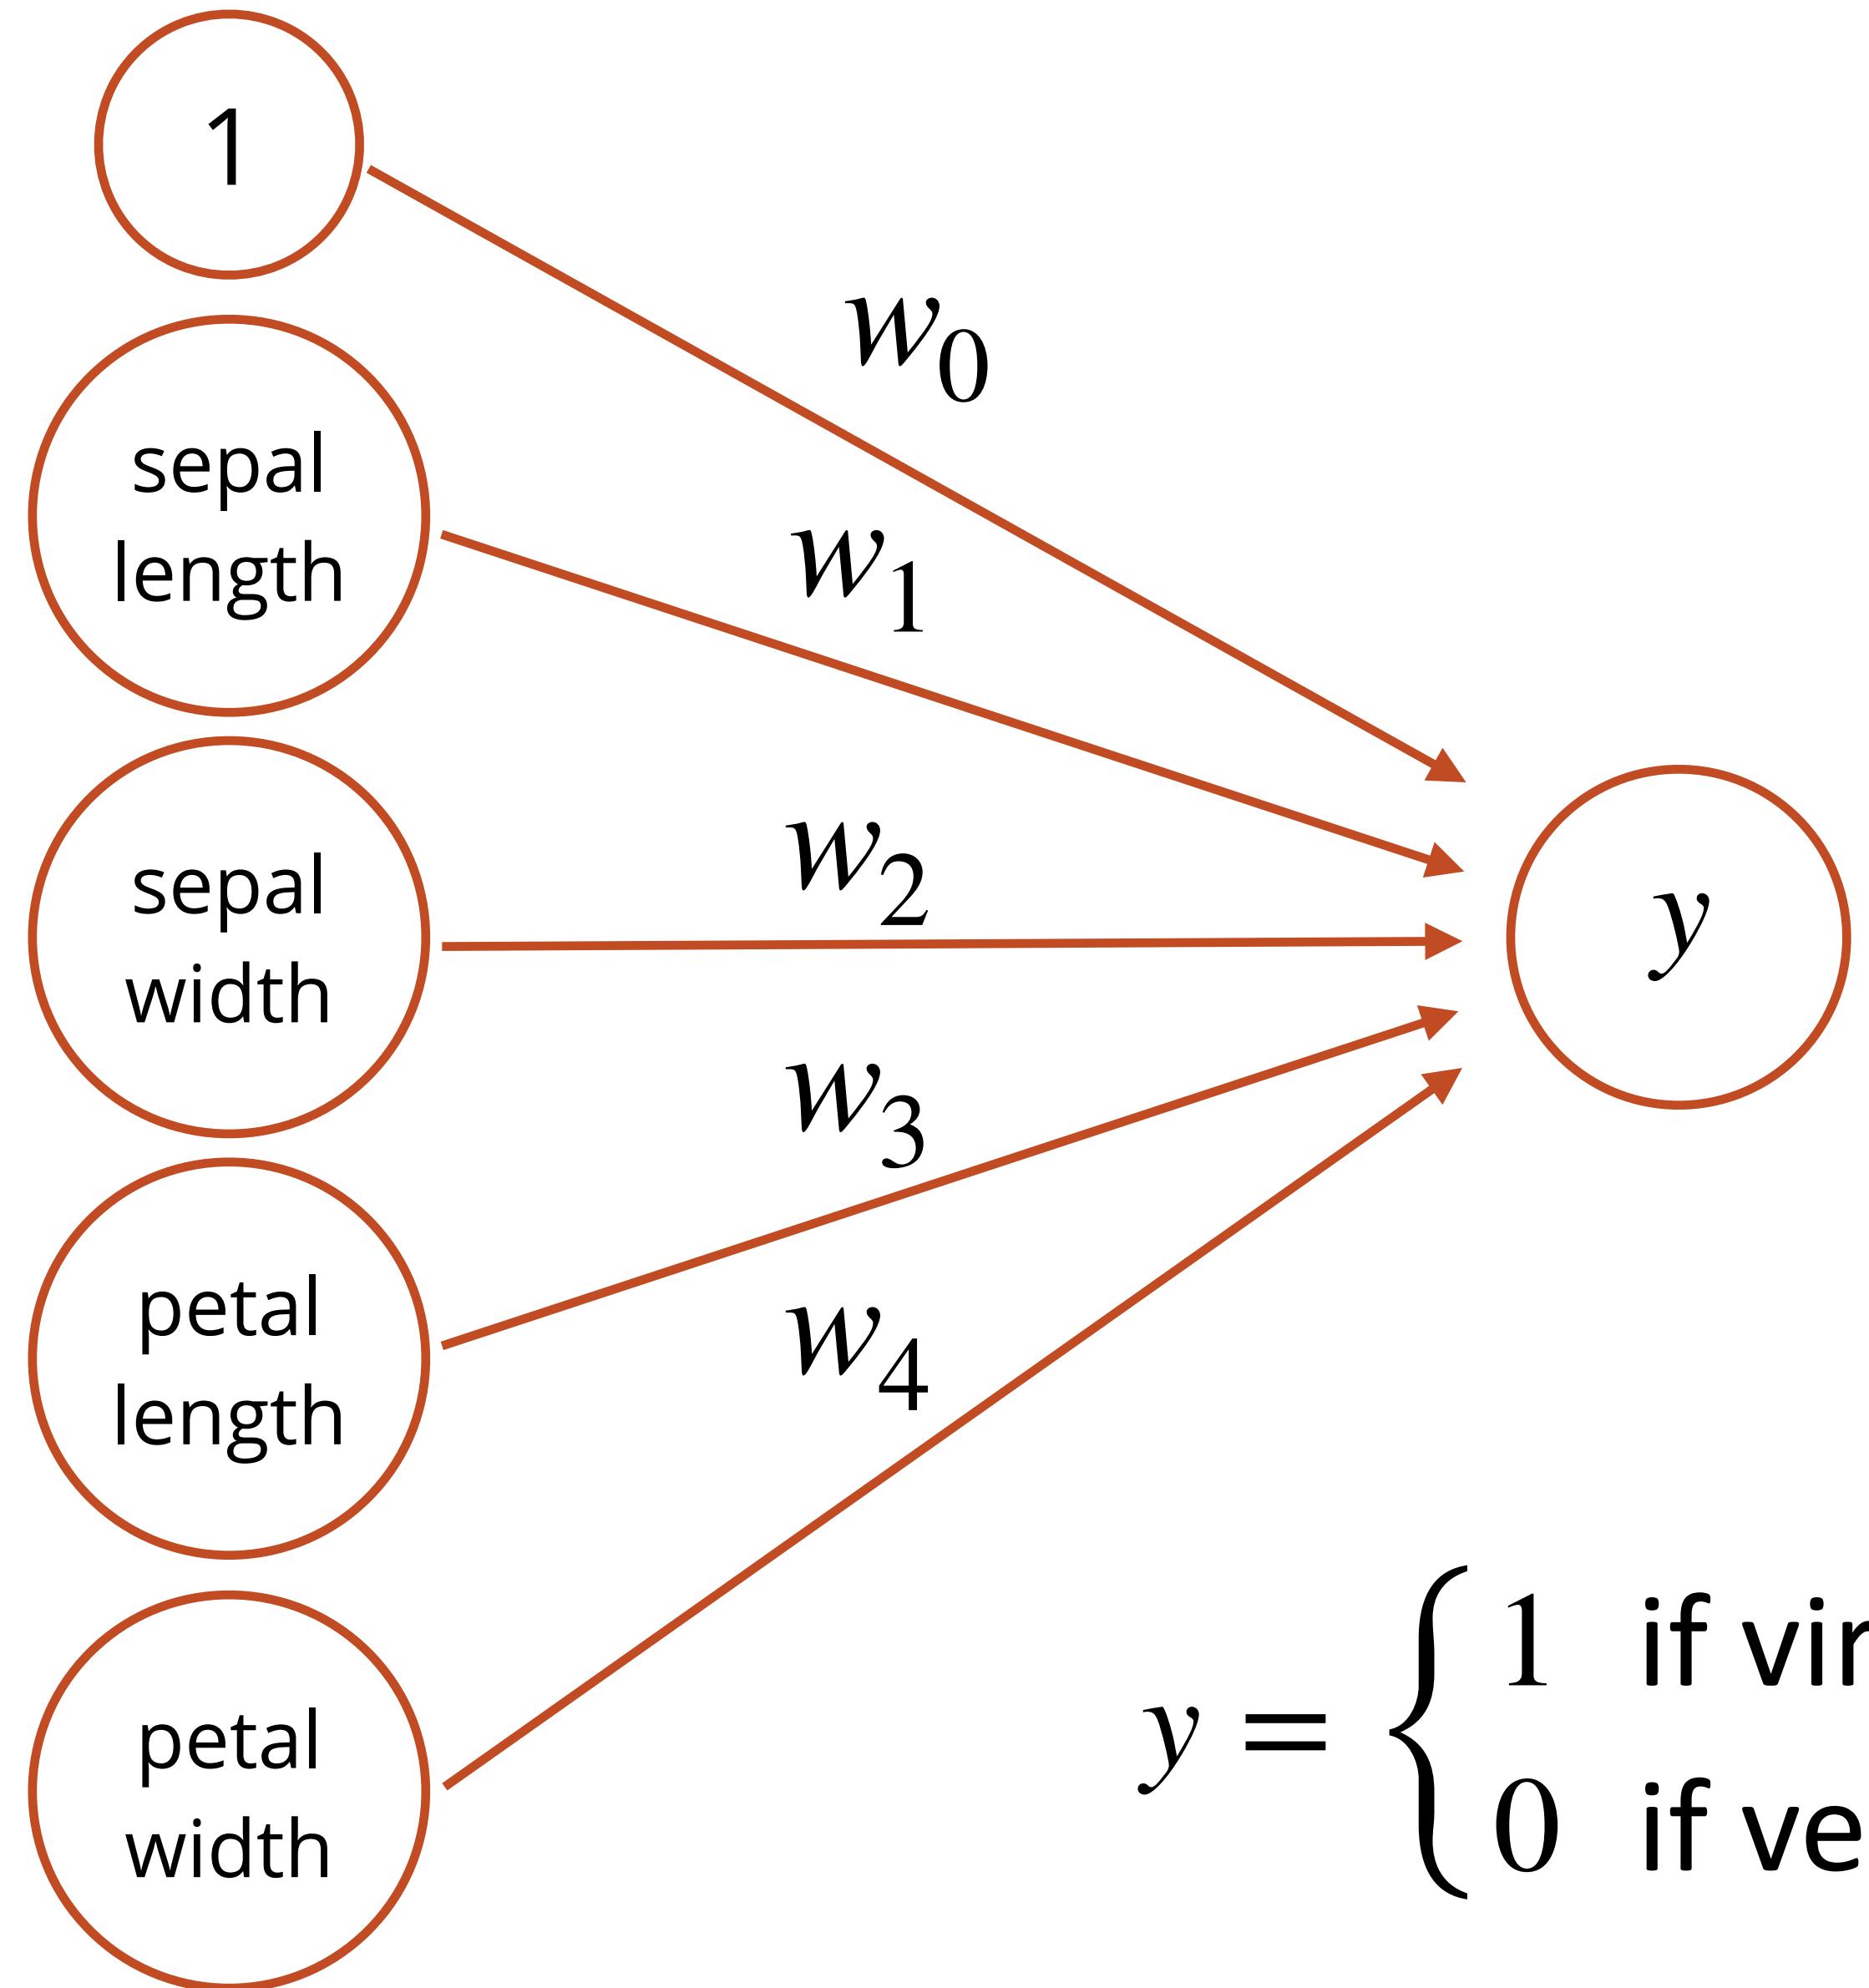


iris versicolor

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
7	3.2	4.7	1.4	0
6.4	3.2	4.5	1.5	0
6.9	3.1	4.9	1.5	0
5.5	2.3	4	1.3	0
6.5	2.8	4.6	1.5	0
5.7	2.8	4.5	1.3	0
6.3	3.3	4.7	1.6	0
4.9	2.4	3.3	1	0
6.6	2.9	4.6	1.3	0
5.2	2.7	3.9	1.4	0
5	2	3.5	1	0

Fisher (1936)

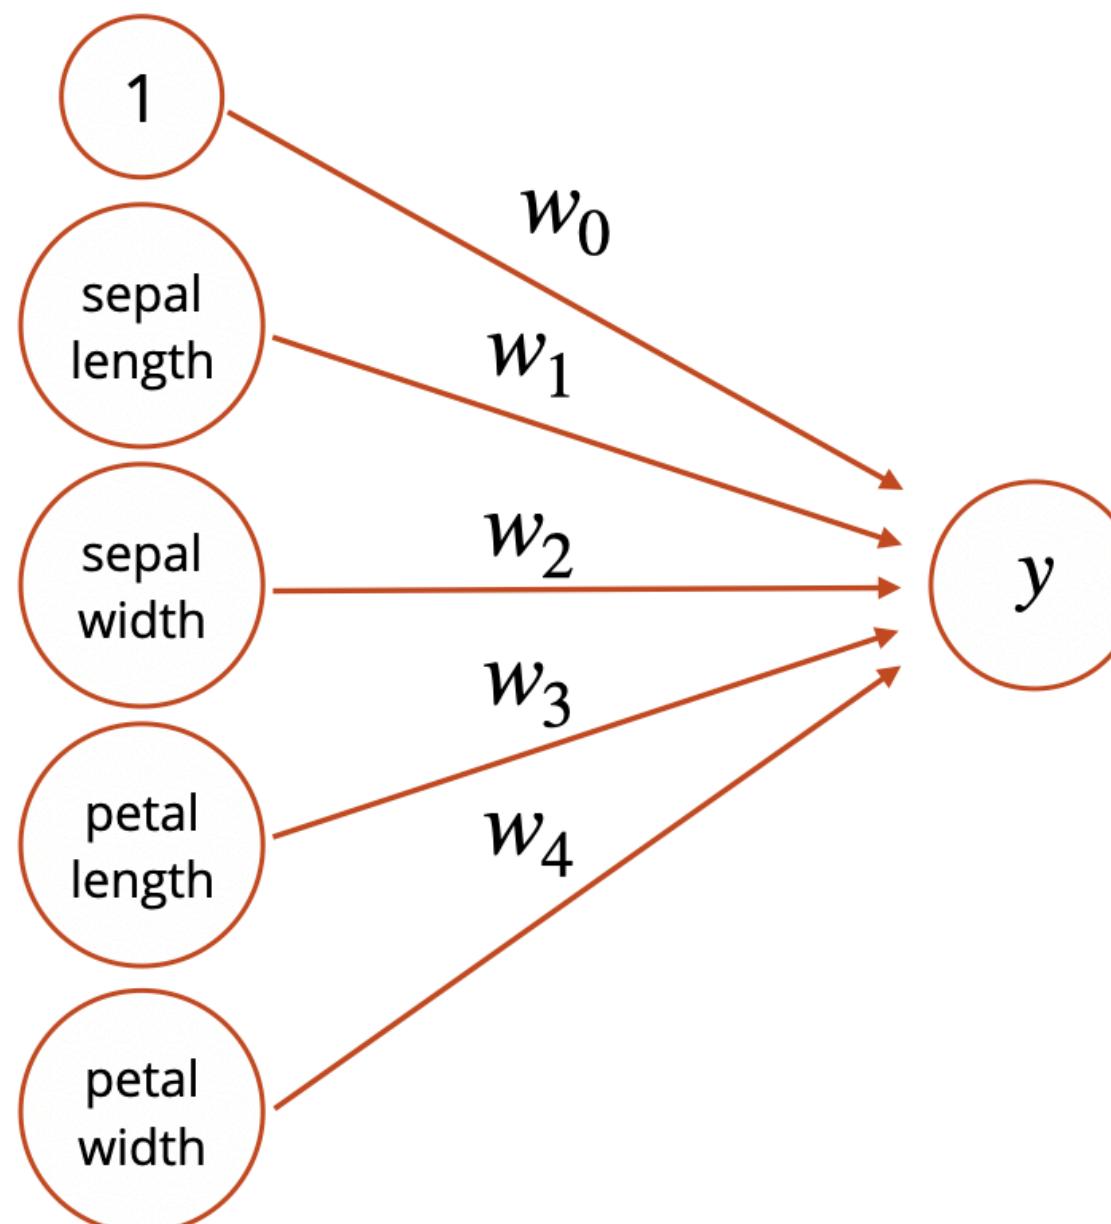
Building an iris classifier



$$y = \begin{cases} 1 & \text{if virginica,} \\ 0 & \text{if versicolor} \end{cases}$$

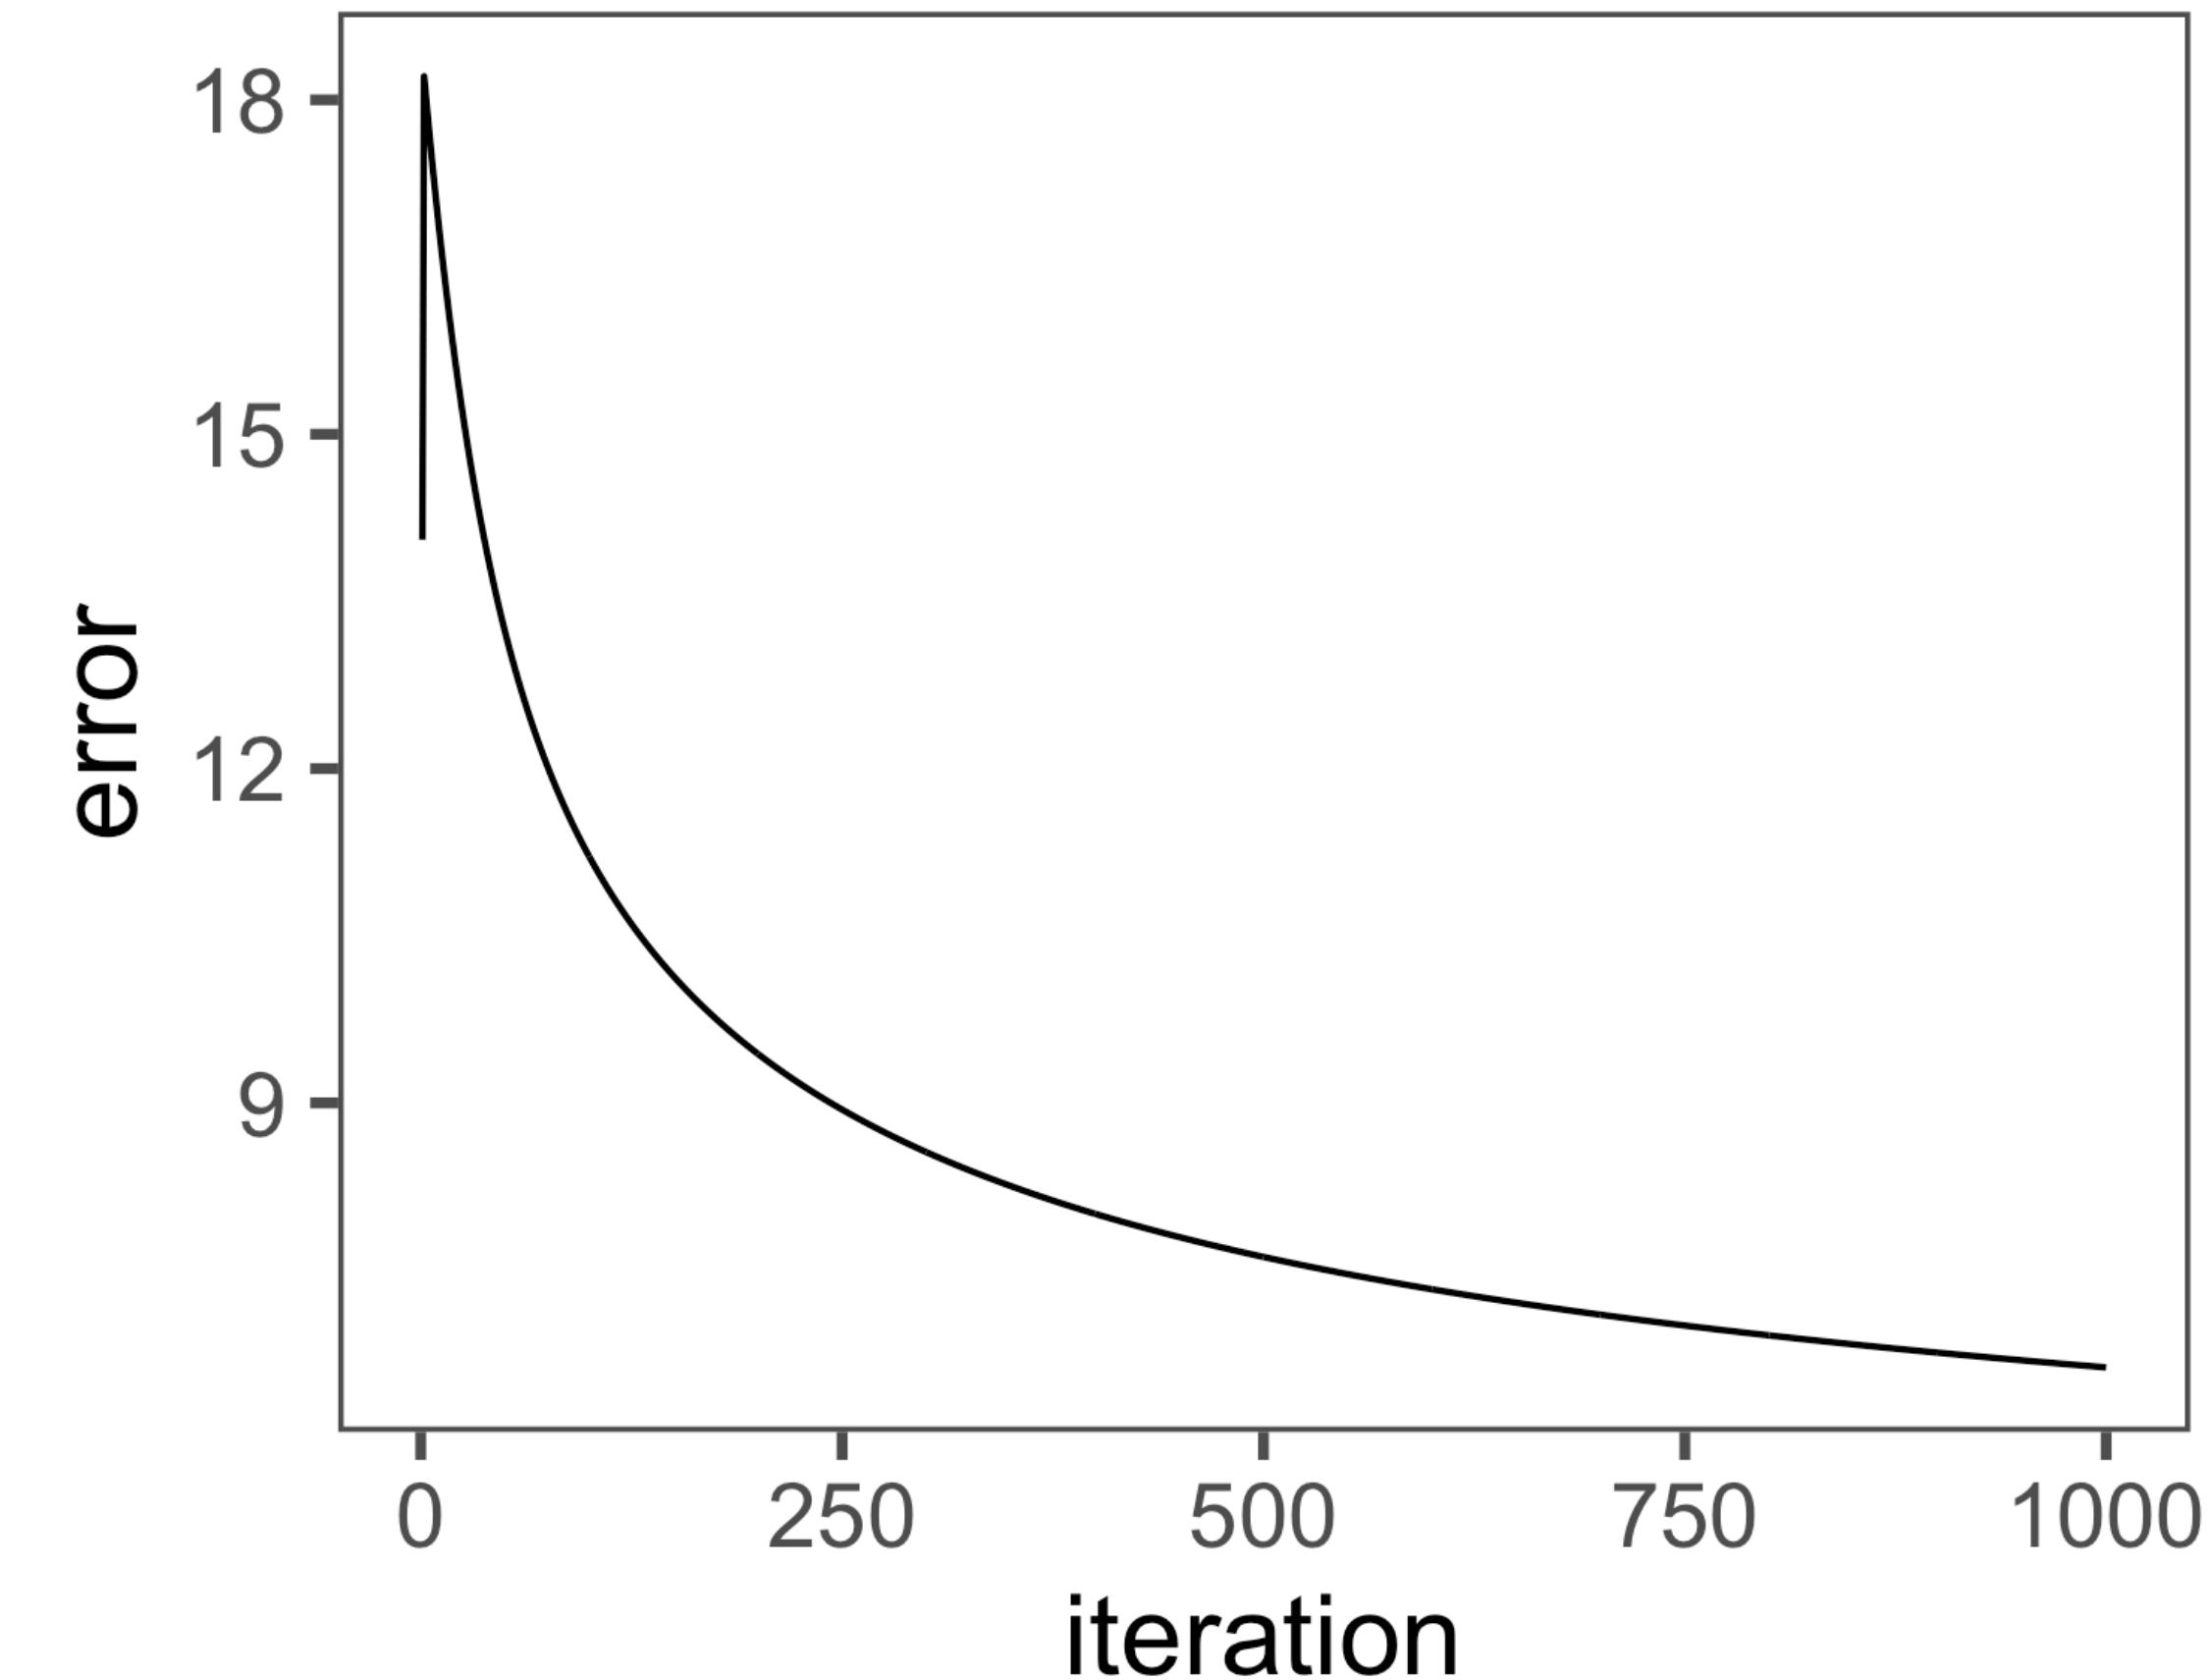
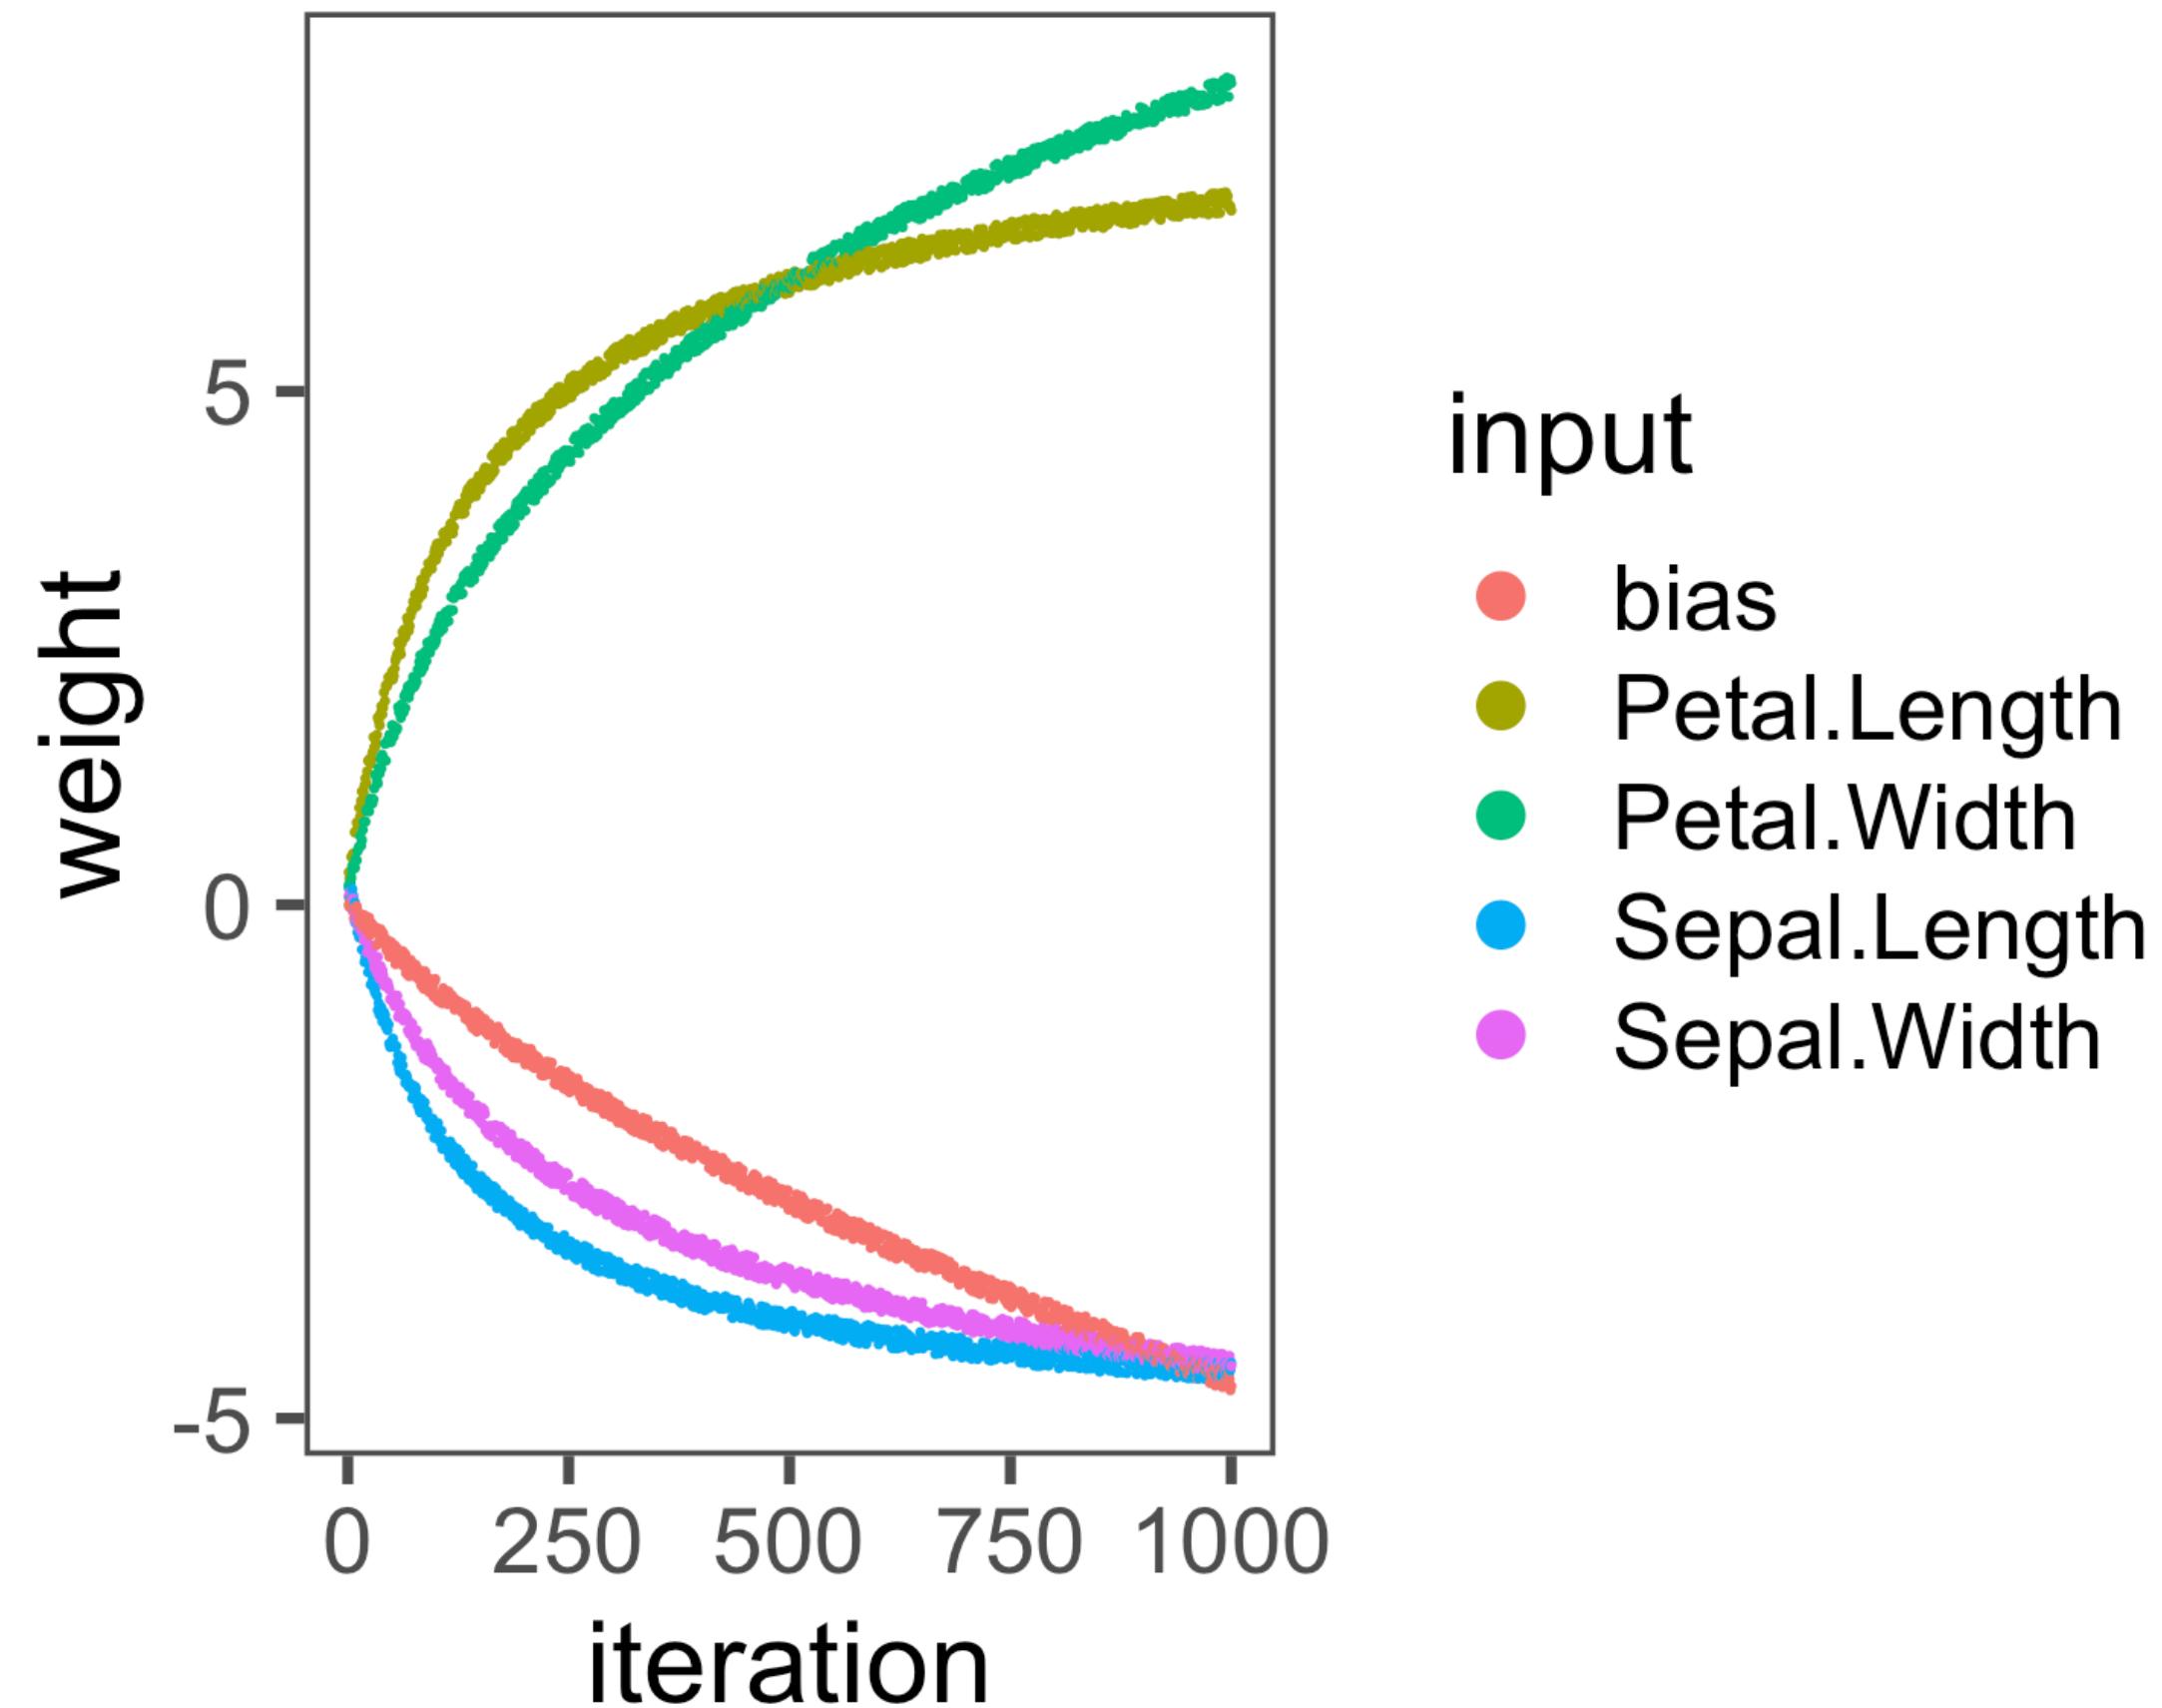
Logistic regression as an iris classifier

```
glm(Species ~ Sepal.Length + Sepal.Width  
+ Petal.Length + Petal.Width,  
family = "binomial")
```

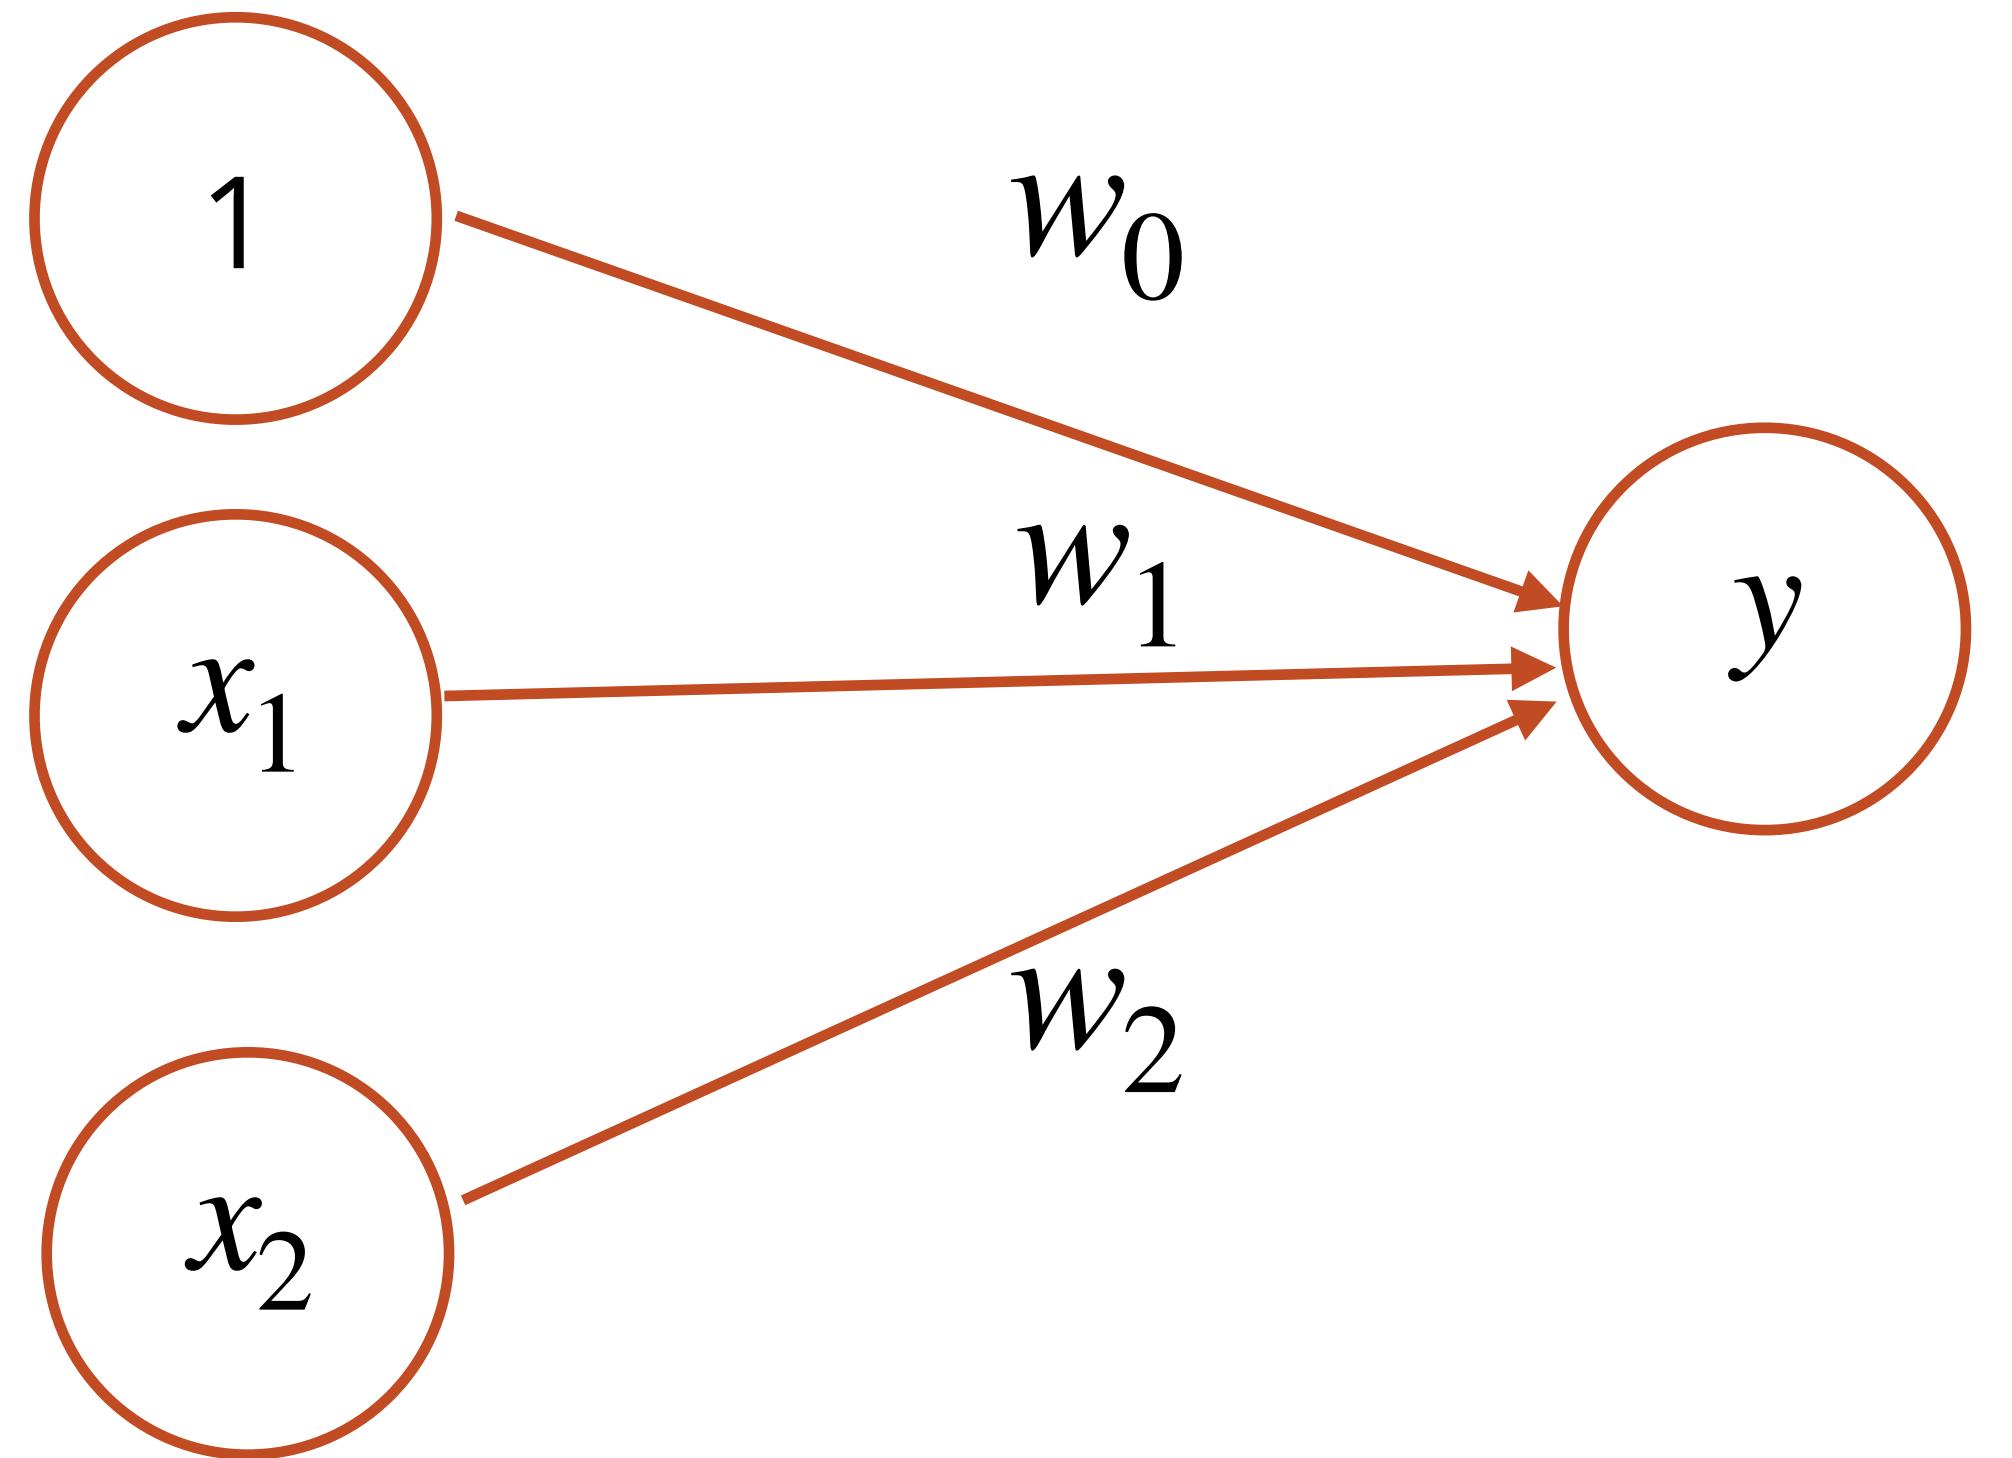


term	estimate	std.error	statistic	p.value
(Intercept)	-42.638	25.707	-1.659	.097
Sepal.Length	-2.465	2.394	-1.030	.303
Sepal.Width	-6.681	4.480	-1.491	.136
Petal.Length	9.429	4.737	1.991	.047
Petal.Width	18.286	9.743	1.877	.061

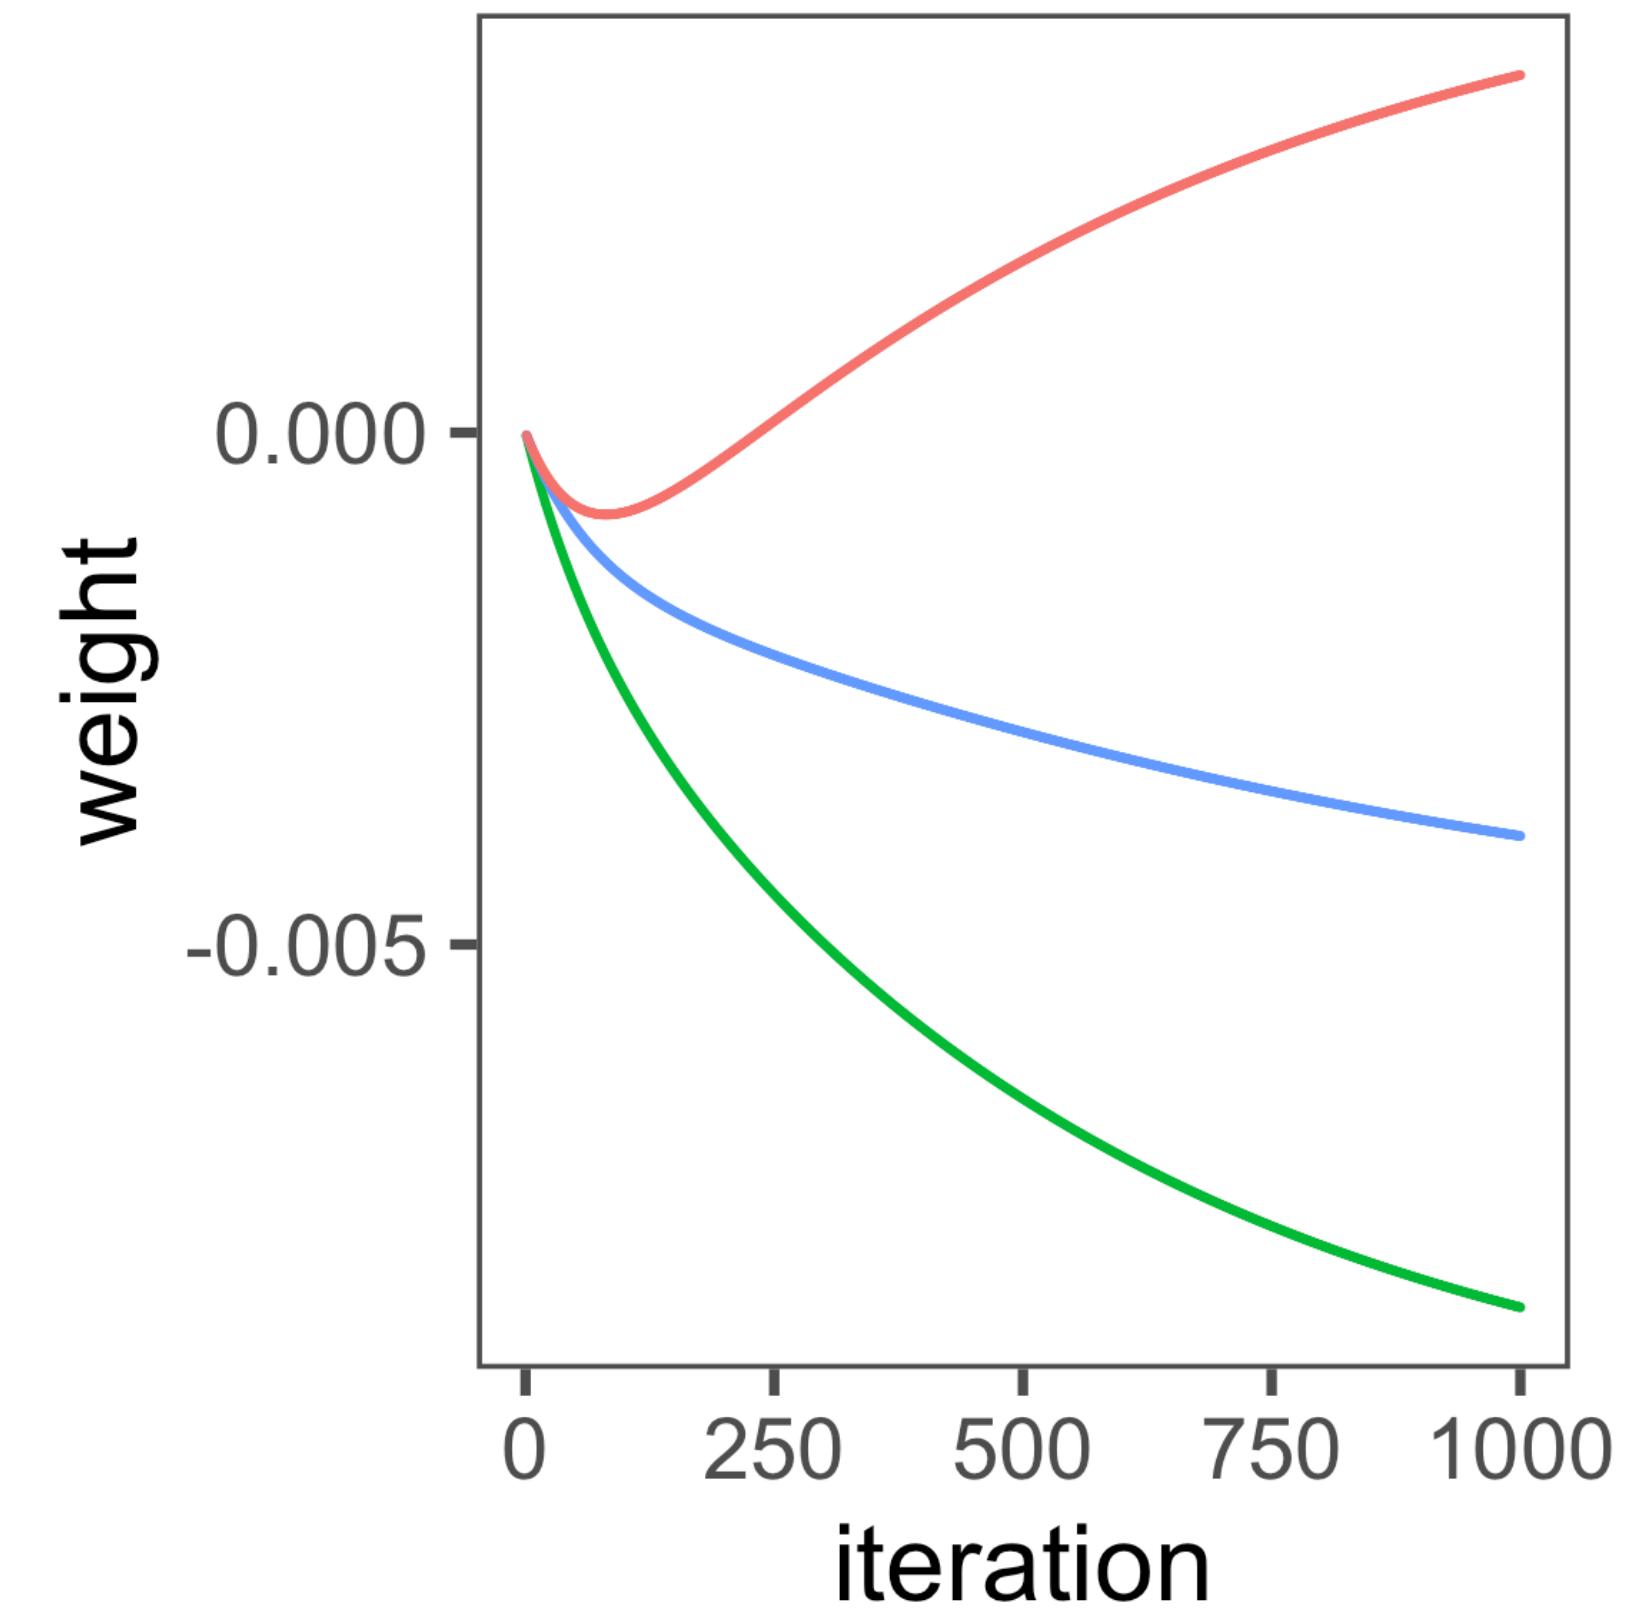
Learning an iris classifier



Building an XOR network

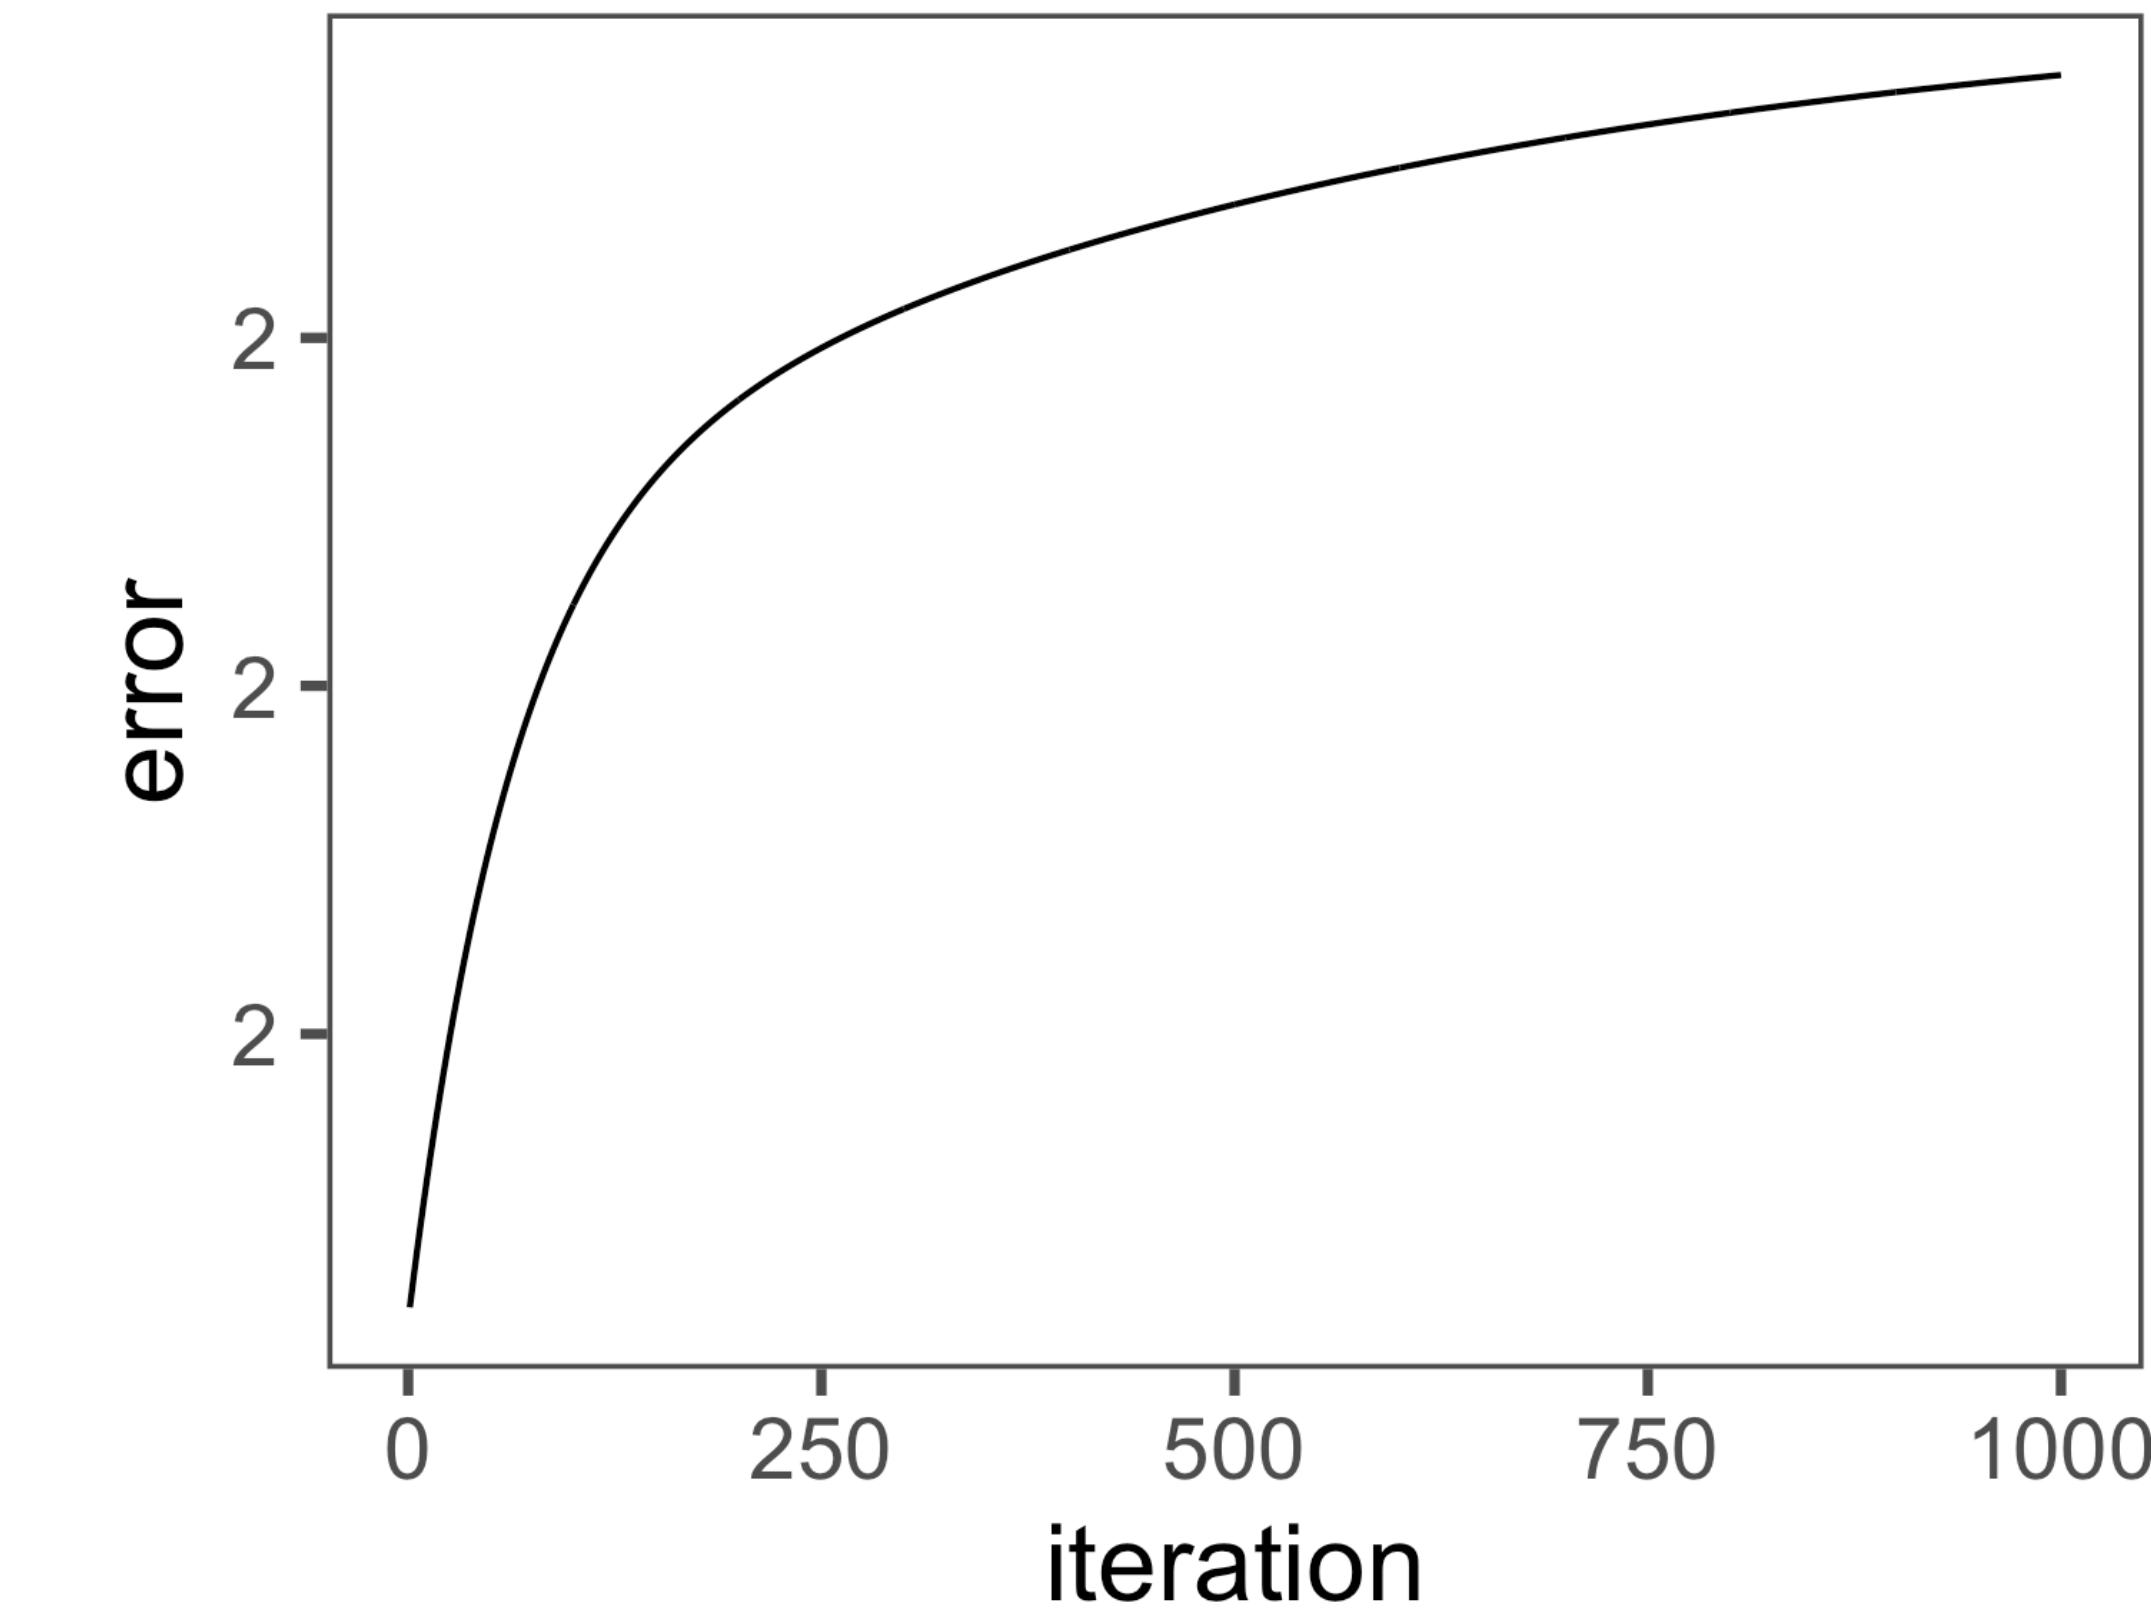


Learning the XOR function



input

- bias
- x_1
- x_2



Why can't this network learn XOR?

QUESTION

ANSWER

How would regression solve xor?

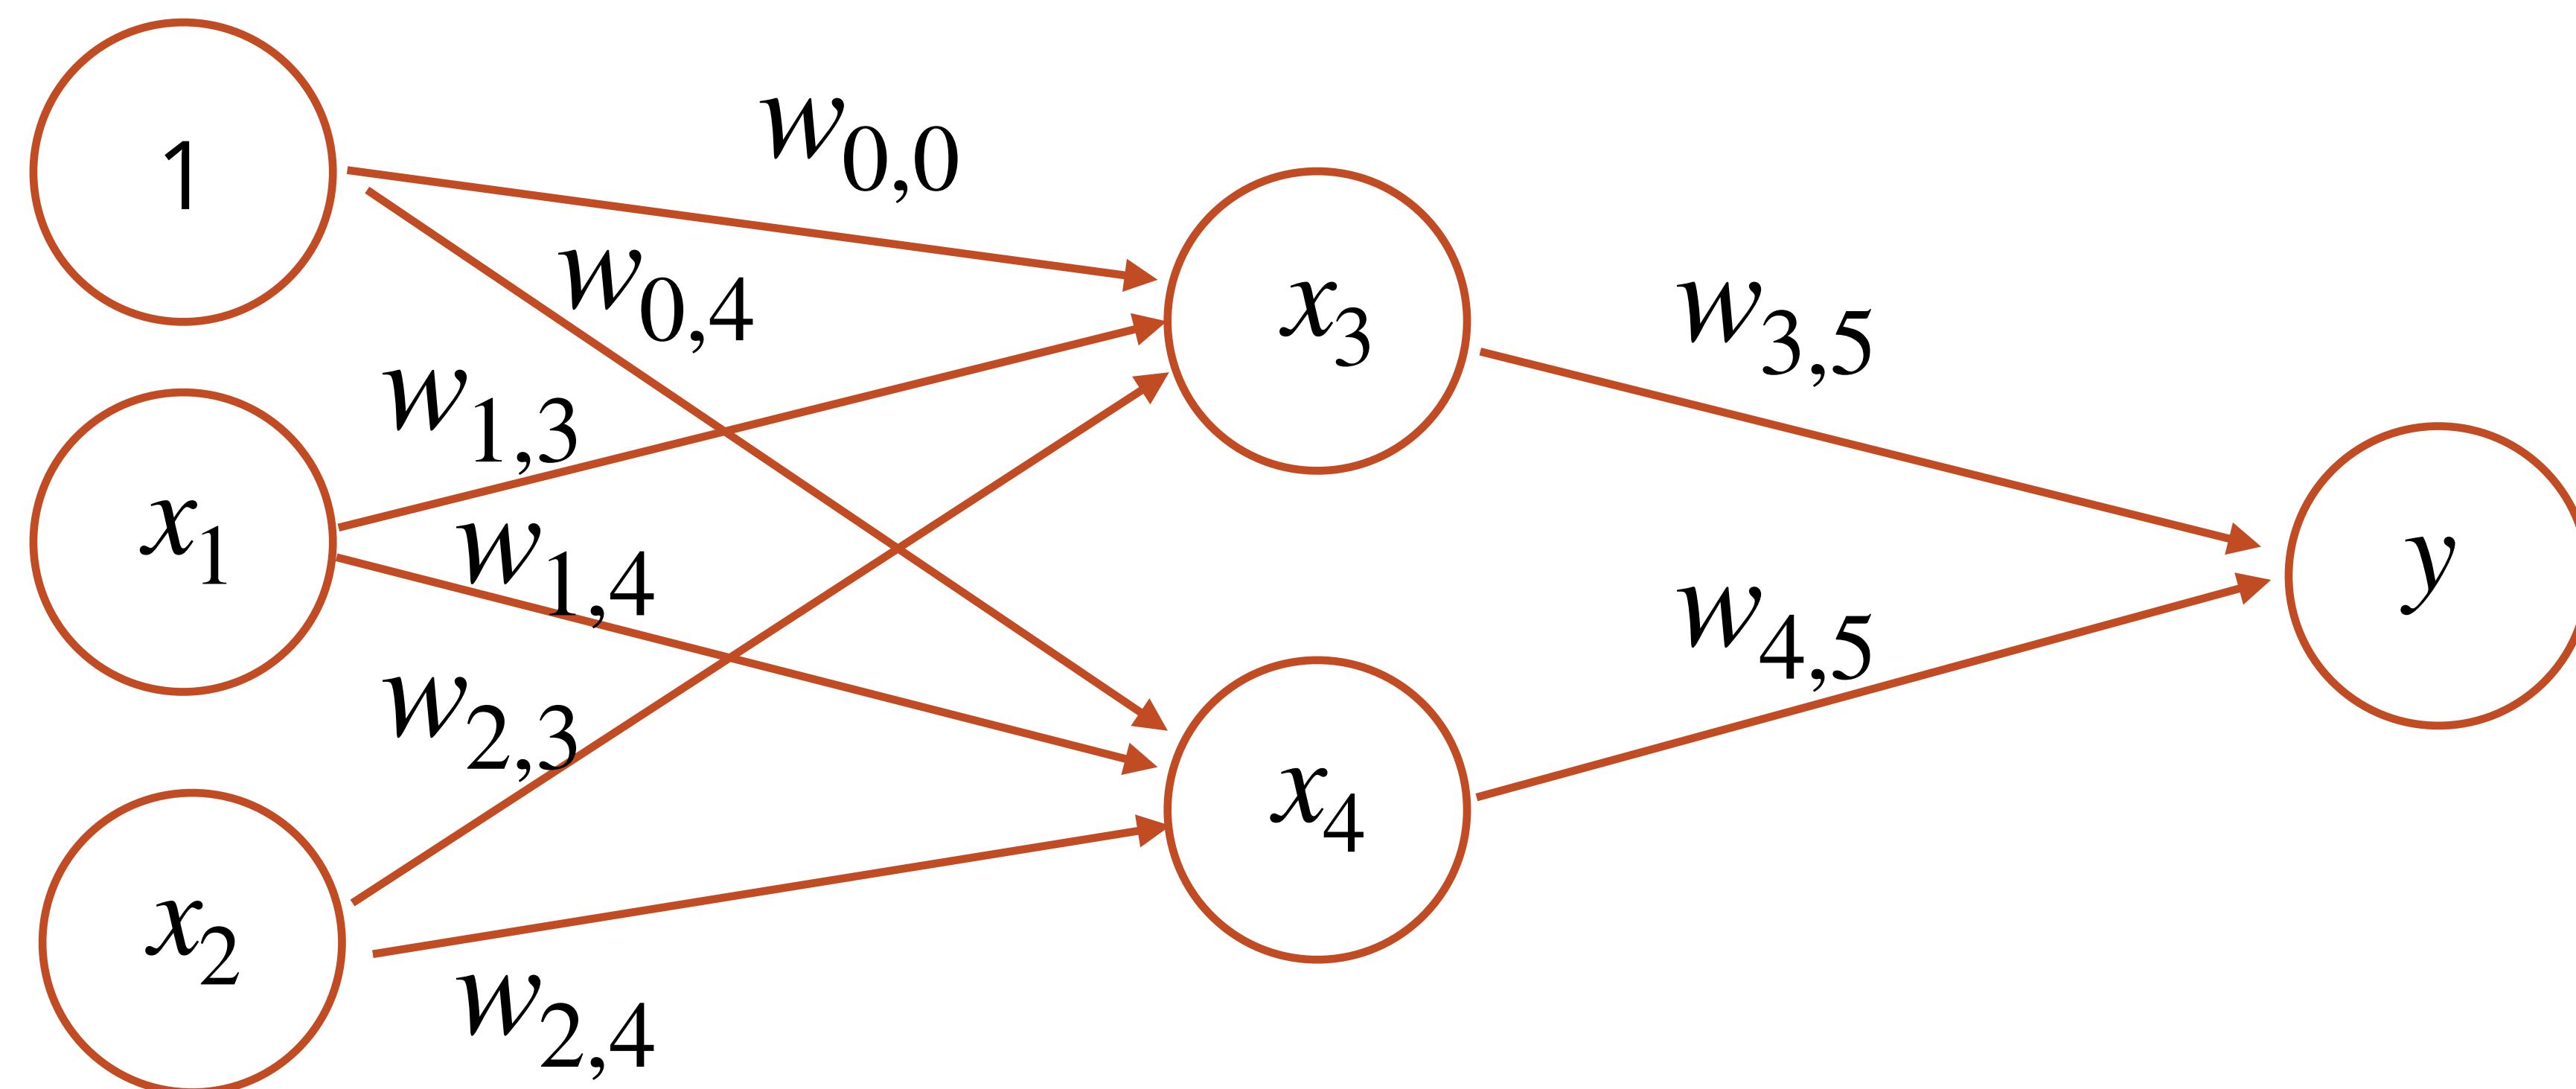
```
glm(y ~ x1 + x2, family = "binomial")
```

```
glm(y ~ x1 * x2, family = "binomial")
```

Need an x_1x_2 term!

x ₁	x ₂	y
0	0	0
0	1	1
1	0	1
1	1	0

How would a network solve xor?



x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0

Perceptrons

- 1. Simple neural networks generalize the Rescorla-Wagner model of associative learning**
- 2. Perceptrons are general-purpose linear classifiers.
They can solve lots of problems**
- 3. But they can't solve all problems...**